

# Computer Vision

DATA.ML.300, 5 study credits

Esa Rahtu

Unit of Computing Sciences, Tampere University

# Model estimation (fitting)

- Least-squares
- Robust fitting
- RANSAC
- Hough transform
  
- These topics are covered in Szeliski's book briefly, but more thoroughly in Chapter 17 of Forsyth & Ponce:
  - <http://courses.cs.washington.edu/courses/cse455/02wi/readings/book-7-revised-a-indx.pdf>

Acknowledgement: many slides from Svetlana Lazebnik, Derek Hoiem, Kristen Grauman, David Forsyth, Marc Pollefeys, and others (detailed credits on individual slides)

# Relevant reading

- These topics are covered in Szeliski's book briefly, but more thoroughly in the following books:
  - Chapter 17 of Forsyth & Ponce:
    - <http://cmuems.com/excap/readings/forsyth-ponce-computer-vision-a-modern-approach.pdf>
  - Chapter 4 of Hartley & Zisserman:
    - [http://cvrs.whu.edu.cn/downloads/ebooks/Multiple%20View%20Geometry%20in%20Computer%20Vision%20\(Second%20Edition\).pdf](http://cvrs.whu.edu.cn/downloads/ebooks/Multiple%20View%20Geometry%20in%20Computer%20Vision%20(Second%20Edition).pdf)


# Fitting

- We've learned how to detect edges, corners, blobs. Now what?
- We would like to form a higher-level, more compact representation of the features in the image by grouping multiple features according to a simple model



# Fitting

- Choose a parametric model to represent a set of features



simple model: lines



simple model: circles



complicated model: car

# Fitting: Issues

- Noise in the measured feature locations
- Extraneous data: clutter (outliers), multiple lines
- Missing data: occlusions

## Case study: Line detection




# Fitting: Overview

- If we know which points belong to the line, how do we find the “optimal” line parameters?
  - Least squares
- What if there are outliers?
  - Robust fitting, RANSAC
- What if there are many lines?
  - Voting methods: RANSAC, Hough transform
- What if we're not even sure it's a line?
  - Model selection (not covered)

# Least squares line fitting

- Data:  $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation:  $y_i = m x_i + b$
- Find  $(m, b)$  to minimize

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$






# Least squares line fitting

- Data:  $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation:  $y_i = m x_i + b$
- Find  $(m, b)$  to minimize

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$



$$E = \|Y - XB\|^2 \quad \text{where} \quad Y = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} \quad X = \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \quad B = \begin{bmatrix} m \\ b \end{bmatrix}$$

$$E = \|Y - XB\|^2 = (Y - XB)^T (Y - XB) = Y^T Y - 2(XB)^T Y + (XB)^T (XB)$$

$$\frac{dE}{dB} = 2X^T XB - 2X^T Y = 0$$

$$X^T XB = X^T Y$$

*Normal equations:*  
least squares solution to  $XB=Y$


# Problem with “vertical” least squares

- Not rotation-invariant
- Fails completely for vertical lines

# Total least squares

- Distance between point  $(x_i, y_i)$  and line  $ax+by=d$  ( $a^2+b^2=1$ ):  $|ax_i + by_i - d|$


$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$



# Total least squares


- Distance between point  $(x_i, y_i)$  and line  $ax+by=d$  ( $a^2+b^2=1$ ):  $|ax_i + by_i - d|$
- Find  $(a, b, d)$  to minimize the sum of squared perpendicular distances

$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$



# Total least squares

- Distance between point  $(x_i, y_i)$  and line  $ax+by=d$  ( $a^2+b^2=1$ ):  $|ax_i + by_i - d|$
- Find  $(a, b, d)$  to minimize the sum of squared perpendicular distances



$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$

$$\frac{\partial E}{\partial d} = \sum_{i=1}^n -2(ax_i + by_i - d) = 0 \quad d = \frac{a}{n} \sum_{i=1}^n x_i + \frac{b}{n} \sum_{i=1}^n y_i = a\bar{x} + b\bar{y}$$


$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = (UN)^T (UN)$$

$$\frac{dE}{dN} = 2(U^T U)N = 0$$

Solution to  $(U^T U)N = 0$ , subject to  $\|N\|^2 = 1$ : eigenvector of  $U^T U$  associated with the smallest eigenvalue

# Least squares: Robustness to noise


- Least squares fit to the red points:



# Least squares: Robustness to noise

- Least squares fit with an outlier:

Problem: squared error heavily penalizes outliers




# Robust estimators

- General approach: find model parameters  $\theta$  that minimize

$$\sum_i \rho(r_i(x_i, \theta); \sigma)$$

$r_i(x_i, \theta)$  – residual of  $i$ th point w.r.t. model parameters  $\theta$


$\rho$  – robust function with scale parameter  $\sigma$



The robust function  $\rho$  behaves like squared distance for small values of the residual  $u$  but saturates for larger values of  $u$




# Choosing the scale: Just right




The effect of the outlier is minimized

## Choosing the scale: Too small



The error value is almost the same for every point and the fit is very poor

## Choosing the scale: Too large



Behaves much the same as least squares

# Robust estimation: Details


- Robust fitting is a nonlinear optimization problem that must be solved iteratively
- Least squares solution can be used for initialization
- Scale of robust function should be chosen adaptively based on median residual

# RANSAC


- Robust fitting can deal with a few outliers – what if we have very many?
- Random sample consensus (RANSAC):  
Very general framework for model fitting in the presence of outliers
- Outline
  - Choose a small subset of points uniformly at random
  - Fit a model to that subset
  - Find all remaining points that are “close” to the model and reject the rest as outliers
  - Do this many times and choose the best model

M. A. Fischler, R. C. Bolles. [Random Sample Consensus: A Paradigm for Model Fitting with Applications to Image Analysis and Automated Cartography](#). Comm. of the ACM, Vol 24, pp 381-395, 1981.

# RANSAC for line fitting example




# RANSAC for line fitting example



Least-squares fit


# RANSAC for line fitting example



1. Randomly select minimal subset of points




# RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model

# RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function

# RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model

# RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

# RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

# RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

# RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

# RANSAC for line fitting

- Repeat  $N$  times:
- Draw  $s$  points uniformly at random
- Fit line to these  $s$  points
- Find inliers to this line among the remaining points (i.e., points whose distance from the line is less than  $t$ )
- If there are  $d$  or more inliers, accept the line and refit using all inliers



# Choosing the parameters

- Initial number of points  $s$ 
  - Typically minimum number needed to fit the model
- Distance threshold  $t$ 
  - Choose  $t$  so probability for inlier is  $p$  (e.g. 0.95)
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2 = 3.84\sigma^2$
- Number of samples  $N$ 
  - Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p=0.99$ ) (outlier ratio:  $e$ )

# Choosing the parameters

- Initial number of points **s**
  - Typically minimum number needed to fit the model
- Distance threshold **t**
  - Choose t so probability for inlier is p (e.g. 0.95)
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2=3.84\sigma^2$
- Number of samples **N**
  - Choose **N** so that, with probability p, at least one random sample is free from outliers (e.g. p=0.99) (outlier ratio: e)

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$


s	proportion of outliers <i>e</i>						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

# Choosing the parameters

- Initial number of points **s**
  - Typically minimum number needed to fit the model
- Distance threshold **t**
  - Choose t so probability for inlier is p (e.g. 0.95)
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2=3.84\sigma^2$
- Number of samples **N**
  - Choose **N** so that, with probability p, at least one random sample is free from outliers (e.g. p=0.99) (outlier ratio: e)

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$



# Choosing the parameters

- Initial number of points  $s$ 
  - Typically minimum number needed to fit the model
- Distance threshold  $t$ 
  - Choose  $t$  so probability for inlier is  $p$  (e.g. 0.95)
  - Zero-mean Gaussian noise with std. dev.  $\sigma$ :  $t^2 = 3.84\sigma^2$
- Number of samples  $N$ 
  - Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p=0.99$ ) (outlier ratio:  $e$ )
- Consensus set size  $d$ 
  - Should match expected inlier ratio

# Adaptively determining the number of samples


- Outlier ratio  $e$  is often unknown a priori, so pick worst case, e.g. 50%, and adapt if more inliers are found, e.g. 80% would yield  $e=0.2$
- Adaptive procedure:
  - $N=\infty$ , sample\_count =0
  - While  $N > \text{sample\_count}$ 
    - Choose a sample and count the number of inliers
    - If inlier ratio is highest of any found so far, set  
 $e = 1 - (\text{number of inliers})/(\text{total number of points})$
    - Recompute  $N$  from  $e$ :

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$

- Increment the sample\_count by 1

# RANSAC pros and cons

- Pros
  - Simple and general
  - Applicable to many different problems
  - Often works well in practice
- Cons
  - Lots of parameters to tune
  - Doesn't work well for low inlier ratios (too many iterations, or can fail completely)
  - Can't always get a good initialization of the model based on the minimum number of samples



# Fitting: Review


- Least squares
- Robust fitting
- RANSAC

# Fitting: Review

- If we know which points belong to the line, how do we find the “optimal” line parameters?
  - Least squares
- What if there are outliers?
  - Robust fitting, RANSAC
- What if there are many lines?
  - Voting methods: RANSAC, Hough transform



# Fitting: The Hough transform



# Voting schemes

- Let each feature vote for all the models that are compatible with it
- Hopefully the noise features will not vote consistently for any single model
- Missing data doesn't matter as long as there are enough features remaining to agree on a good model

# Hough transform

- An early type of voting scheme
- General outline:
  - Discretize parameter space into bins
  - For each feature point in the image, put a vote in every bin in the parameter space that could have generated this point
  - Find bins that have the most votes






Image space




Hough parameter space

P.V.C. Hough, *Machine Analysis of Bubble Chamber Pictures*, Proc. Int. Conf. High Energy Accelerators and Instrumentation, 1959


# Parameter space representation

- A line in the image corresponds to a point in Hough space




# Parameter space representation

- What does a point  $(x_0, y_0)$  in the image space map to in the Hough space?




# Parameter space representation

- What does a point  $(x_0, y_0)$  in the image space map to in the Hough space?
  - Answer: the solutions of  $b = -x_0m + y_0$
  - This is a line in Hough space




# Parameter space representation

- Where is the line that contains both  $(x_0, y_0)$  and  $(x_1, y_1)$ ?



# Parameter space representation

- Where is the line that contains both  $(x_0, y_0)$  and  $(x_1, y_1)$ ?
  - It is the intersection of the lines  $b = -x_0m + y_0$  and  $b = -x_1m + y_1$






# Parameter space representation

- Problems with the  $(m,b)$  space:
  - Unbounded parameter domains
  - Vertical lines require infinite  $m$

# Parameter space representation

- Problems with the (m,b) space:
  - Unbounded parameter domains
  - Vertical lines require infinite m
- Alternative: polar representation




Each point (x,y) will add a sinusoid in the  $(\theta, \rho)$  parameter space


# Algorithm outline

- Initialize accumulator  $H$  to all zeros
- For each feature point  $(x,y)$  in the image
  - For  $\theta = 0$  to  $180$ 
    - $\rho = x \cos \theta + y \sin \theta$
    - $H(\theta, \rho) = H(\theta, \rho) + 1$
- end
- end
- Find the value(s) of  $(\theta, \rho)$  where  $H(\theta, \rho)$  is a local maximum
  - The detected line in the image is given by
    - $\rho = x \cos \theta + y \sin \theta$


H: accumulator array (votes)



# Basic illustration




features





votes

[Hough transform demo](#)


## Square




# Several lines



# A more complicated image




# Effect of noise






# Effect of noise

- Peak gets fuzzy and hard to locate




features



votes


# Effect of noise

- Number of votes for a line of 20 points with increasing noise:




# Random points

- Uniform noise can lead to spurious peaks in the array




features



votes

# Random points

- As the level of uniform noise increases, the maximum number of votes increases too:




# Dealing with noise

- Choose a good grid / discretization
  - Too coarse: large votes obtained when too many different lines correspond to a single bucket
  - Too fine: miss lines because some points that are not exactly collinear cast votes for different buckets
- Increment neighboring bins (smoothing in accumulator array)
- Try to get rid of irrelevant features
  - E.g., take only edge points with significant gradient magnitude

# Incorporating image gradients

- Recall: when we detect an edge point, we also know its gradient direction
- But this means that the line is uniquely determined!
- Modified Hough transform:  
For each edge point (x,y)  
 $\theta$  = gradient orientation at (x,y)  
 $\rho = x \cos \theta + y \sin \theta$   
 $H(\theta, \rho) = H(\theta, \rho) + 1$   
end



$$\theta = \tan^{-1} \left( \frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

# Hough transform: Pros and cons

- Pros

- Can deal with non-locality and occlusion
- Can detect multiple instances of a model
- Some robustness to noise: noise points unlikely to contribute consistently to any single bin

- Cons

- Complexity of search time increases exponentially with the number of model parameters
- Non-target shapes can produce spurious peaks in parameter space
- It's hard to pick a good grid size