Computer Vision

SGN-45007, 5 study credits

Esa Rahtu
Laboratory of Signal Processing, Tampere University

Image classification

Image classification – Fundamental task in vision

Classify image using given set of discrete labels {dog, cat, truck, plane, …}

cat

Challenges – Semantic gap

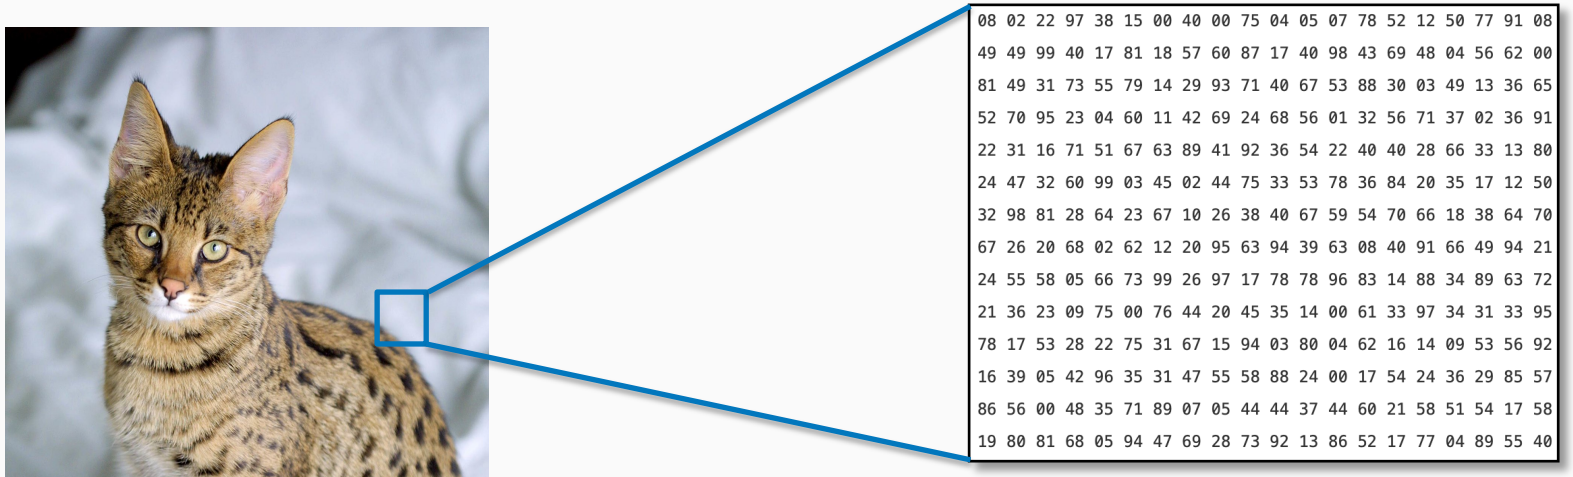

What you see

What the computer sees

For a computer the image is just an array of numbers (pixels) between [0, 255].
E.g. 1024 x 1024 x 3 (RGB)

Challenges – Viewpoint

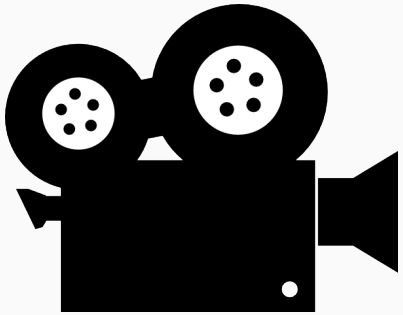

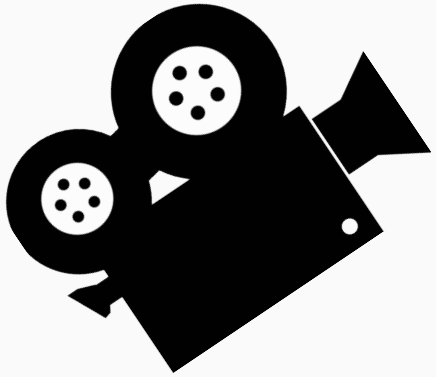

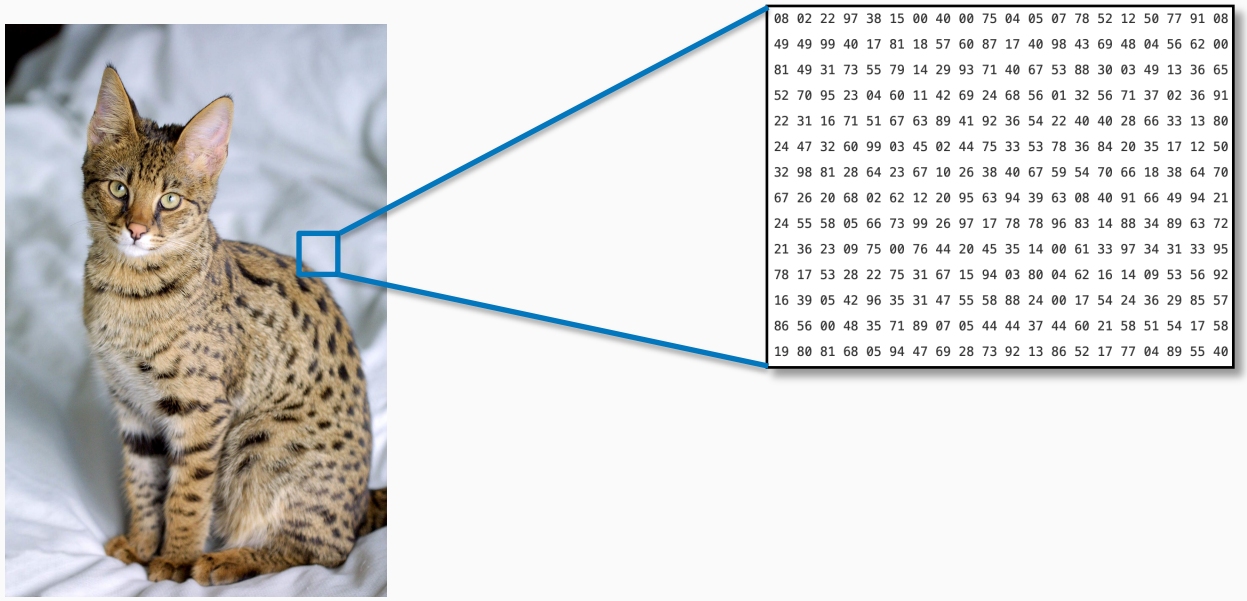

All pixel values change if camera moves!

Challenges – Illumination

This image is CC0 1.0 public domain

This image is CC0 1.0 public domain

This image is CC0 1.0 public domain

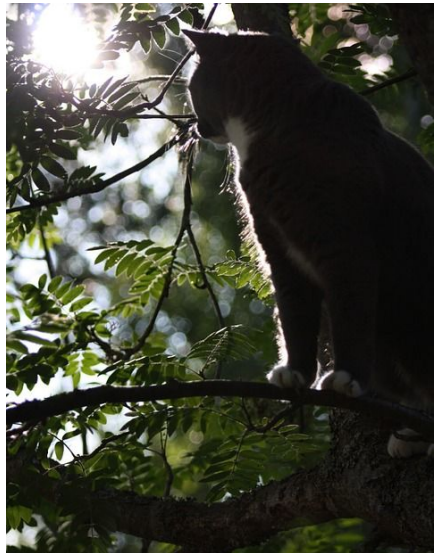

This image is CC0 1.0 public domain

Challenges – Deformation

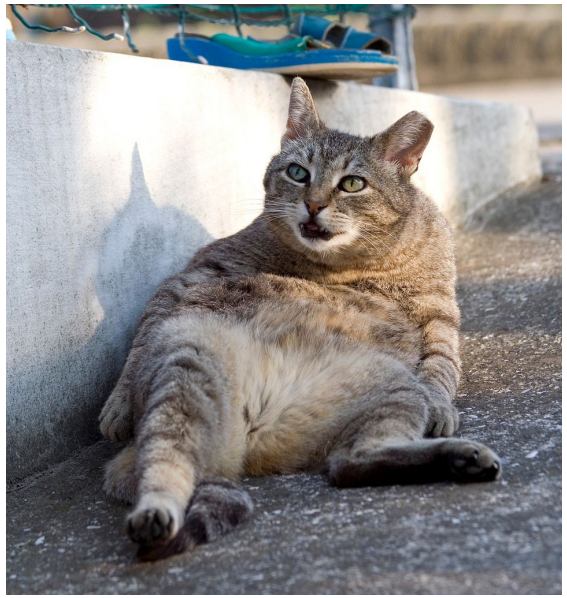

This image by Umberto Salvagnin
is licensed under CC-BY 2.0

This image by Umberto Salvagnin
is licensed under CC-BY 2.0

This image by sare bear is
licensed under CC-BY 2.0

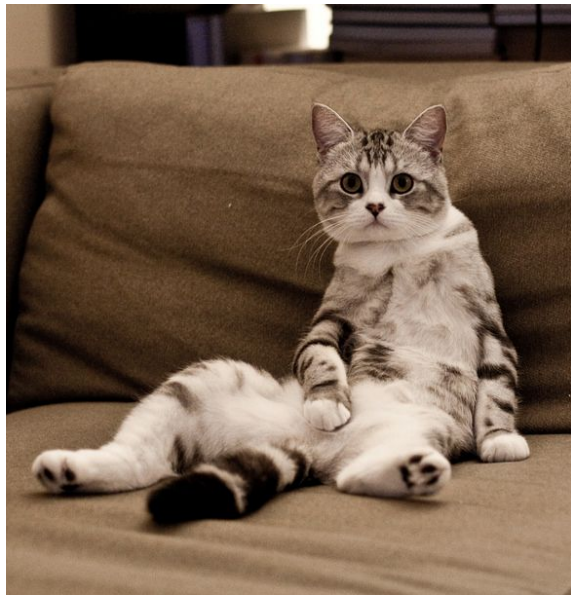

This image by Tom Thai is
licensed under CC-BY 2.0

Challenges – Occlusion

This image is CC0 1.0 public domain

This image is CC0 1.0 public domain

This image by jonsson is licensed under CC-BY 2.0

Challenges – Background clutter

This image is CC0 1.0 public domain

This image is CC0 1.0 public domain

Challenges – Intraclass variation

This image is CC0 1.0 public domain

How to build image classifier?

- Extremely hard task!
- No obvious way to ”hard-code” the algorithm for recognizing image categories (e.g. cat). (c.f. other problems like sorting a list of numbers)

```
def classify_image(image):
  # Some magic here?
  return class_label
```

Figure from Fei-Fei Li

How to build image classifier?

- Extremely hard task!
- No obvious way to ”hard-code” the algorithm for recognizing image categories (e.g. cat). (c.f. other problems like sorting a list of numbers)
- Many attempts have been made over the years:

Image classification with machine learning

- Instead of hard coding, learn from examples!
- Data driven approach:

Figure from Fei-Fei Li

Image classification with machine learning

- Instead of hard coding, learn from examples!
- Data driven approach:

1. Collect a dataset of images and labels

Example training set

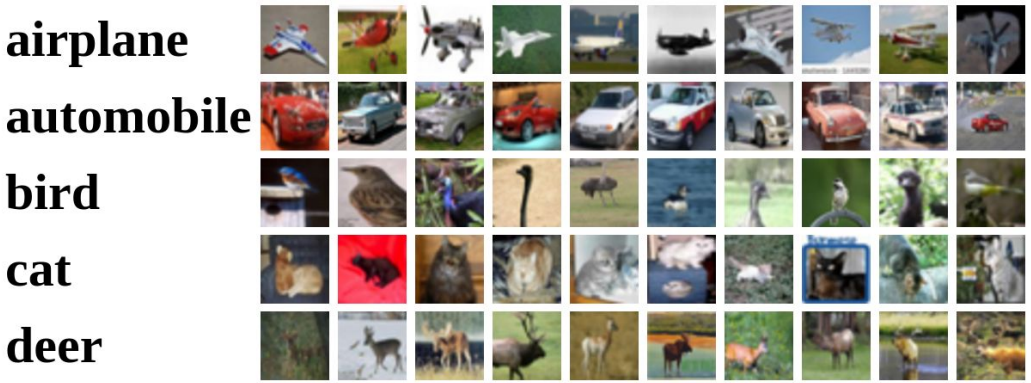

Figure from Fei-Fei Li

Image classification with machine learning

- Instead of hard coding, learn from examples!
- Data driven approach:

1. Collect a dataset of images and labels
2. Use machine learning to train a classifier

```
def train(images, labels):
  # Machine learning!
  return model
```

Example training set

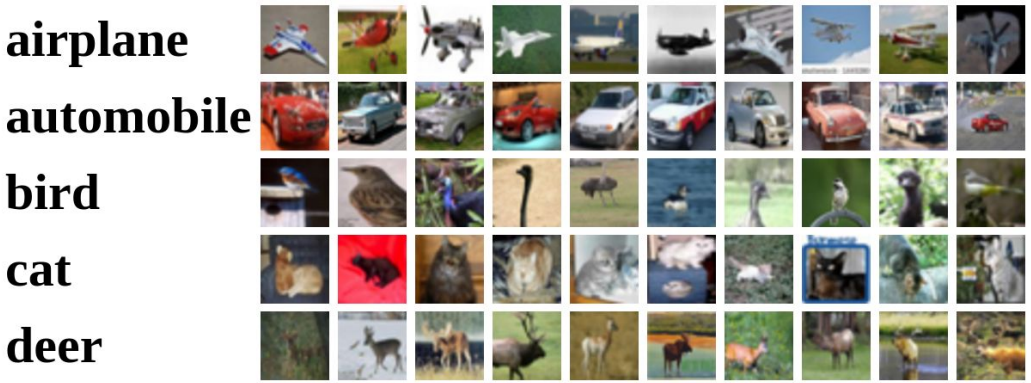

Figure from Fei-Fei Li

Image classification with machine learning

- Instead of hard coding, learn from examples!
- Data driven approach:

1. Collect a dataset of images and labels
2. Use machine learning to train a classifier
3. Evaluate the classifier on new images

Example training set

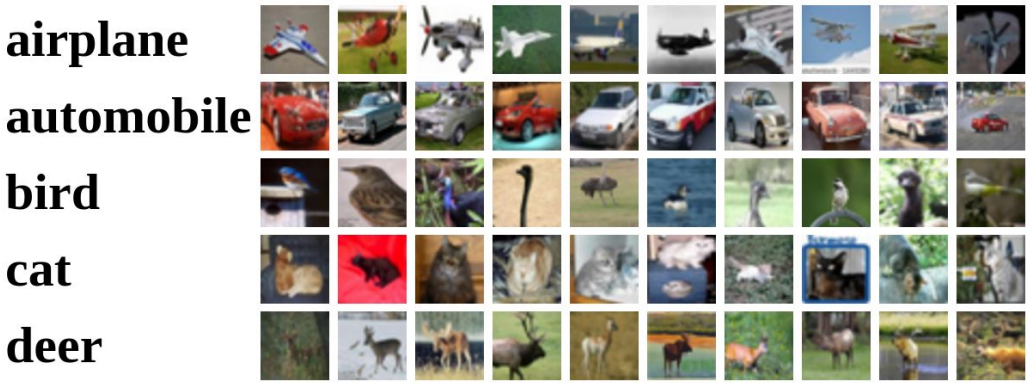


```
def train(images, labels):
  # Machine learning!
  return model
```

```
def predict(model, test_images):
  # Use model to predict labels
  return test_labels
```

Figure from Fei-Fei Li

Nearest neighbour classifier

```
def train(images, labels):
  # Machine learning!
  return model
```

Memorize all data and labels

Figure from Fei-Fei Li

Nearest neighbour classifier

```
def train(images, labels):
  # Machine learning!
  return model
```

Memorize all data and labels

```
def predict(model, test_images):
  # Use model to predict labels
  return test_labels
```

Predict the label of the most similar training image

Figure from Fei-Fei Li

Example: CIFAR 10

10 classes
50 000 training images
10 000 testing images

Learning multiple layers of features from tiny images, Krizhevsky, Technical report 2009

Example: CIFAR 10

10 classes
50 000 training images
10 000 testing images

Test images and nearest neighbours

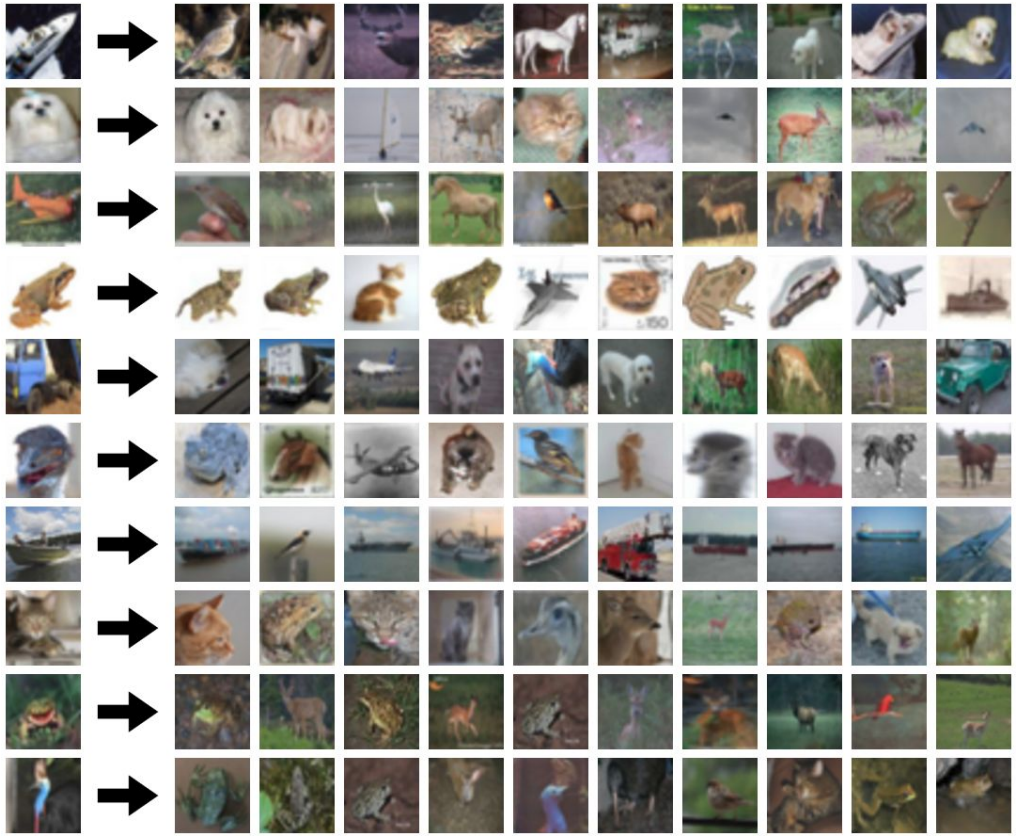

Learning multiple layers of features from tiny images, Krizhevsky, Technical report 2009

How to determine the most similar image?

- Determine distance metric between images
- L1 distance: $$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

Test image

56	32	10	18
90	23	128	133
24	26	178	200
2	0	255	220

-

Training image

10	20	24	17
8	10	89	100
12	16	178	170
4	32	233	112

=

Pixel-wise absolute difference

46	12	14	1
82	13	39	33
12	10	0	30
2	32	22	108

$\xrightarrow{\text{add}}$ 456

Figure from Fei-Fei Li

How does this look like?

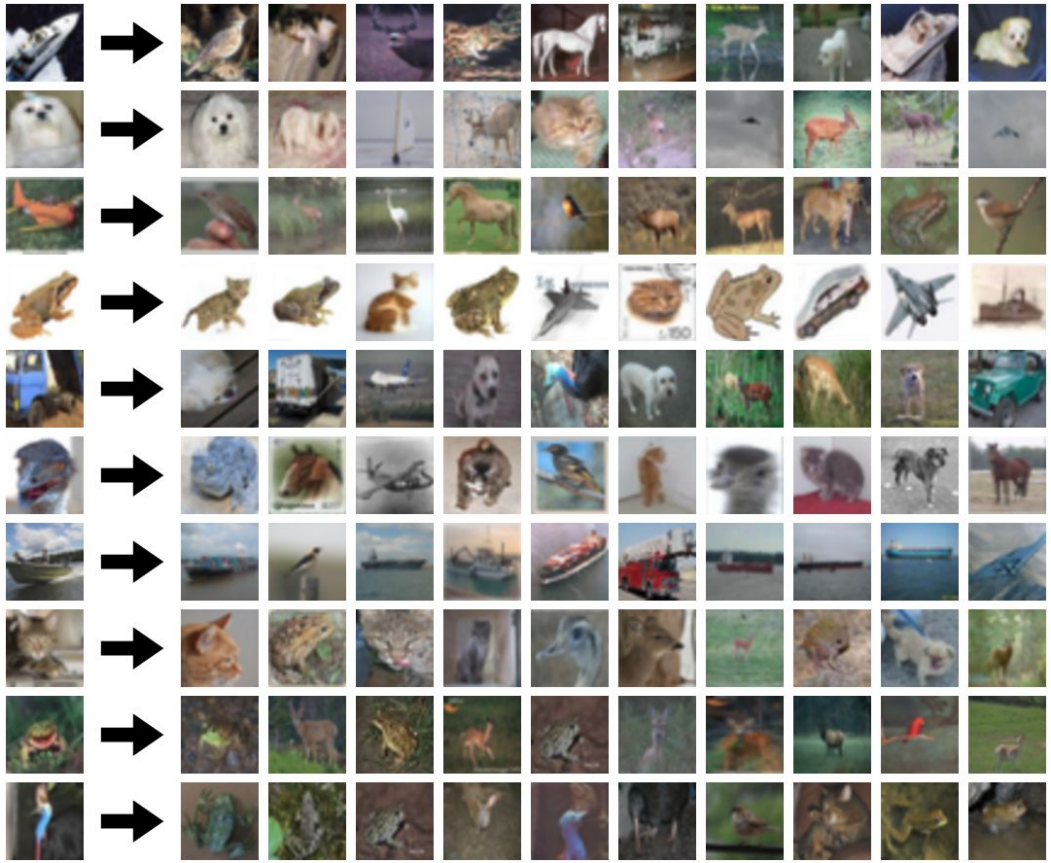

How does this look like?

How does this look like?

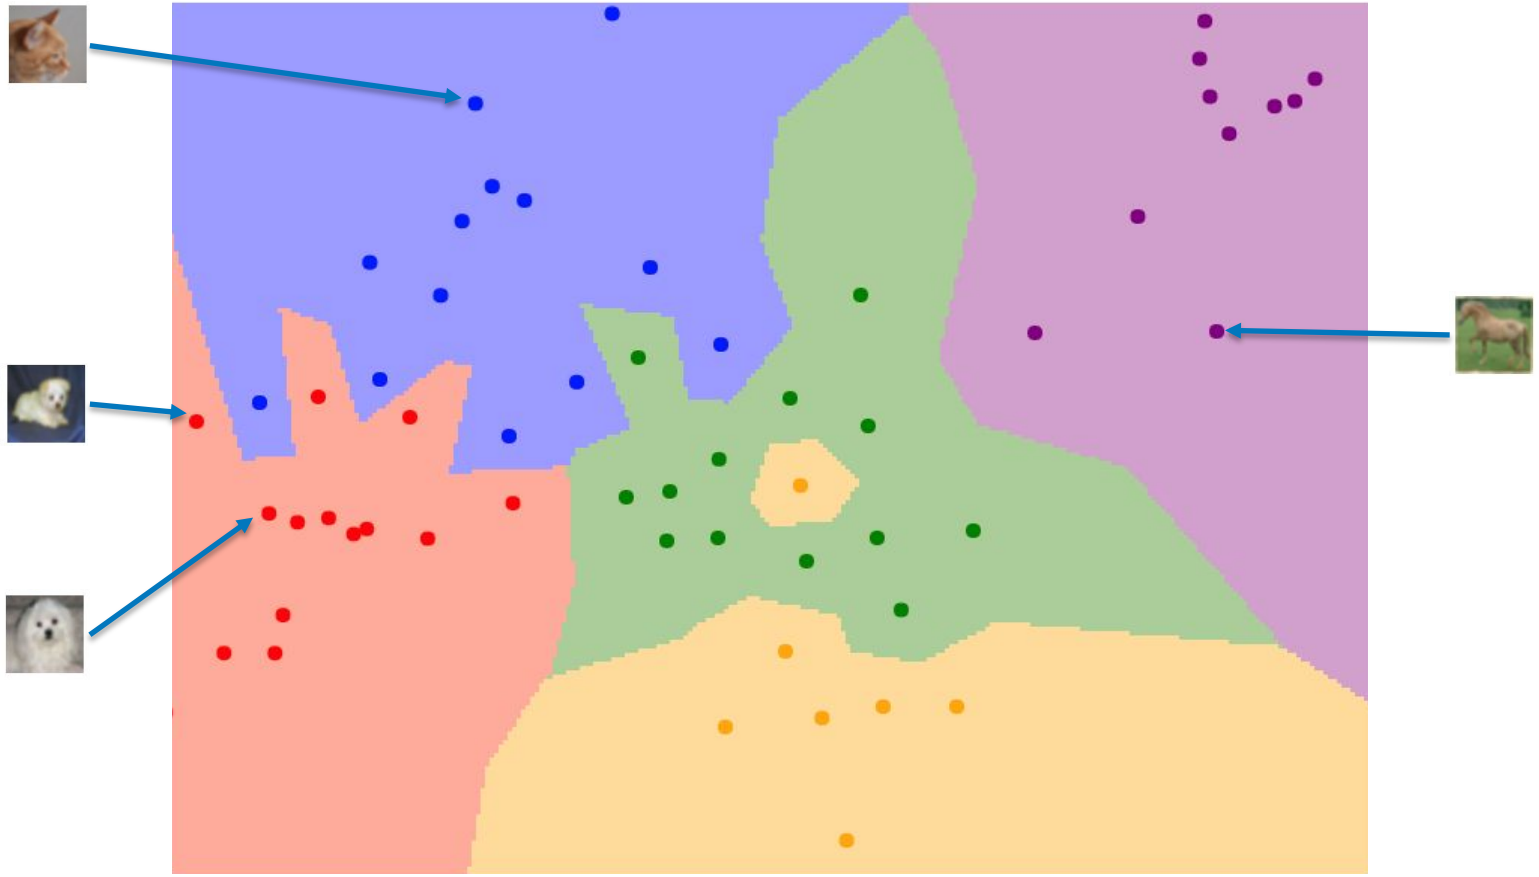

Credits: Fei-Fei Li

How does this look like?

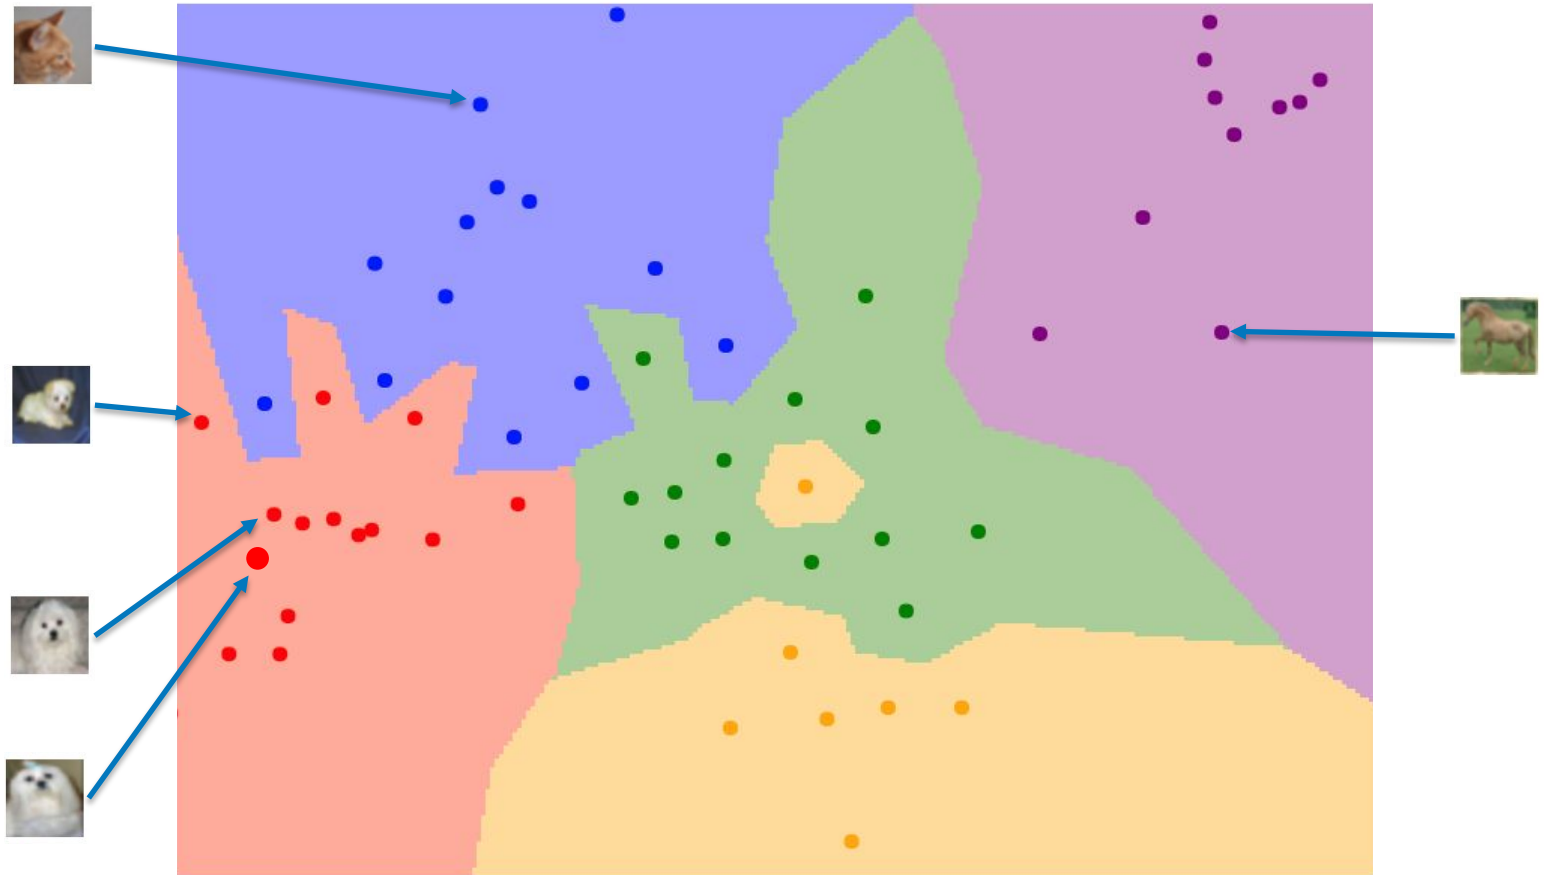

Credits: Fei-Fei Li

K-Nearest neighbours

- Instead of copying label from the nearest sample, take **majority vote** from K closest

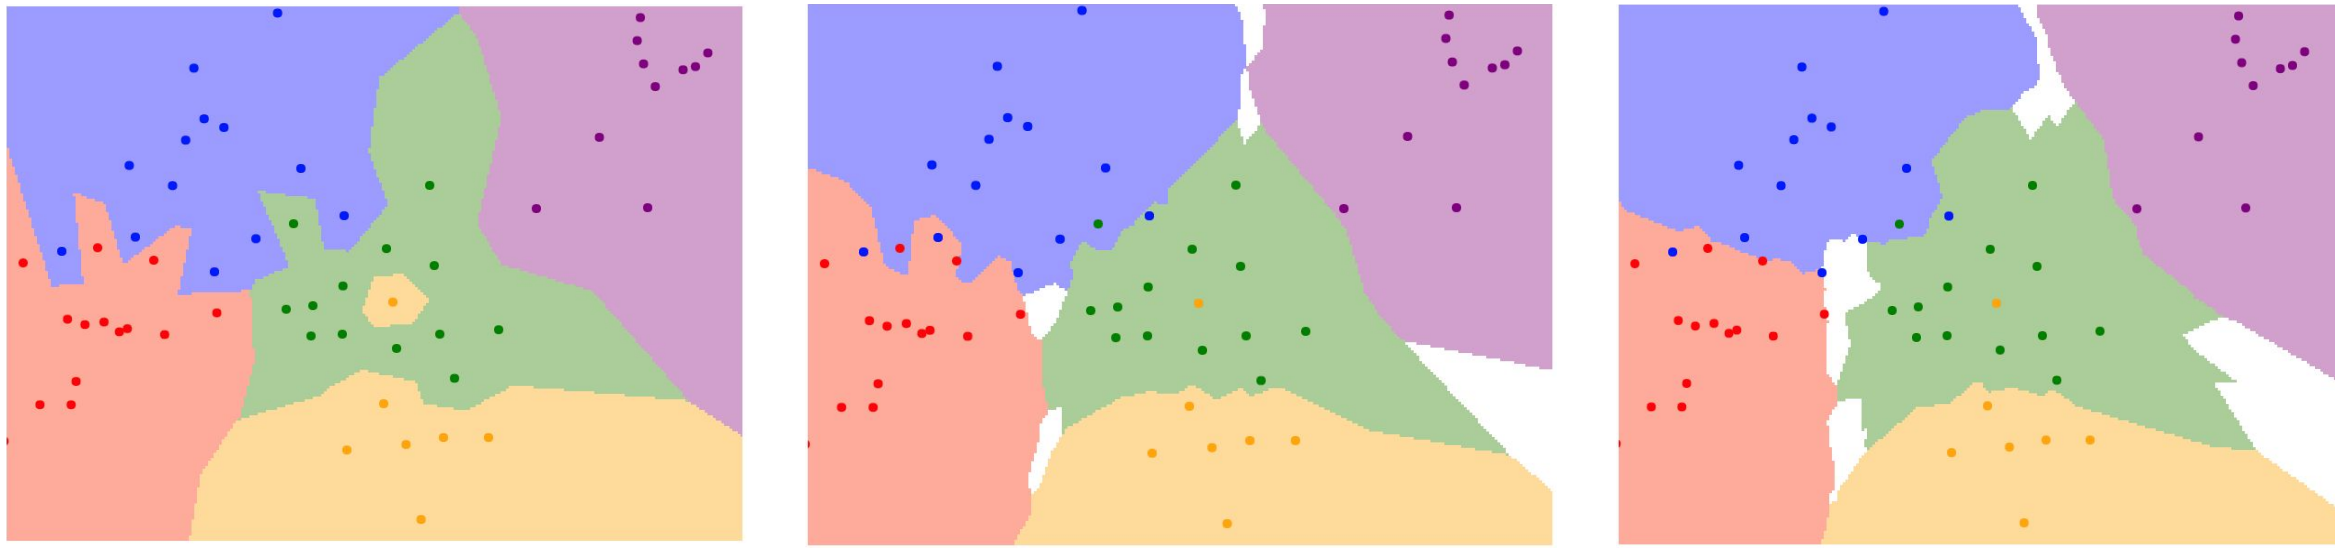

K = 1 K = 3 K = 5

Figure from Fei-Fei Li

How to select hyper parameters?

- How to choose an optimal value for K?
- How to choose optimal distance metric?
- These are **hyperparameters**: choices about the algorithm that we set rather than learn.
- Very problem dependent and often need to optimized by trial and error.

How to select hyper parameters?

- Common approach for parameter setting:

1. Split data into training, validation, and test sets
2. Choose hyperparameters on validation set
3. Evaluate generalization using test set

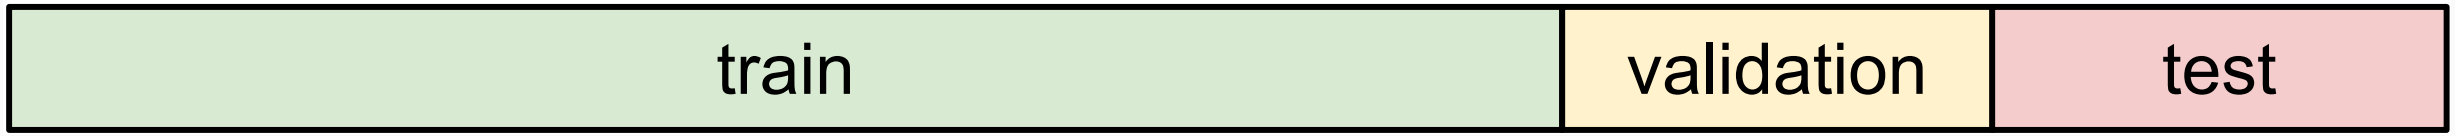

Figure from Fei-Fei Li

Image classification with neural networks

Recap: Supervised learning

```
def train(images, labels):
  # Machine learning!
  return model
```

```
def predict(model, test_images):
  # Use model to predict labels
  return test_labels
```

Recap: Supervised learning

Functions $\mathcal{F}$

$$f : \mathcal{X} \to \mathcal{Y}$$

Training data

$$\{(x_i, y_i) \in \mathcal{X} \times \mathcal{Y}\}$$

```
def train(images, labels):
  # Machine learning!
  return model
```

$$\text{find } \hat{f} \in \mathcal{F}$$
$$\text{s.t. } y_i \approx \hat{f}(x_i)$$

Learning machine

```
def predict(model, test_images):
  # Use model to predict labels
  return test_labels
```

$$y = \hat{f}(x)$$

New data

$$x$$

Recap: Supervised learning

Functions $\mathcal{F}$

$f : \mathcal{X} \to \mathcal{Y}$

Training data

$\{(x_i, y_i) \in \mathcal{X} \times \mathcal{Y}\}$

LEARNING

$$\text{find } \hat{f} \in \mathcal{F}$$
$$\text{s.t. } y_i \approx \hat{f}(x_i)$$

Find f that minimises given **loss** for data

$$\frac{1}{N}\sum_{i=1}^{N}\underbrace{[\mathbf{y}_i \neq f(\mathbf{x}_i)]}_{\text{loss function}}$$

New data

PREDICTION

$$y = \hat{f}(x)$$

$$x$$

Deep convolutional networks for image classification

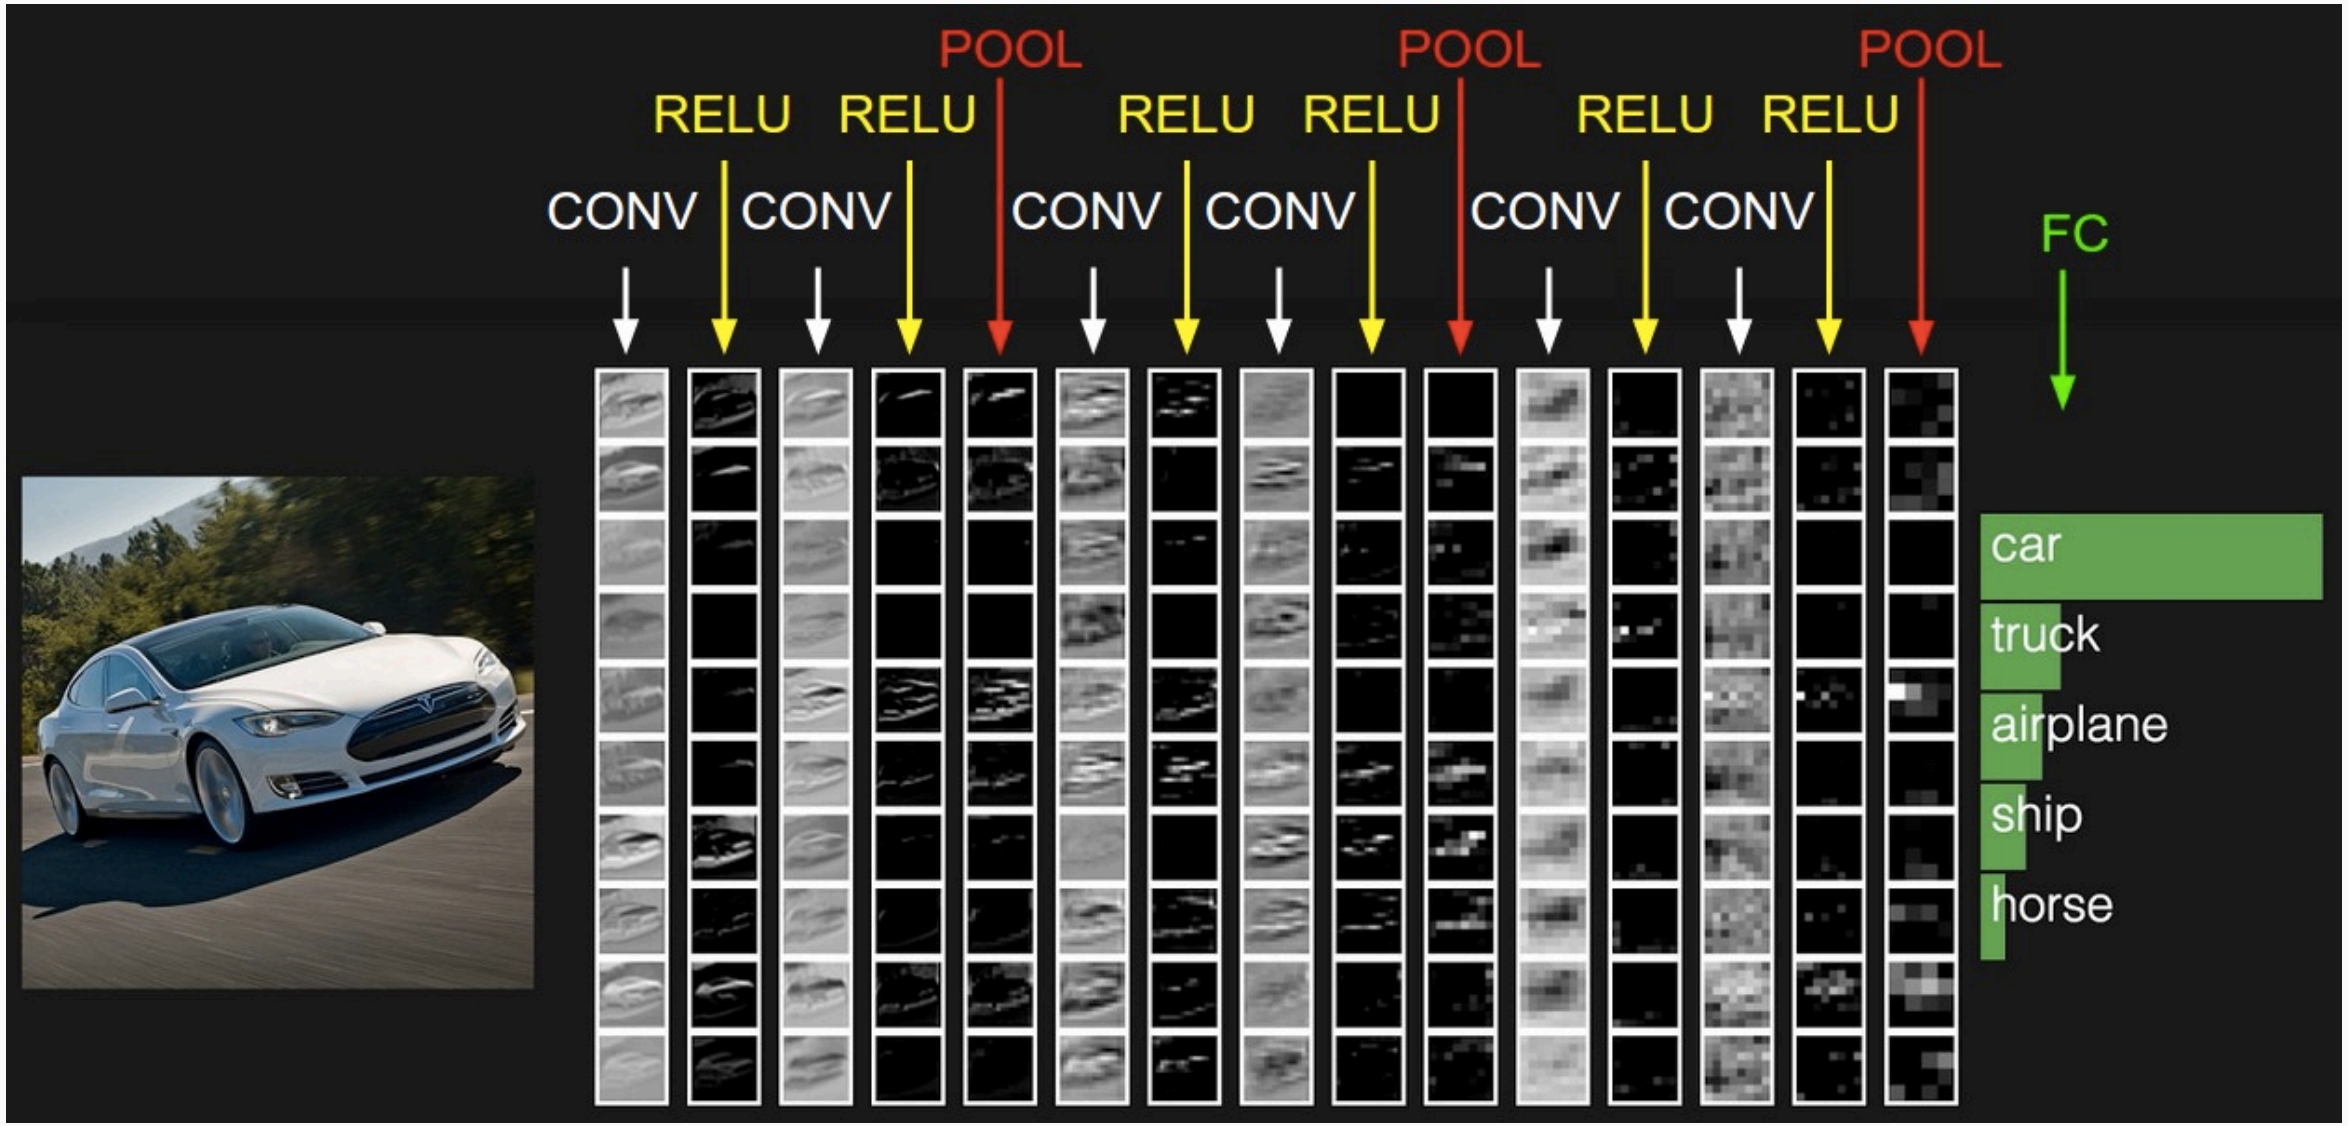

Deep learning

- Learn a feature hierarchy all the way from pixels to classifier
- Each layer extracts features from the output of previous layer
- Train all layers jointly

Background: Linear classifiers

- When the data is linearly separable, there may be more than one separator (hyperplane)

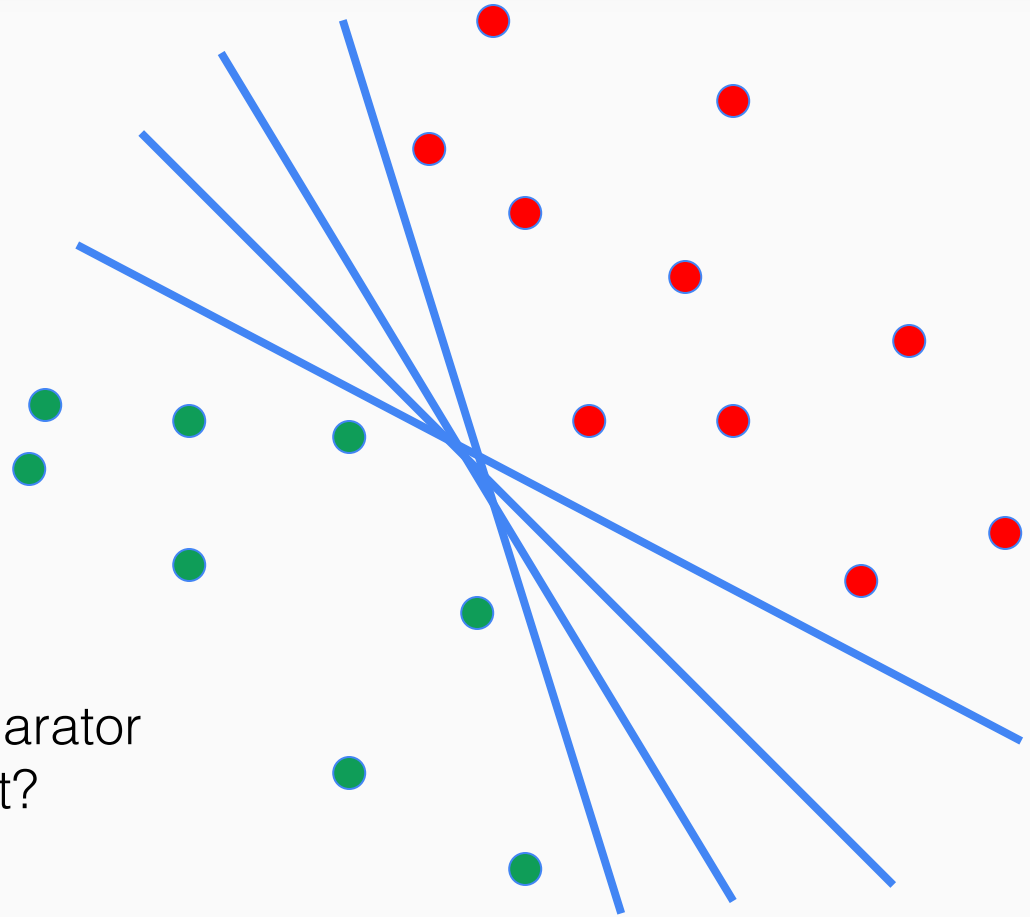

Which separator is best?

Perceptron

Perceptron

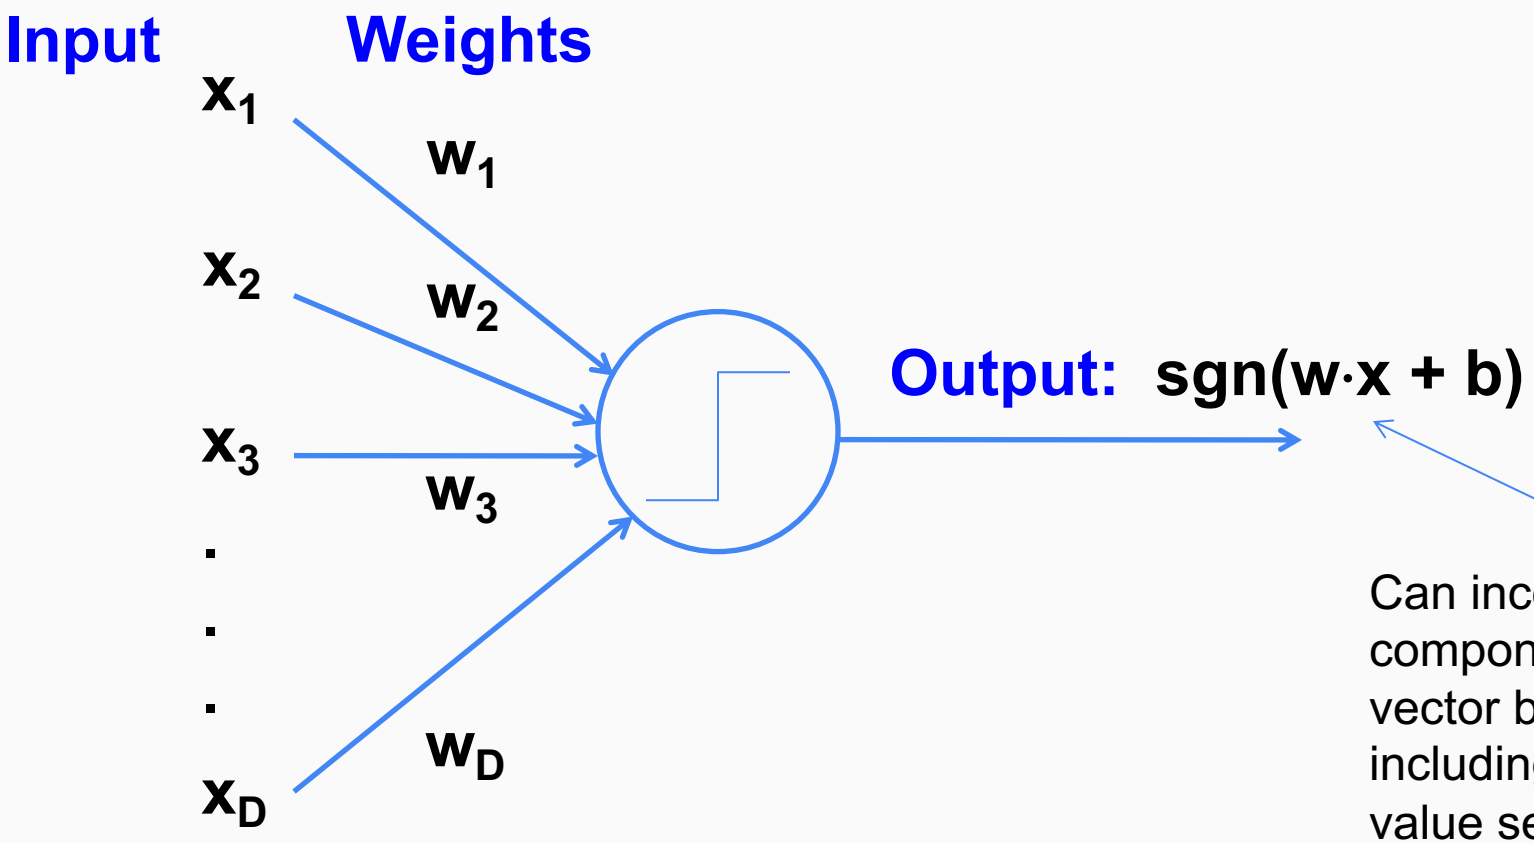

Can incorporate bias as component of the weight vector by always including a feature with value set to 1

Loose inspiration: Human neurons

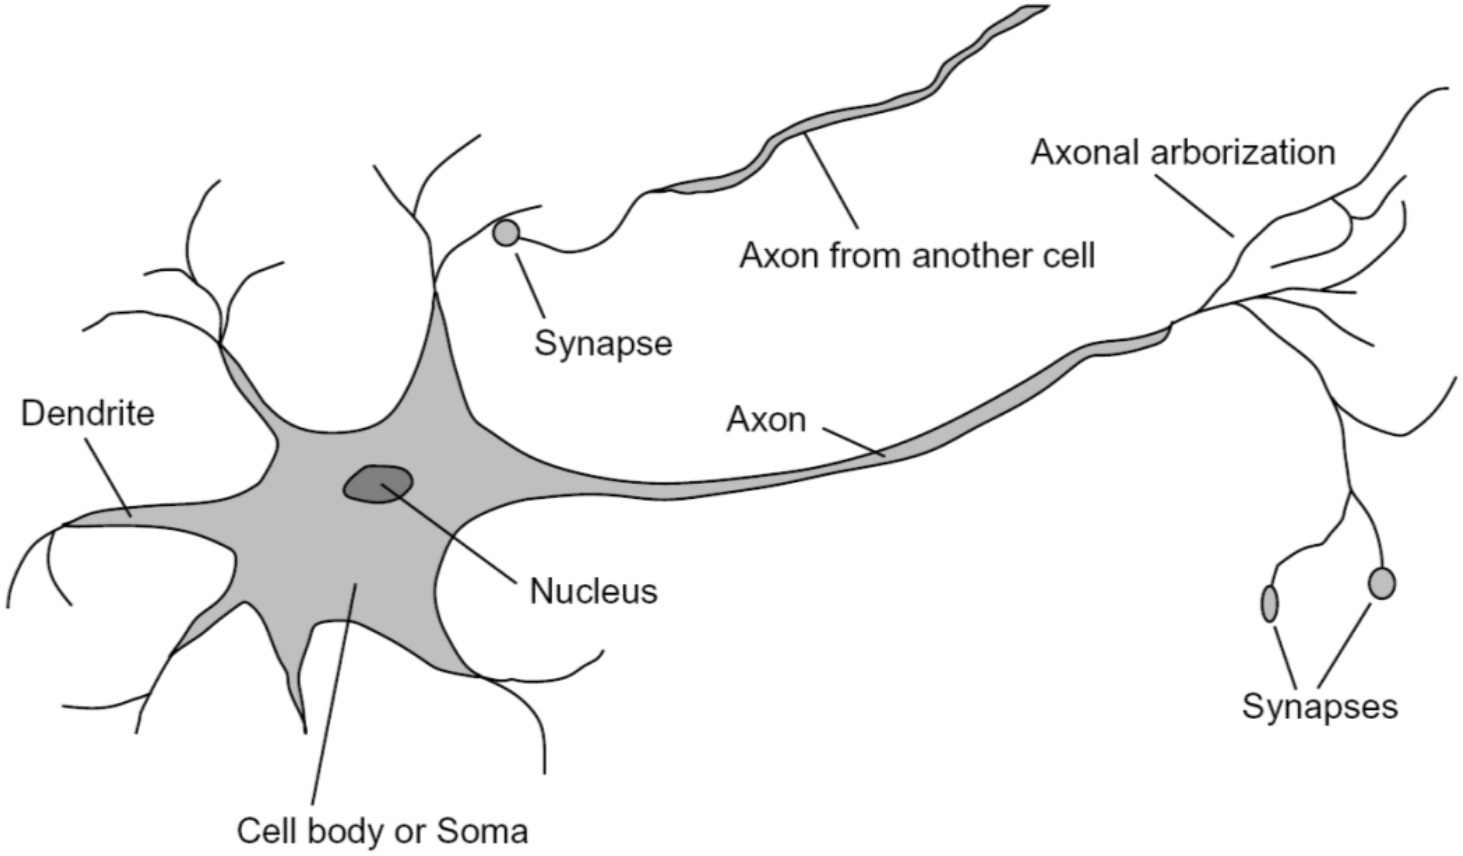

History

1958 New York Times…

NEW NAVY DEVICE LEARNS BY DOING

Psychologist Shows Embryo of Computer Designed to Read and Grow Wiser

WASHINGTON, July 7 (UPI) —The Navy revealed the embryo of an electronic computer today that it expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence.

The embryo—the Weather Bureau's $2,000,000 "704" computer—learned to differentiate between right and left after fifty attempts in the Navy's demonstration for newsmen.

The service said it would use this principle to build the first of its Perceptron thinking machines that will be able to read and write. It is expected to be finished in about a year at a cost of $100,000.

Dr. Frank Rosenblatt, designer of the Perceptron, conducted the demonstration. He said the machine would be the first device to think as the human brain. As do human beings, Perceptron will make mistakes at first, but will grow wiser as it gains experience, he said.

Dr. Rosenblatt, a research psychologist at the Cornell Aeronautical Laboratory, Buffalo, said Perceptrons might be fired to the planets as mechanical space explorers.

Without Human Controls

The Navy said the perceptron would be the first non-living mechanism "capable of receiving, recognizing and identifying its surroundings without any human training or control."

The "brain" is designed to remember images and information it has perceived itself. Ordinary computers remember only what is fed into them on punch cards or magnetic tape.

Later Perceptrons will be able to recognize people and call out their names and instantly translate speech in one language to speech or writing in another language, it was predicted.

Mr. Rosenblatt said in principle it would be possible to build brains that could reproduce themselves on an assembly line and which would be conscious of their existence.

In today's demonstration, the "704" was fed two cards, one with squares marked on the left side and the other with squares on the right side.

Learns by Doing

In the first fifty trials, the machine made no distinction between them. It then started registering a "Q" for the left squares and "O" for the right squares.

Dr. Rosenblatt said he could explain why the machine learned only in highly technical terms. But he said the computer had undergone a "self-induced change in the wiring diagram."

The first Perceptron will have about 1,000 electronic "association cells" receiving electrical impulses from an eye-like scanning device with 400 photo-cells. The human brain has 10,000,000,000 responsive cells, including 100,000,000 connections with the eyes.

Perceptron update rule

- Initialize weights randomly
- Cycle through training examples in multiple passes (epochs)
- For each training instance x with label y:
  - Classify with current weights: $y' = \text{sgn}(w \cdot x)$
  - Update weights: $w \leftarrow w + \alpha(y-y')x$
- α is a learning rate that should decay as 1/t (t is the epoch)
- What happens if y' is correct?
- Otherwise, consider what happens to individual weights
  $wi \leftarrow w_i + \alpha(y-y')x_i$
  - If y = 1 and y' = -1, w_i will be increased if x_i is positive or decreased if x_i is negative → w·x will get bigger
  - If y = -1 and y' = 1, w_i will be decreased if x_i is positive or increased if x_i is negative → w·x will get smaller

Convergence of perceptron update rule

- **Linearly separable data:** converges to a perfect solution
- **Non-separable data:** converges to a minimum-error solution assuming learning rate decays as O(1/t) and examples are presented in random sequence

Multi-Layer Neural Networks

- Network with a hidden layer:

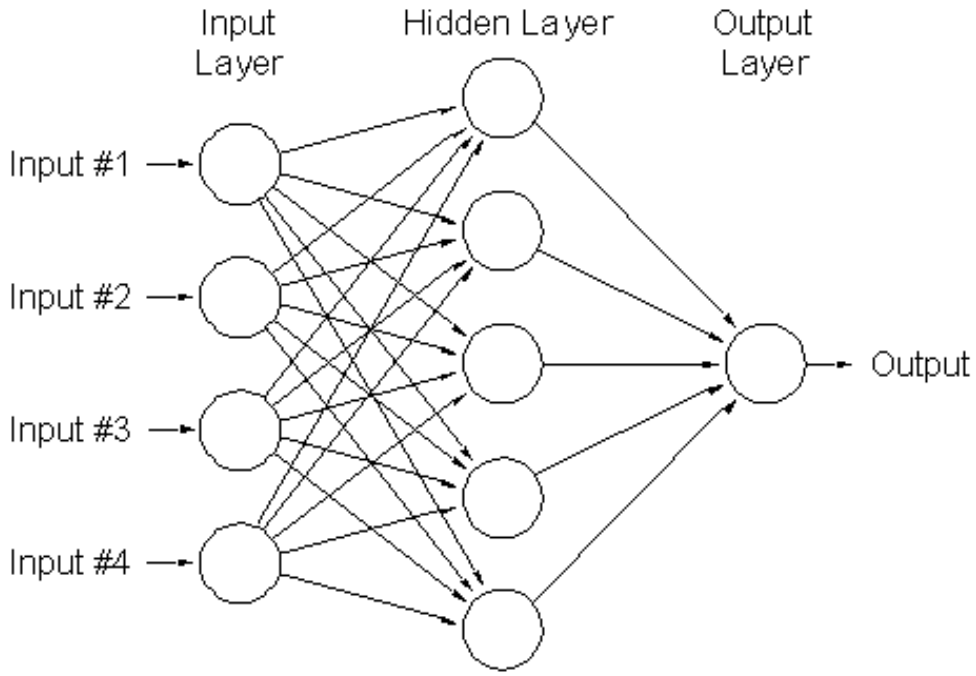

- Can represent nonlinear functions (provided each perceptron has a nonlinearity)

Multi-Layer Neural Networks

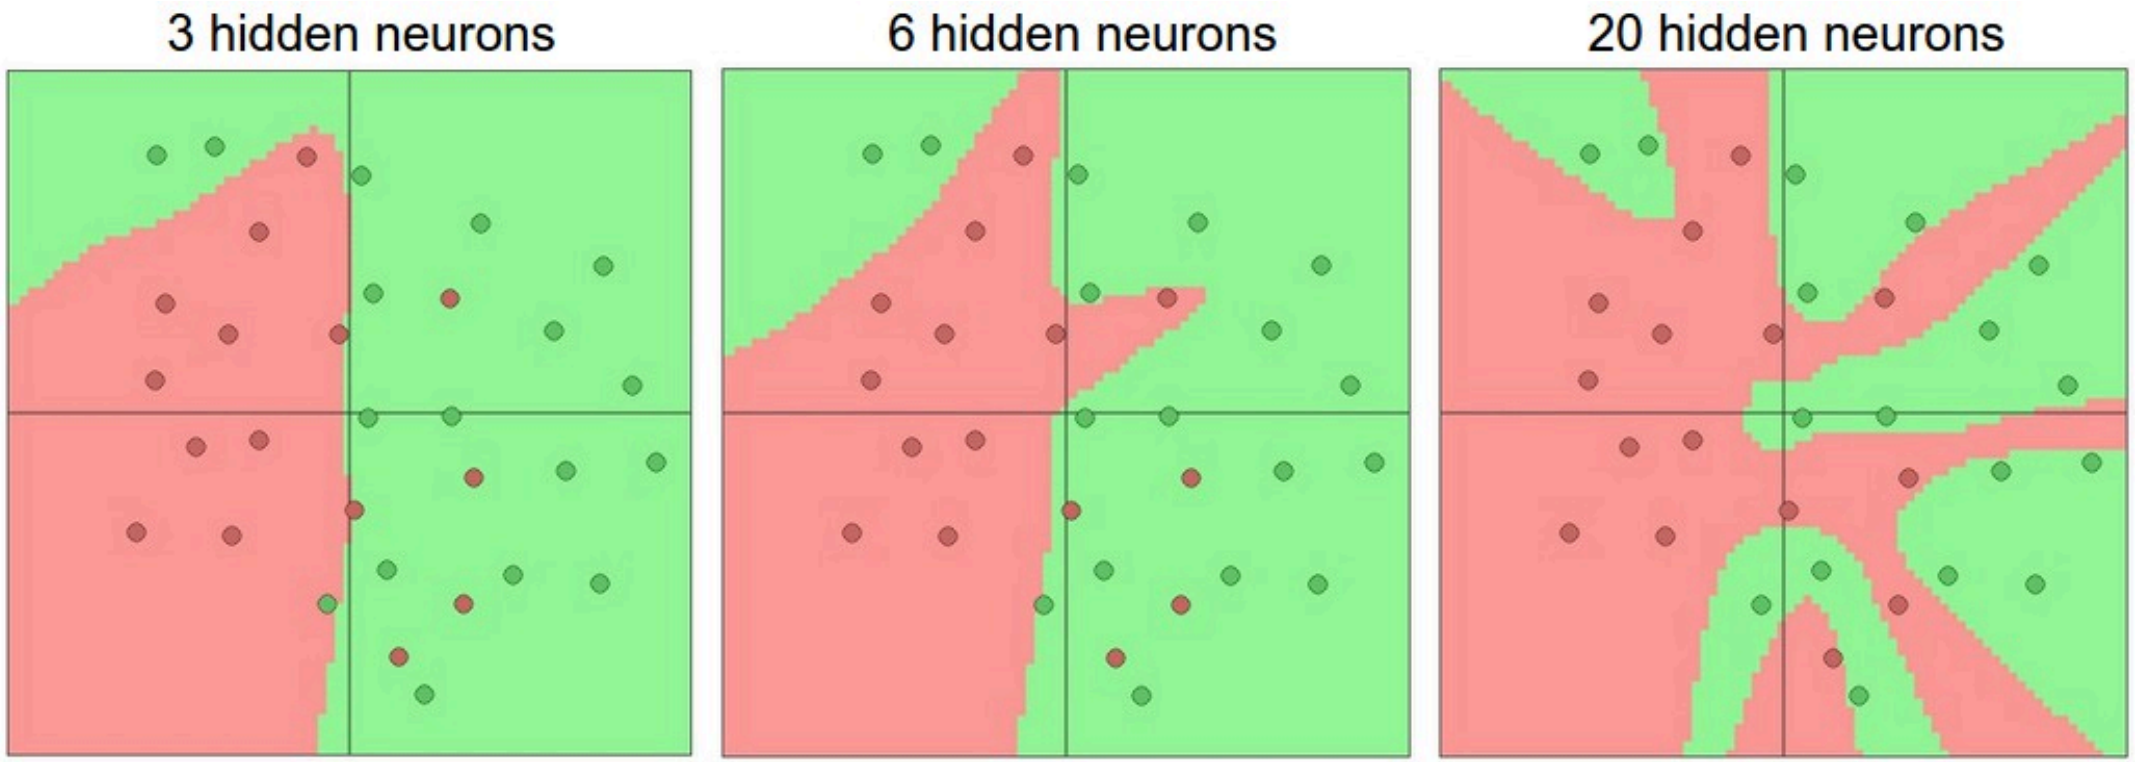

Source: http://cs231n.github.io/neural-networks-1/

Multi-Layer Neural Networks

- Beyond a single hidden layer:

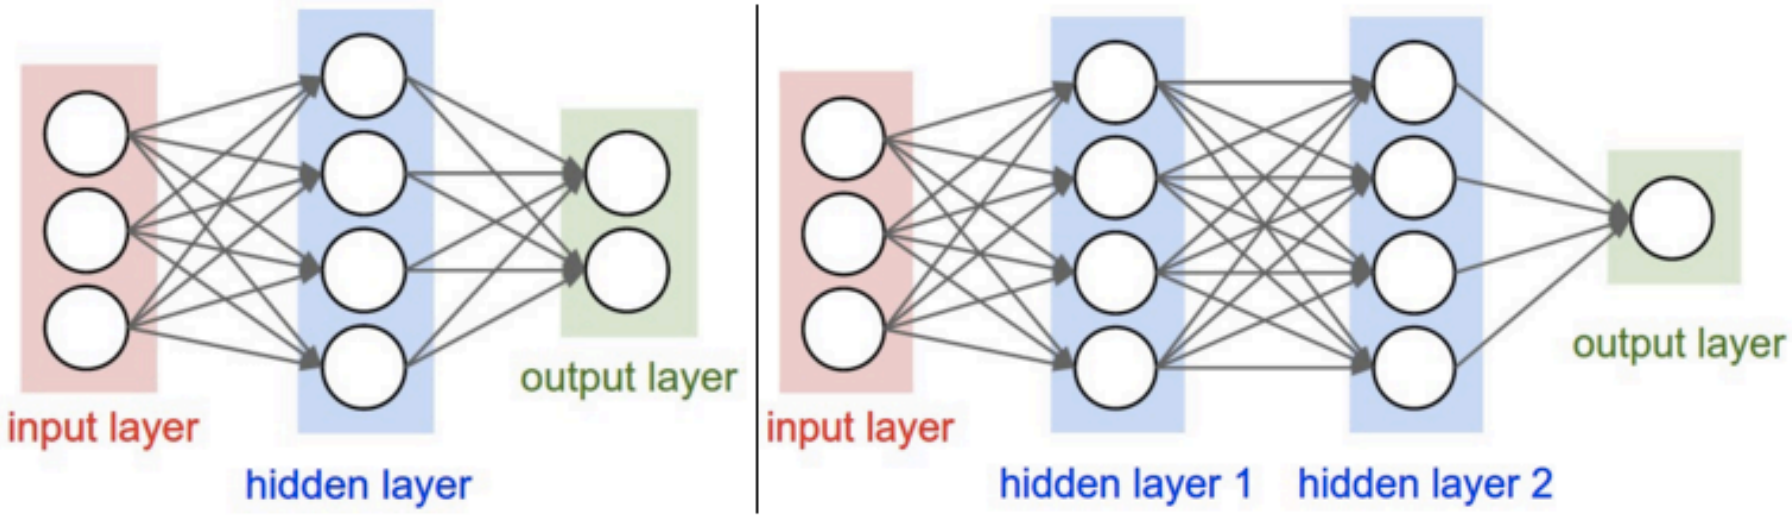

Figure source: http://cs231n.github.io/neural-networks-1/

Training of multi-layer networks

- Find network weights to minimize the error between true and estimated labels of training examples:

$$E(\mathbf{w}) = \sum_{j=1}^{N} \left(y_j - f_{\mathbf{w}}(\mathbf{x}_j) \right)^2$$

- Update weights by gradient descent:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial E}{\partial \mathbf{w}}$$

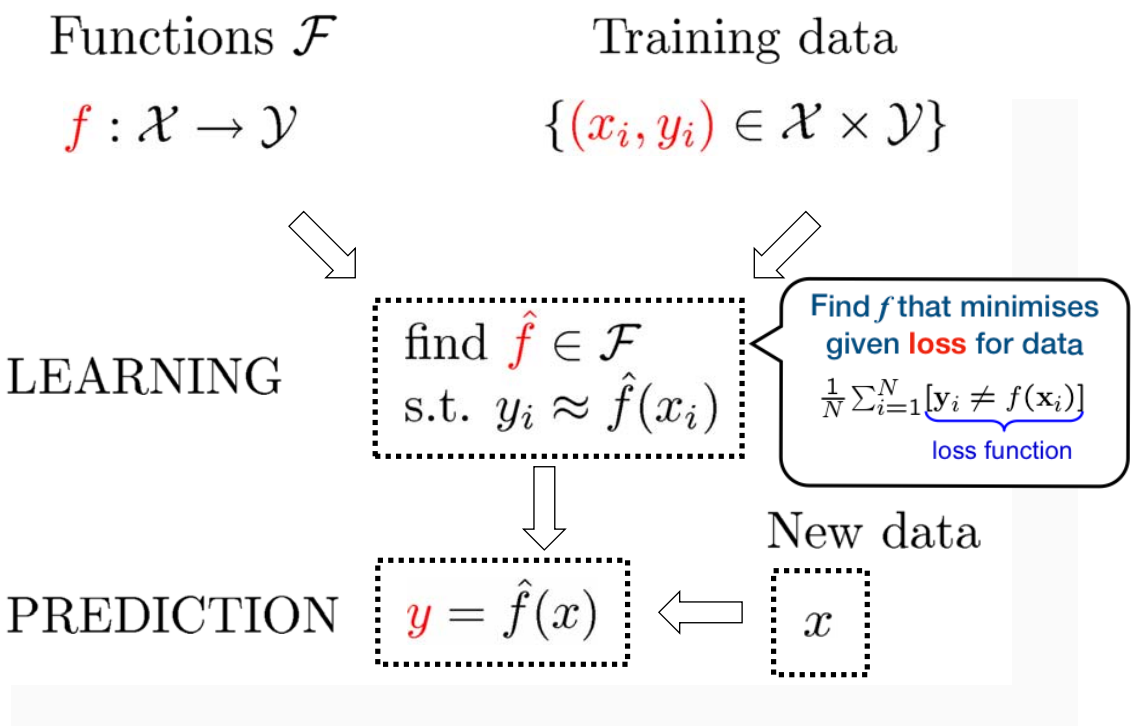

Training of multi-layer networks

- Find network weights to minimize the error between true and estimated labels of training examples:

$$E(\mathbf{w}) = \sum_{j=1}^{N} \left(y_j - f_{\mathbf{w}}(\mathbf{x}_j) \right)^2$$

- Update weights by gradient descent:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial E}{\partial \mathbf{w}}$$

- This requires perceptrons with a differentiable nonlinearity

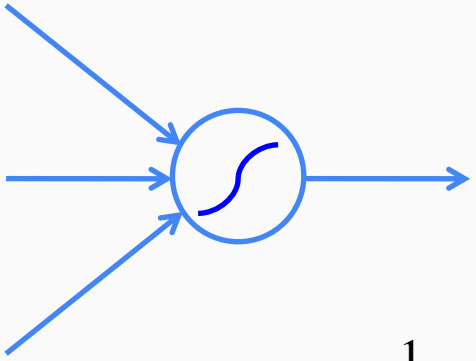

Sigmoid: $g(t) = \dfrac{1}{1+e^{-t}}$

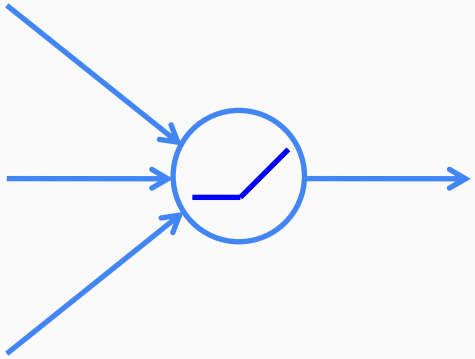

Rectified linear unit (ReLU):
$g(t) = \max(0,t)$

Training of multi-layer networks

- Find network weights to minimize the error between true and estimated labels of training examples:

$$E(\mathbf{w}) = \sum_{j=1}^{N} \left(y_j - f_{\mathbf{w}}(\mathbf{x}_j) \right)^2$$

- Update weights by gradient descent: $\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial E}{\partial \mathbf{w}}$
- Back-propagation:
  gradients are computed from output to input layers and combined using chain rule
- Stochastic gradient descent:
  compute the weight update w.r.t. one training example (or a small batch of examples) at a time, cycle through training examples in random order in multiple epochs

Multi-Layer Network Demo

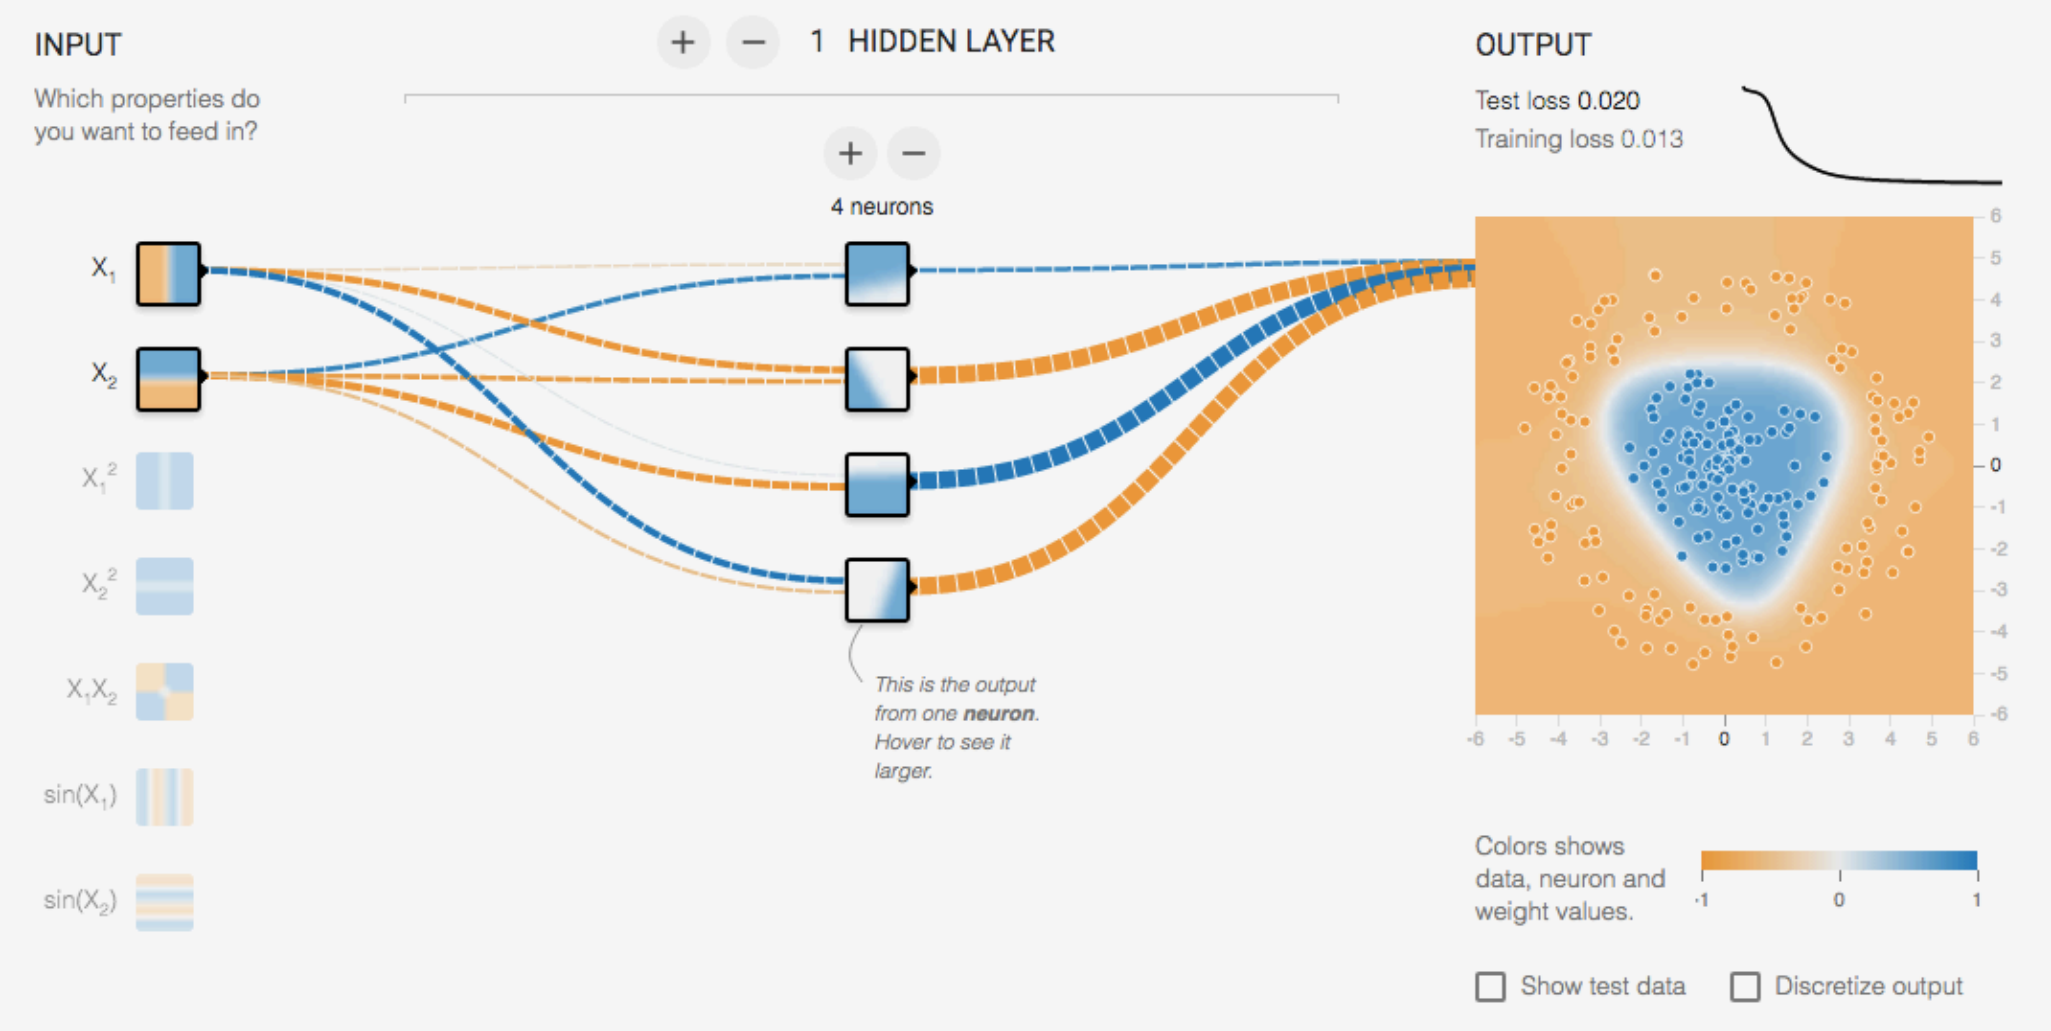

http://playground.tensorflow.org/

Neural networks: Pros and cons

- Pros
  - Flexible and general function approximation framework
  - Can build extremely powerful models by adding more layers
- Cons
  - Hard to analyze theoretically (e.g., training is prone to local optima)
  - Huge amount of training data and computing power may be required to get good performance
  - The space of implementation choices is huge (network architectures, parameters)

Convolutional neural network

Neural networks for images

Neural network

weights

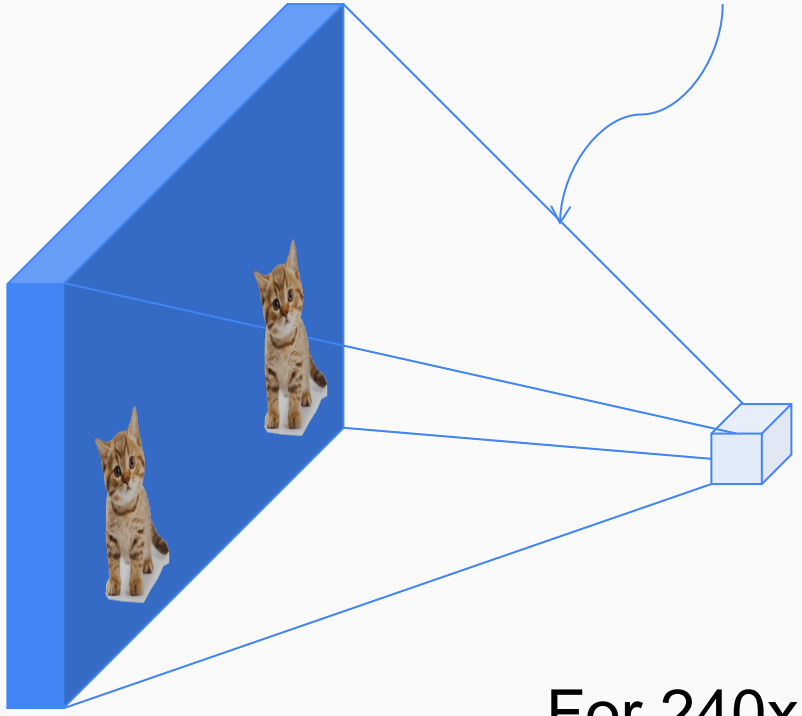

Output: weighted sum
and non-linearity

image

For 240x320 image -> 76 800 weights!
Requires plenty of training samples to learn translation invariance.

Neural networks for images

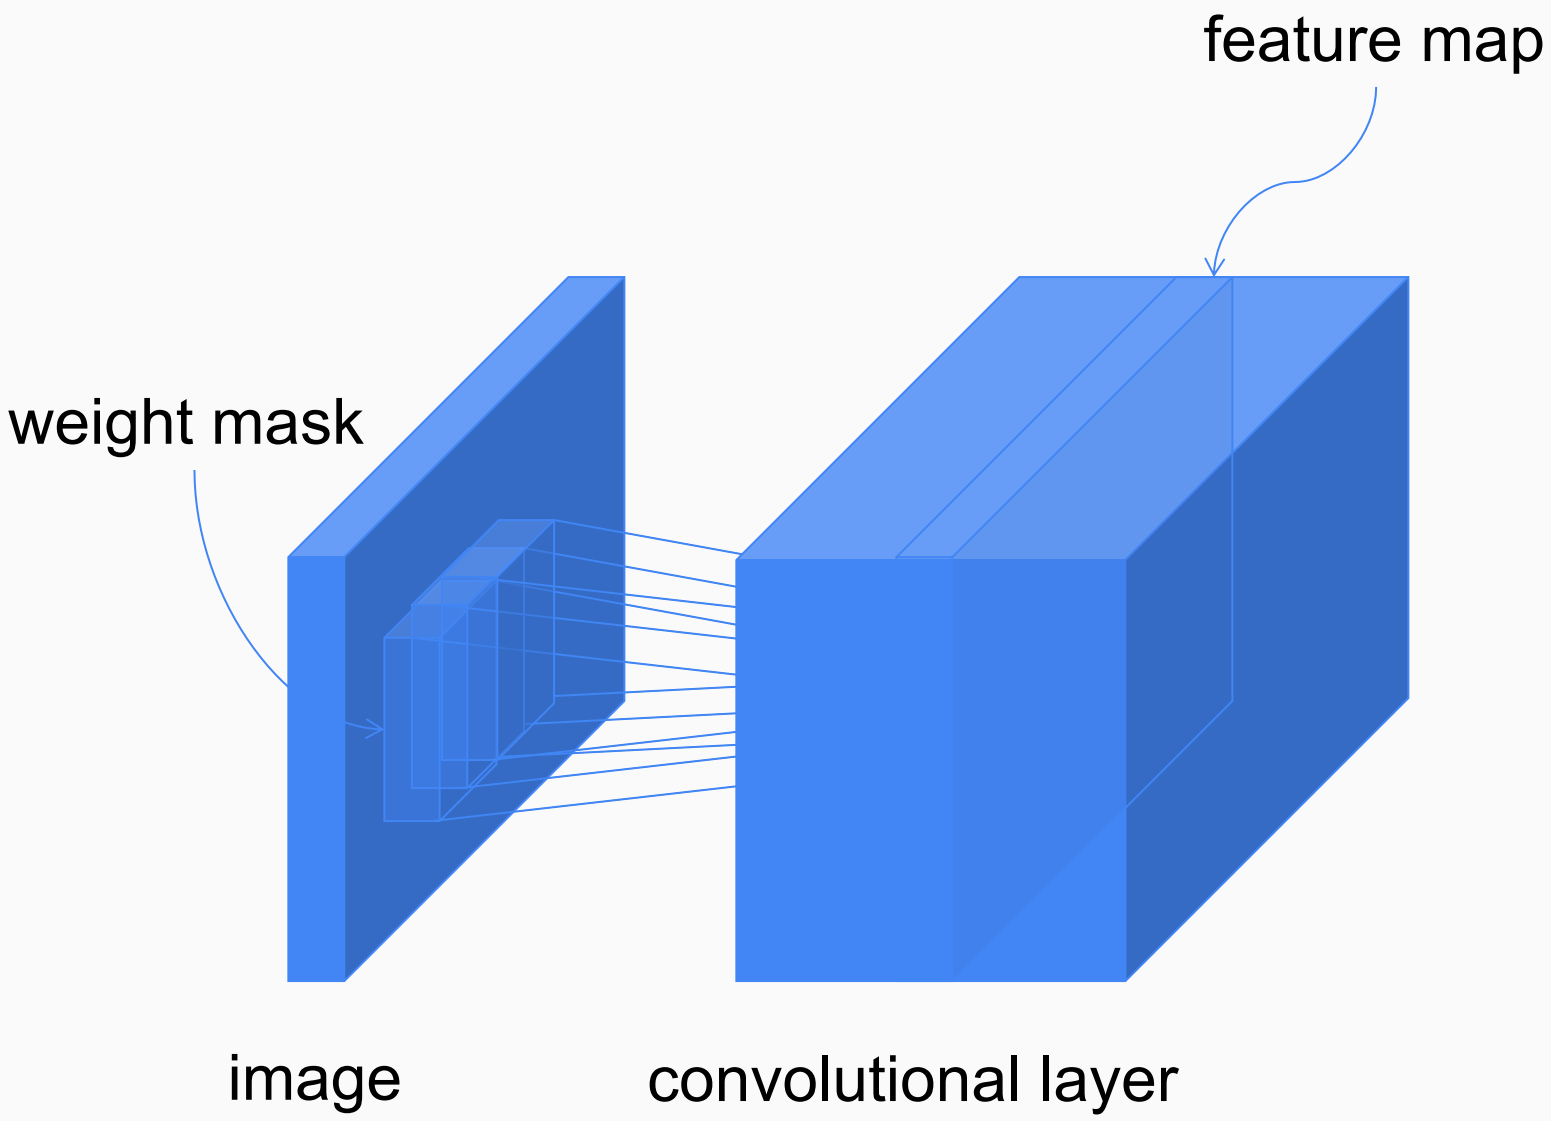

Neural networks for images

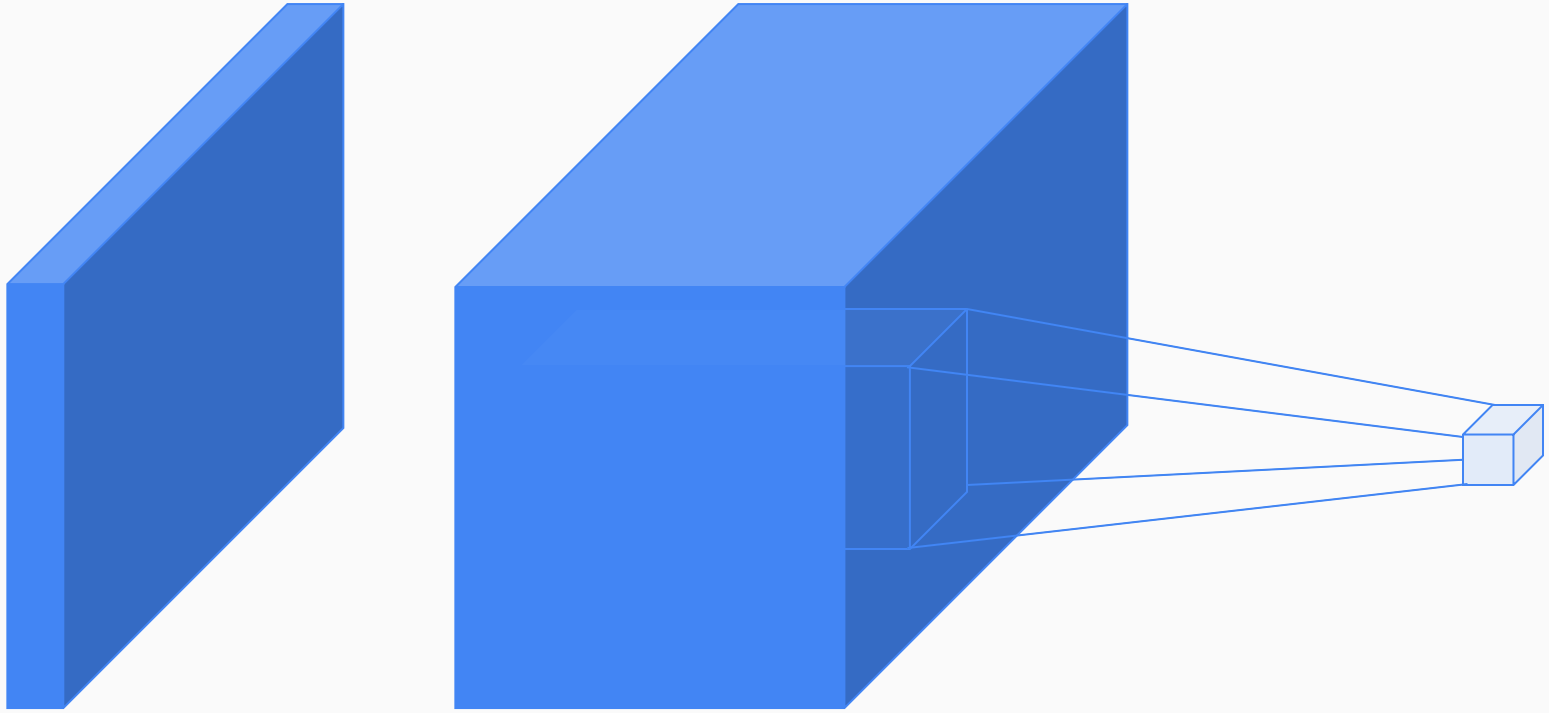

image convolutional layer

Convolution as feature extraction

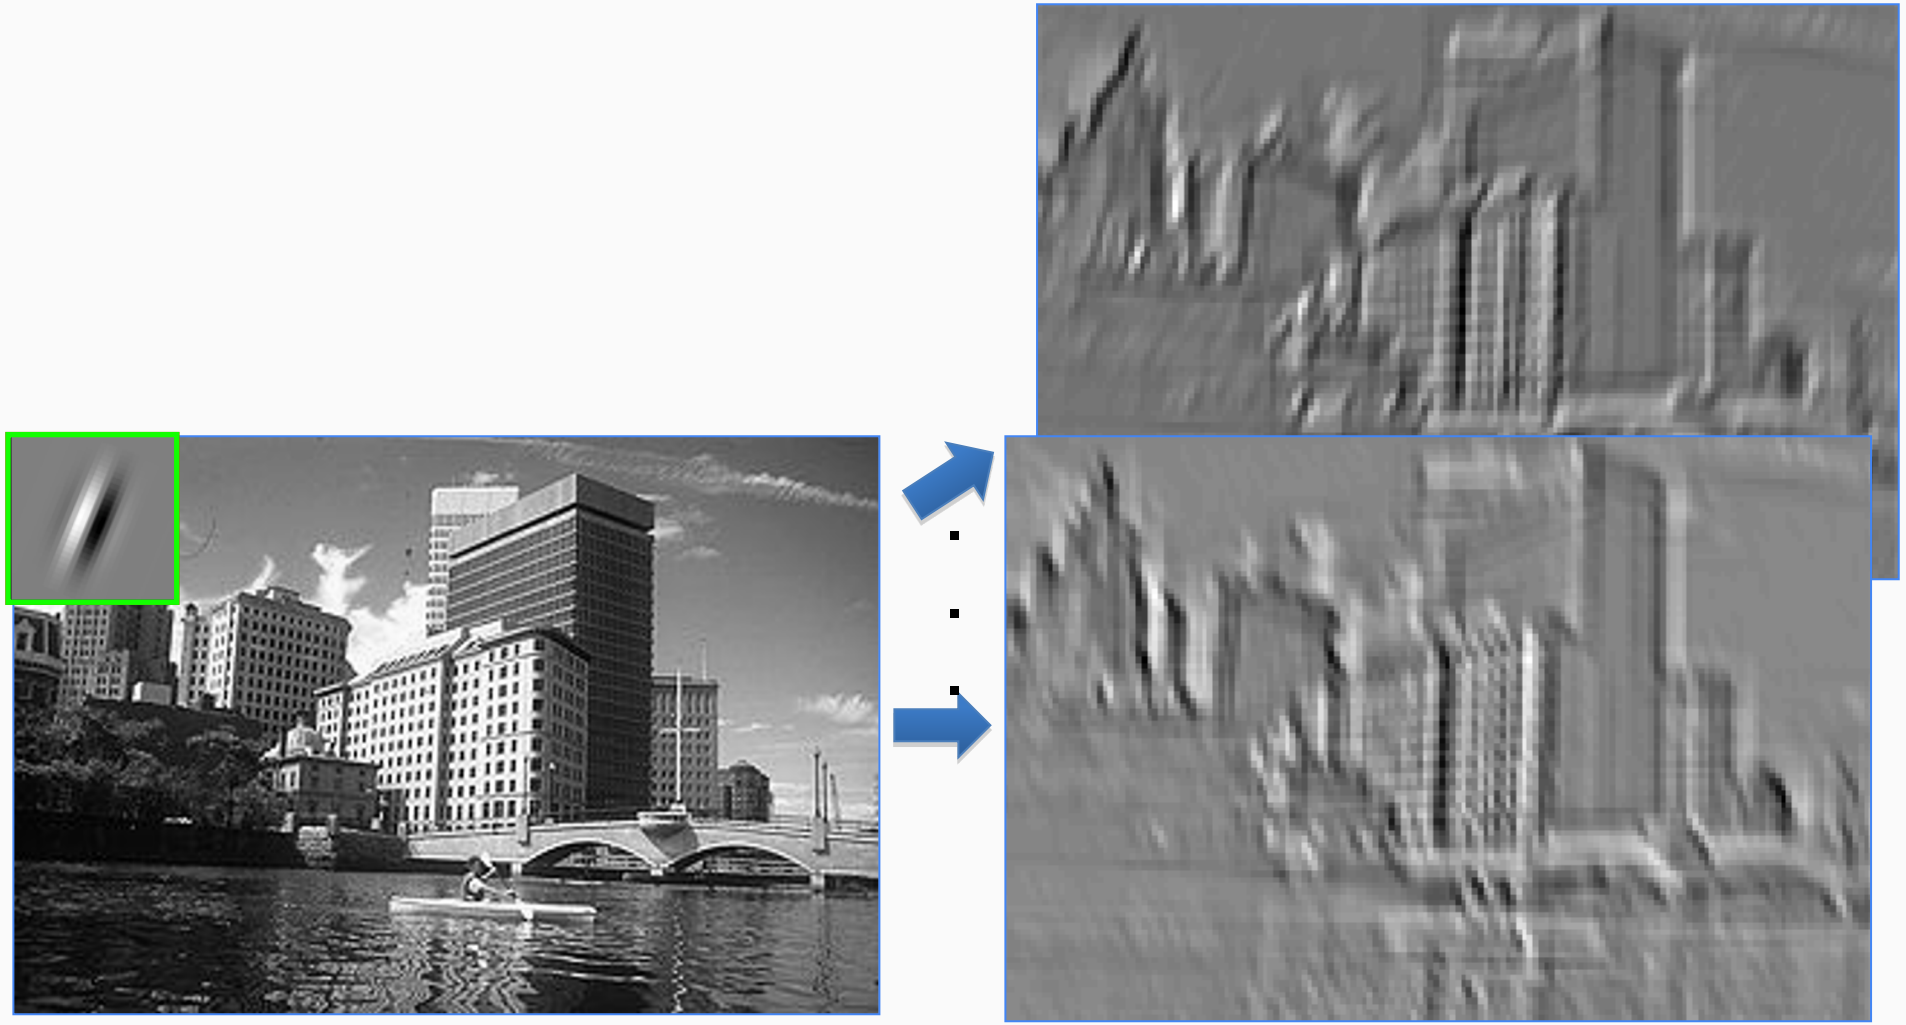

Convolutional Neural Networks

- Neural network with specialized connectivity structure
- Stack multiple stages of feature extractors
- Higher stages compute more global, more invariant features
- Classification layer at the end

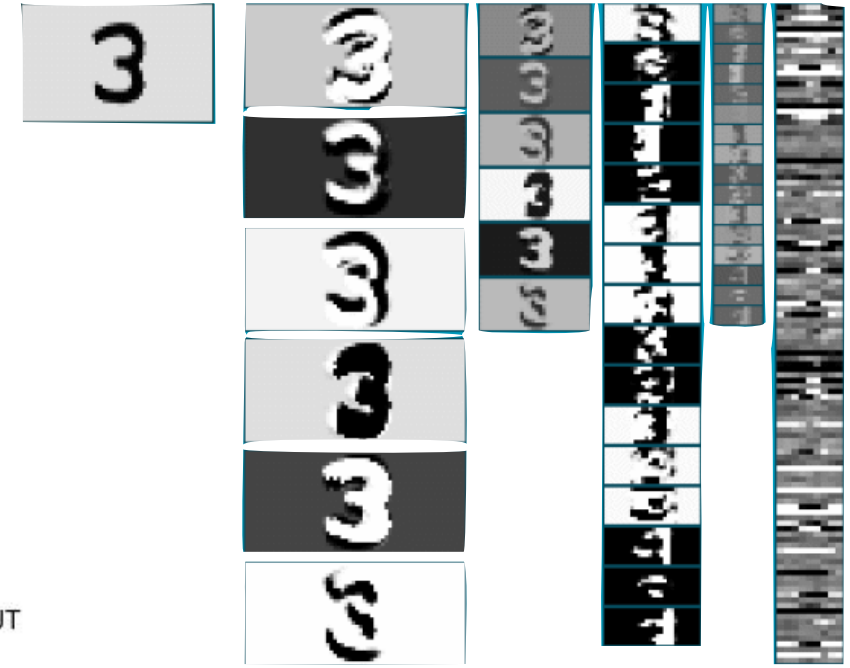

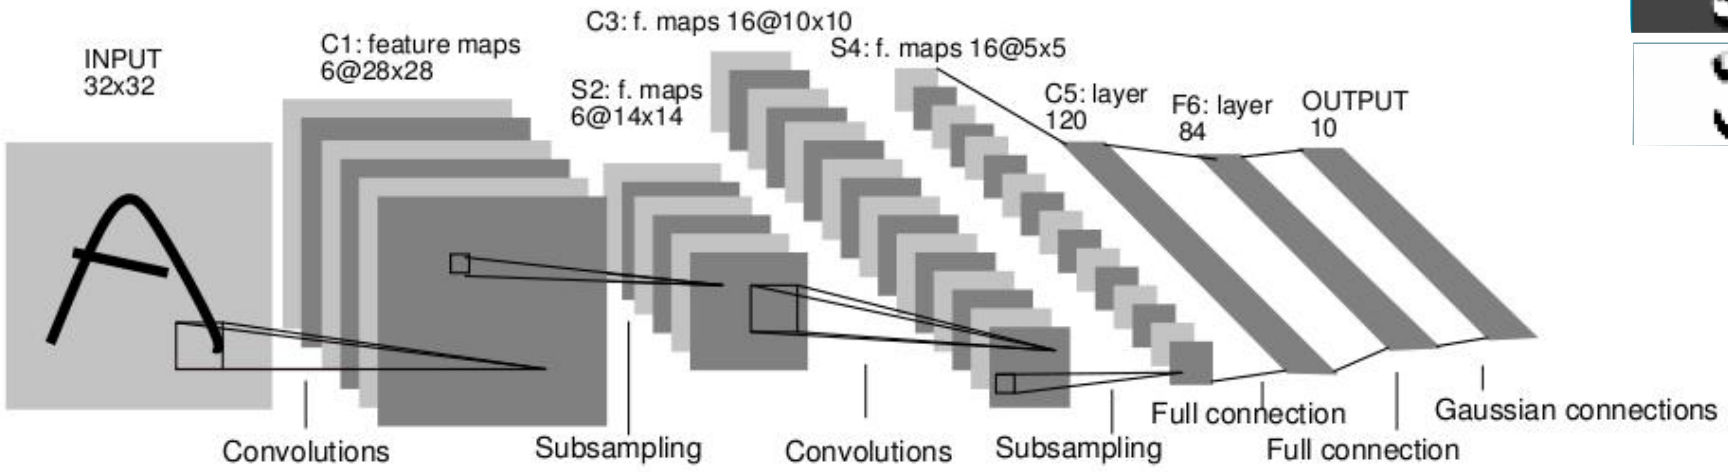

Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, [Gradient-based learning applied to document recognition](), Proc. IEEE 86(11): 2278–2324, 1998.

Biological inspiration

- D. Hubel and T. Wiesel (1959, 1962, Nobel Prize 1981)
  - Visual cortex consists of a hierarchy of simple, complex, and hyper-complex cells

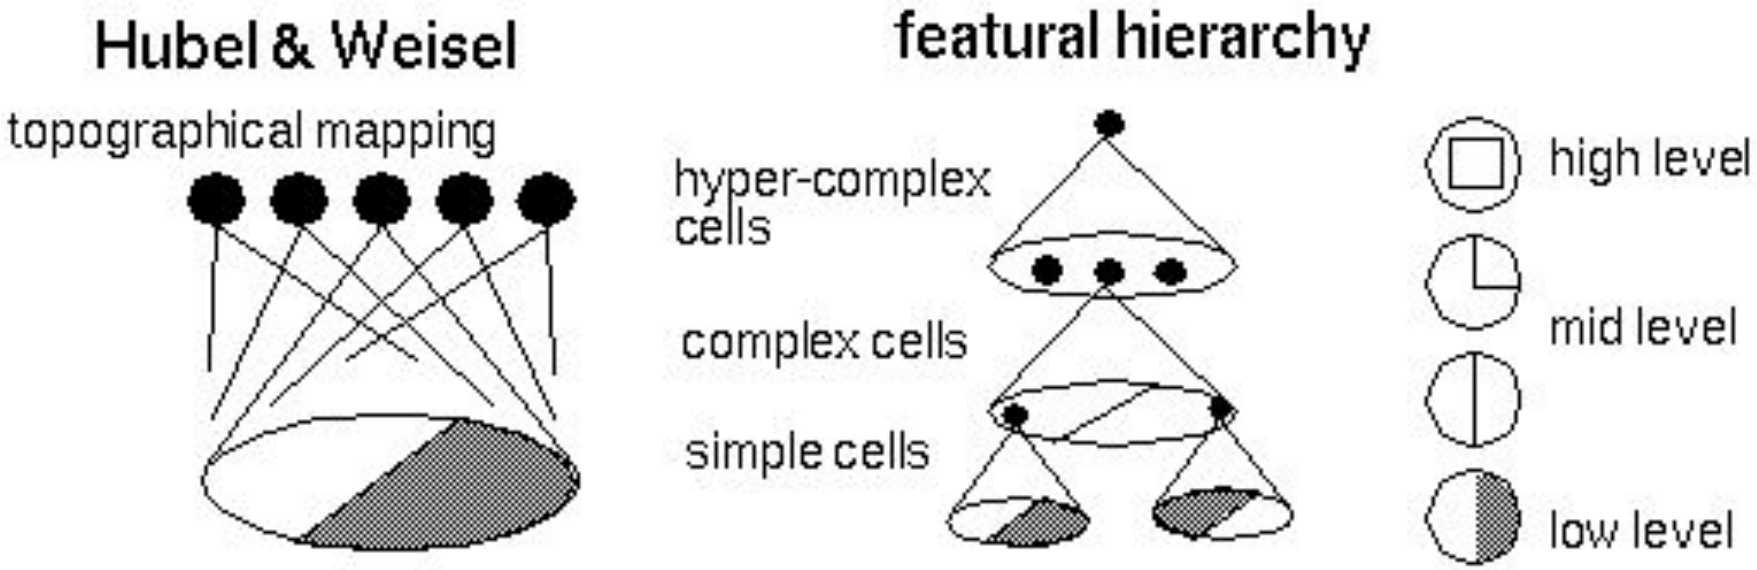

Source

Convolutional Neural Networks

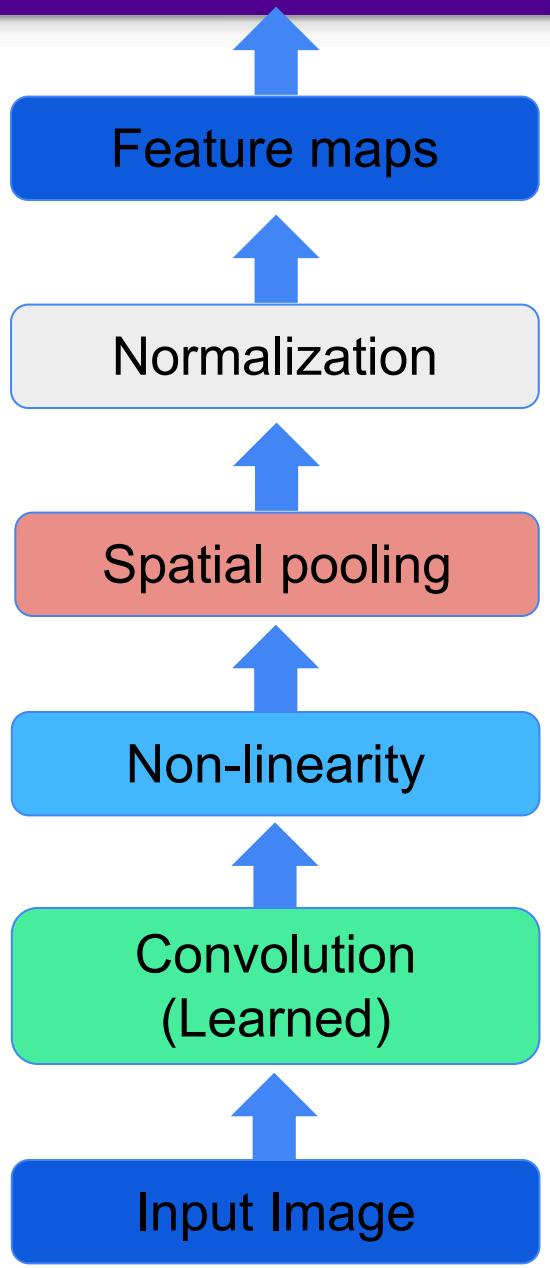

Convolutional Neural Networks

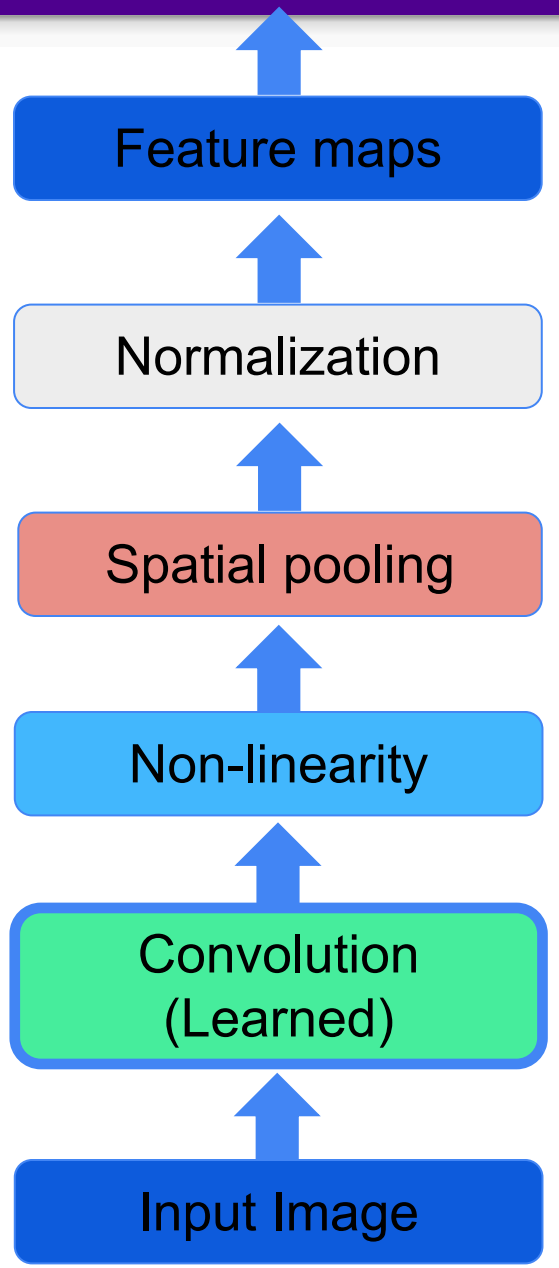

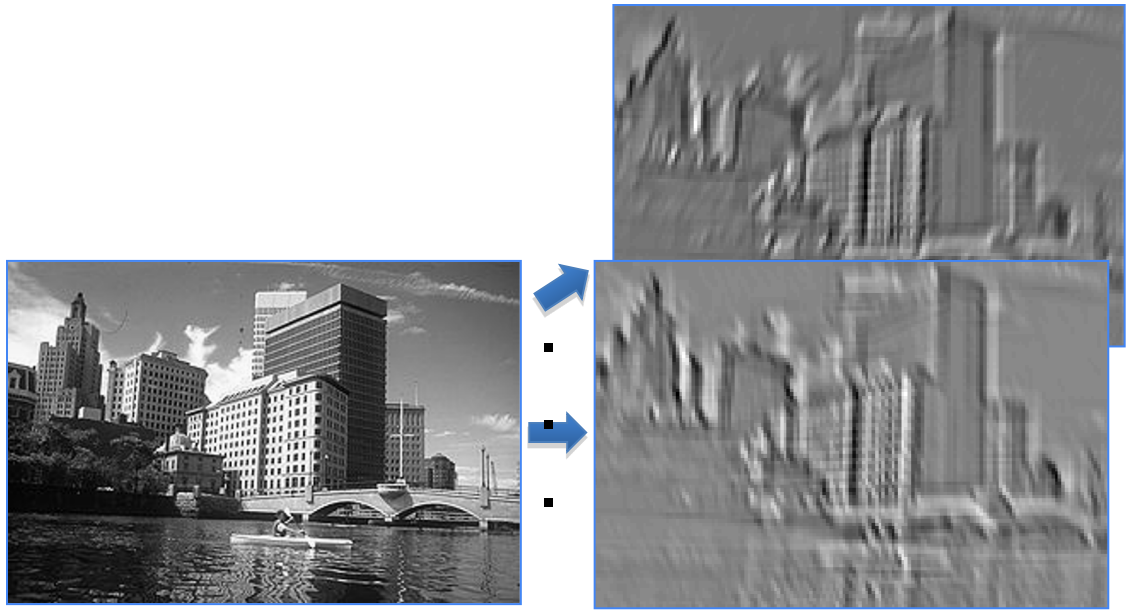

Convolutional Neural Networks

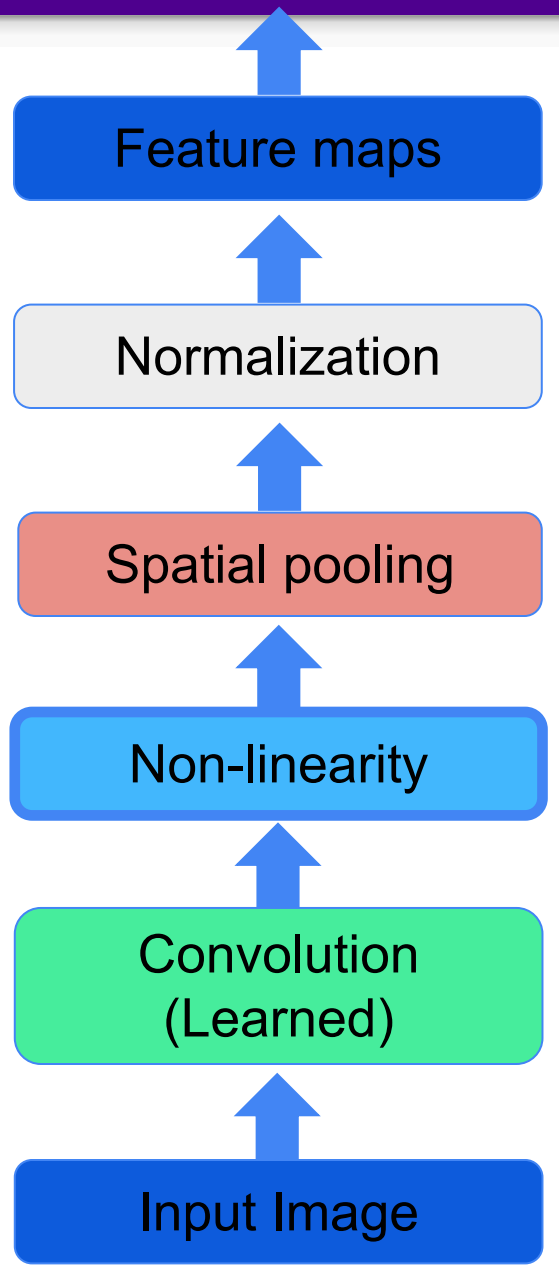

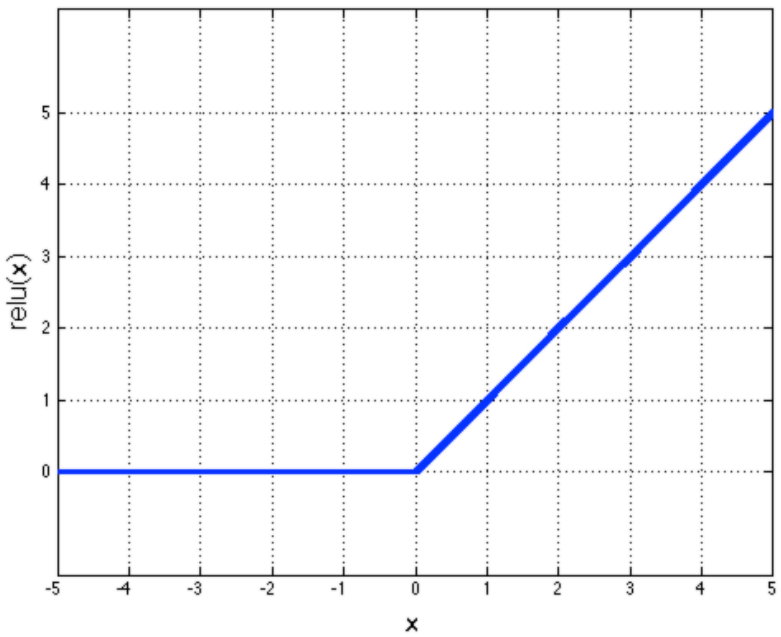

Convolutional Neural Networks

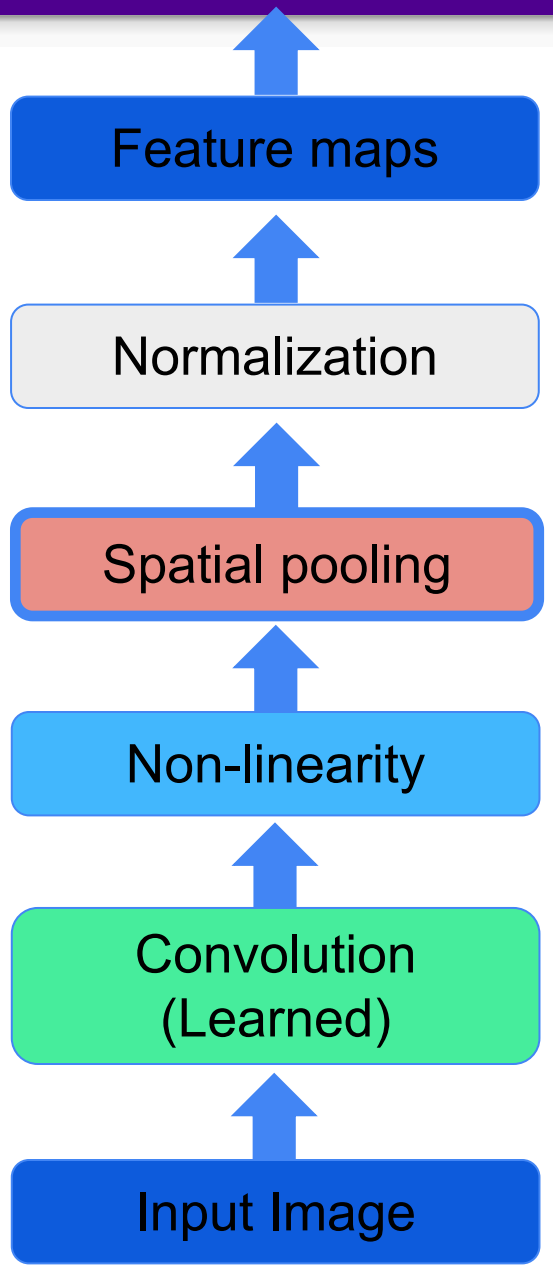

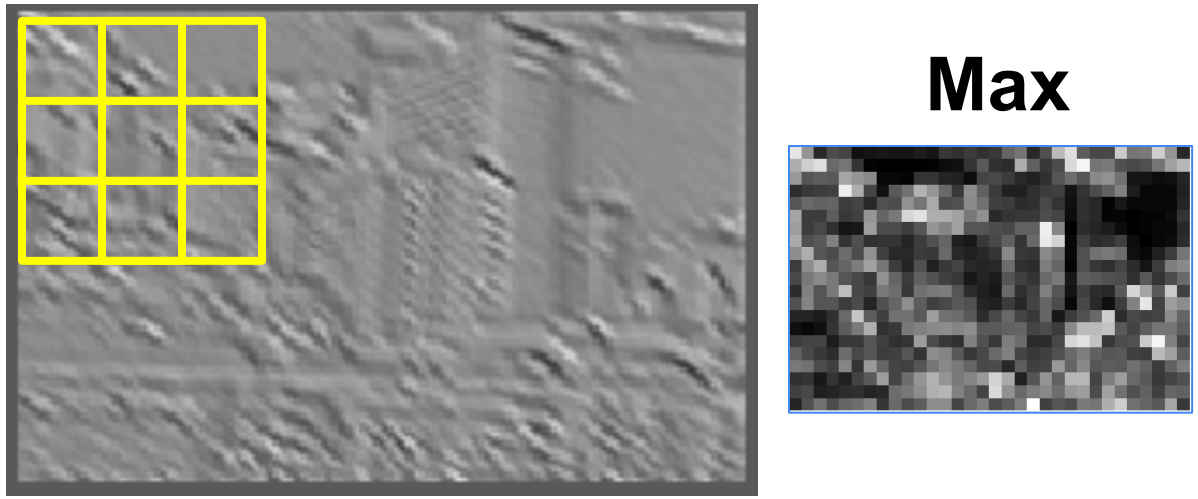

Convolutional Neural Networks

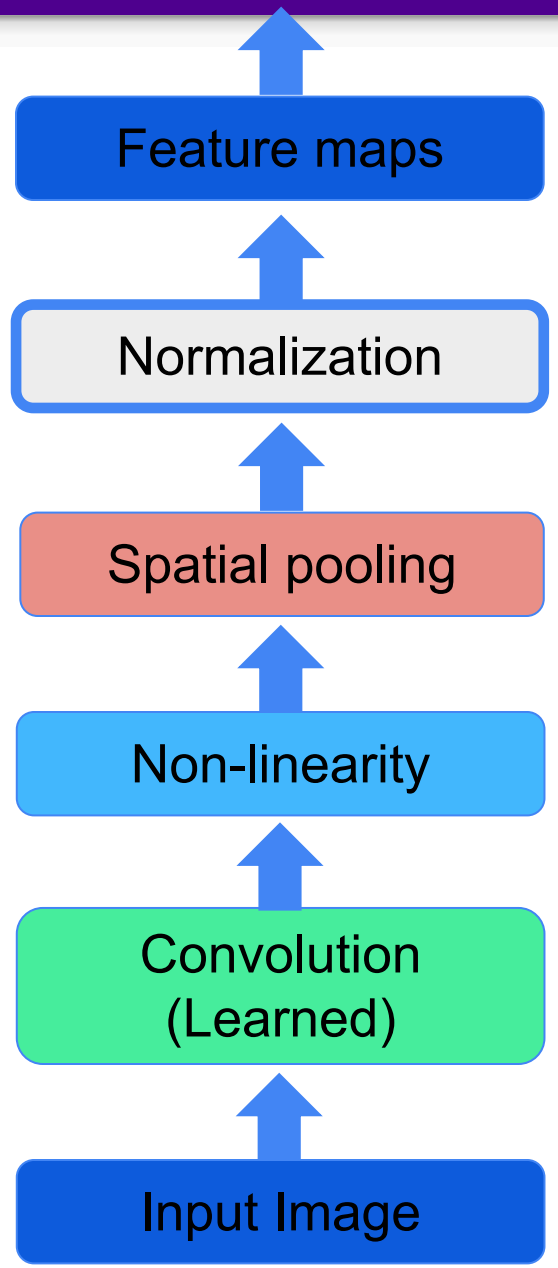

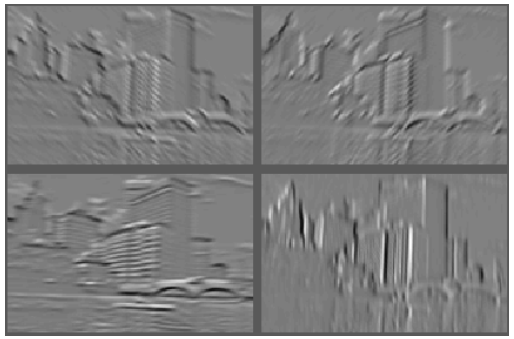

Feature Maps

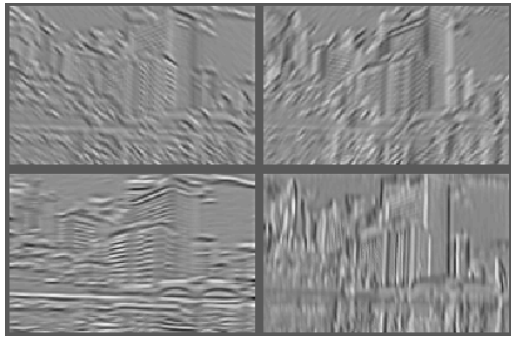

Feature Maps After Contrast Normalization

Convolutional Neural Networks

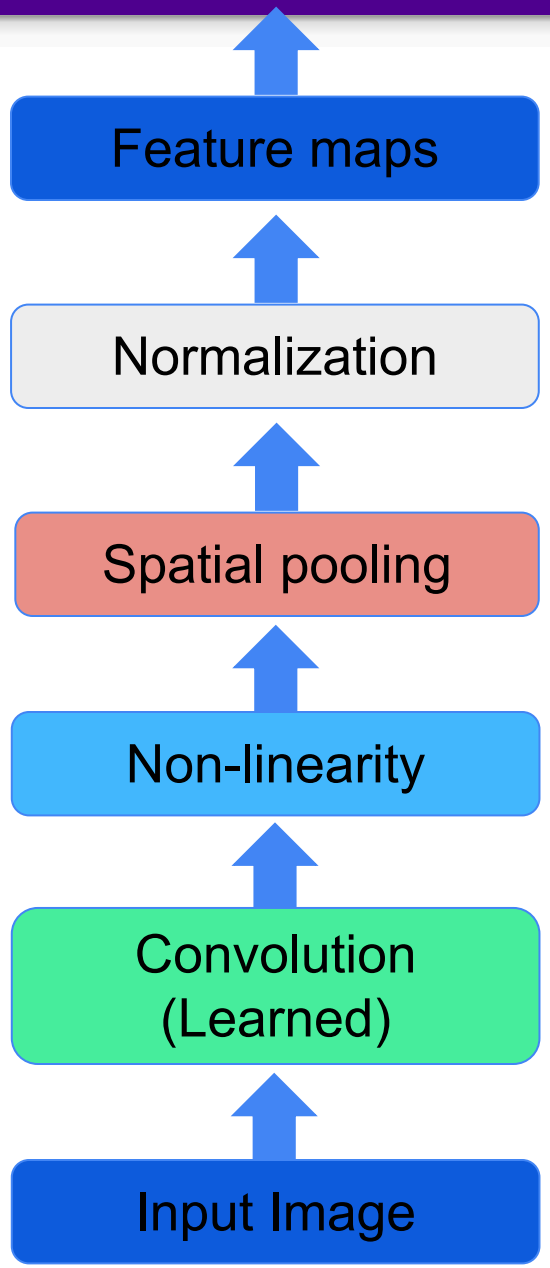

Convolutional filters are trained in a supervised manner by back-propagating classification error

Simplified architecture

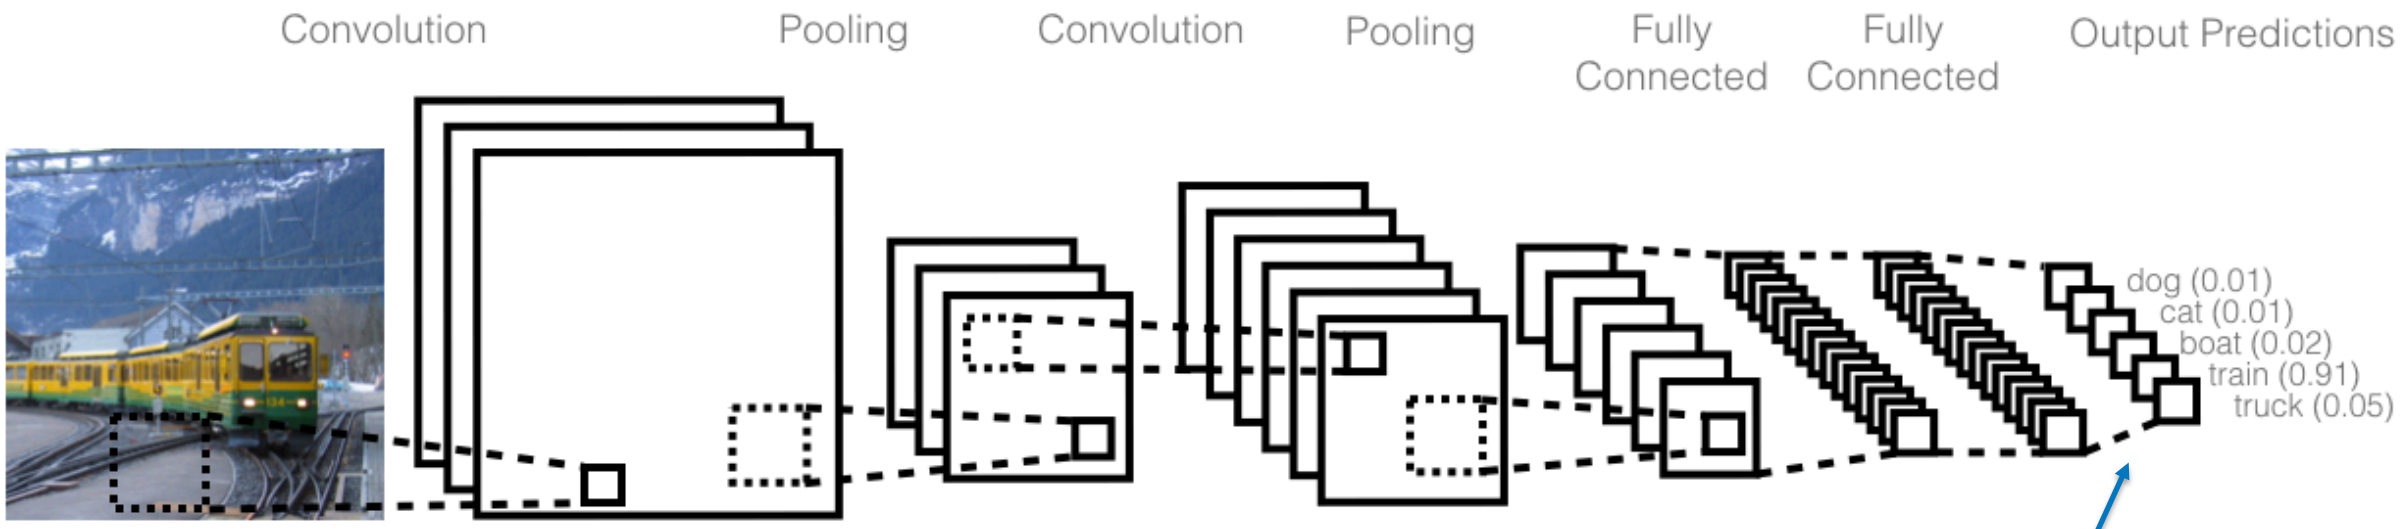

Softmax layer:

$$P(c \mid \mathbf{x}) = \frac{\exp(\mathbf{w}_c \cdot \mathbf{x})}{\sum_{k=1}^{C} \exp(\mathbf{w}_k \cdot \mathbf{x})}$$

AlexNet

- Similar framwork to LeCun'98, but:
  - Larger model (7 hidden layers, 650 000 units, 60 000 000 parameters)
  - More data (10^6 vs. 10^3 images)
  - GPU implementation (50x speedup over CPU, trained on two GPUs for a week)

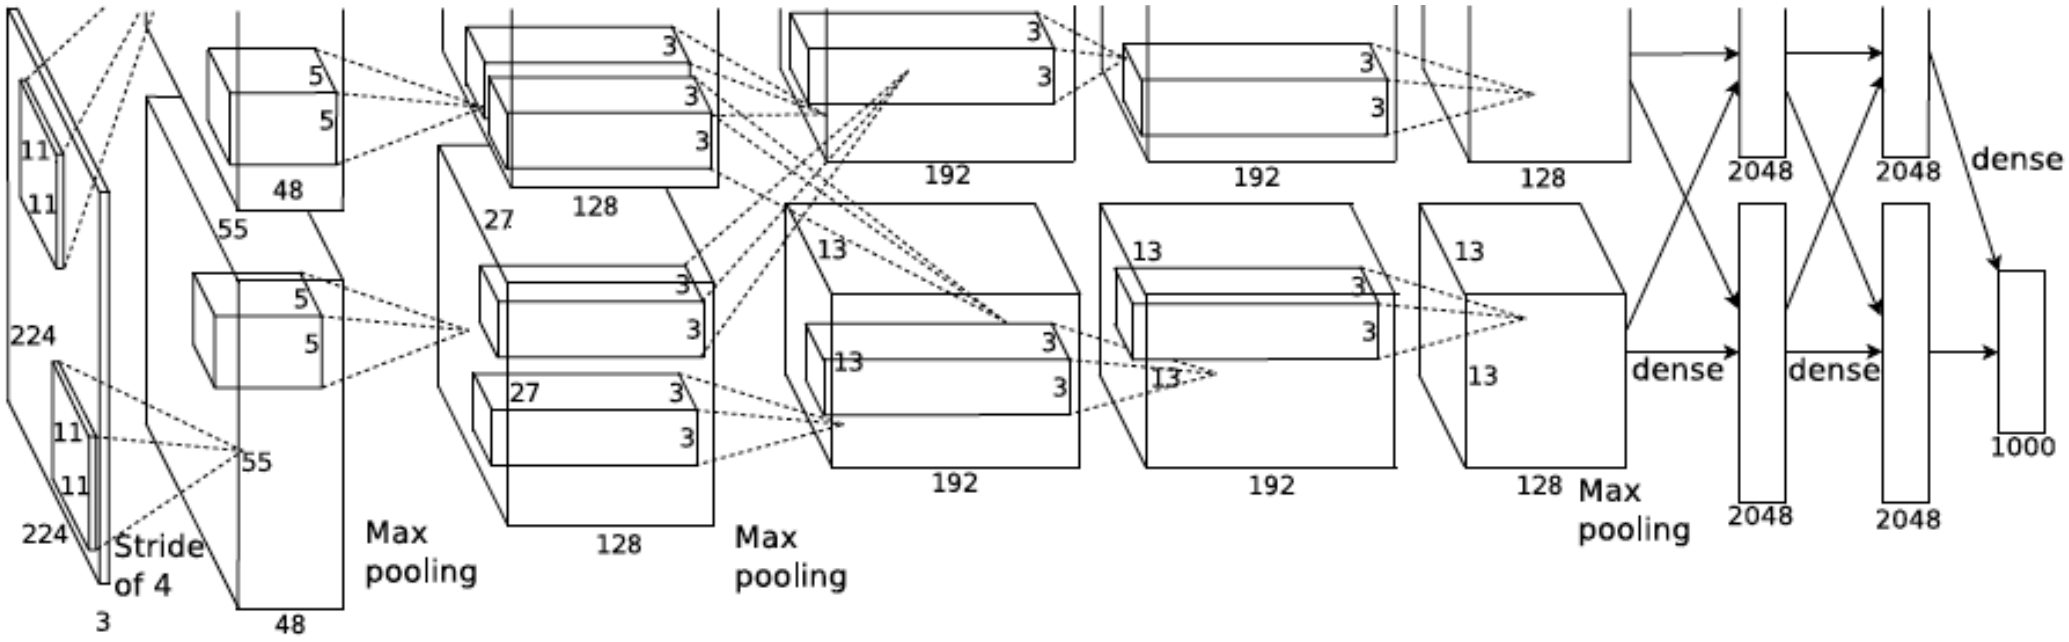

A. Krizhevsky, I. Sutskever, and G. Hinton, [ImageNet Classification with Deep Convolutional Neural Networks](), NIPS 2012

Refinement of AlexNet architecture

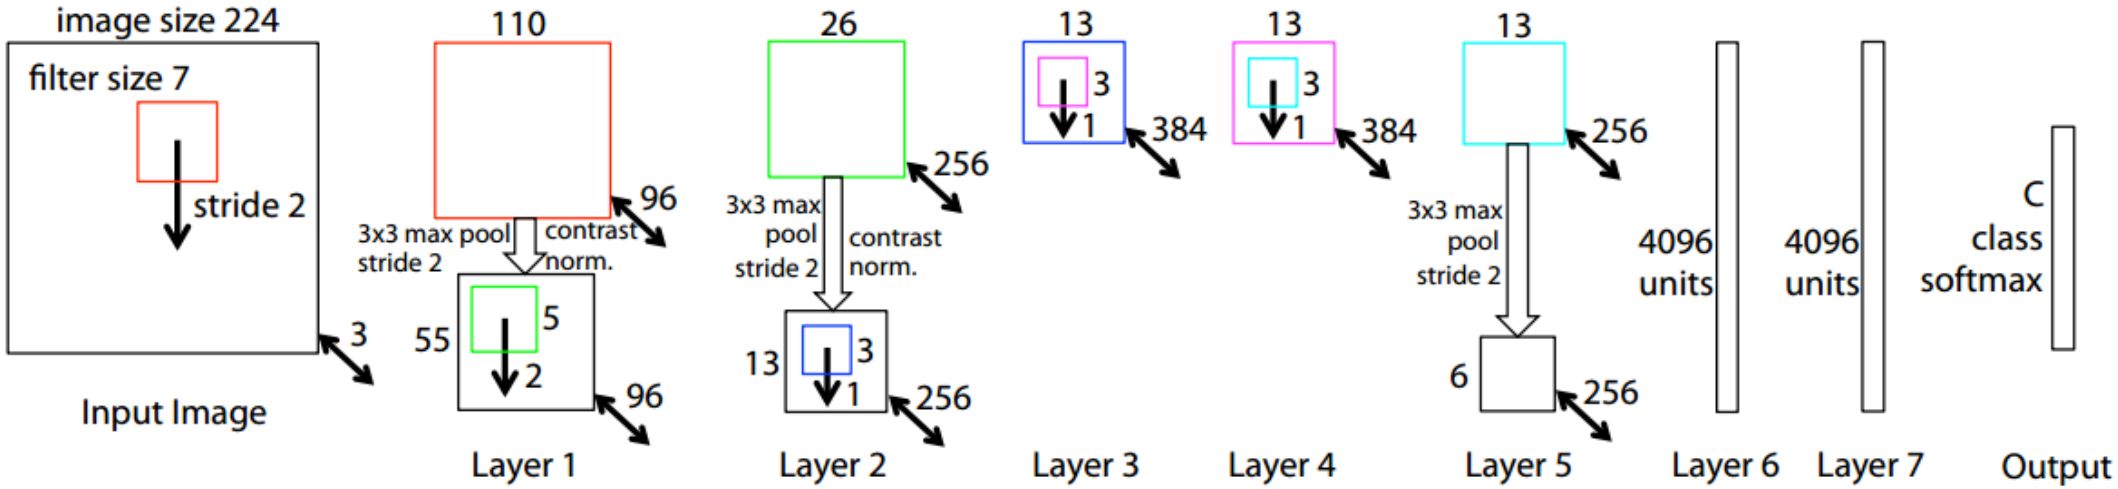

Figure 3. Architecture of our 8 layer convnet model. A 224 by 224 crop of an image (with 3 color planes) is presented as the input. This is convolved with 96 different 1st layer filters (red), each of size 7 by 7, using a stride of 2 in both x and y. The resulting feature maps are then: (i) passed through a rectified linear function (not shown), (ii) pooled (max within 3x3 regions, using stride 2) and (iii) contrast normalized across feature maps to give 96 different 55 by 55 element feature maps. Similar operations are repeated in layers 2,3,4,5. The last two layers are fully connected, taking features from the top convolutional layer as input in vector form ($6 \cdot 6 \cdot 256 = 9216$ dimensions). The final layer is a C-way softmax function, C being the number of classes. All filters and feature maps are square in shape.

M. Zeiler and R. Fergus, Visualizing and Understanding Convolutional Networks, arXiv preprint, 2013

Layer 1 Filters

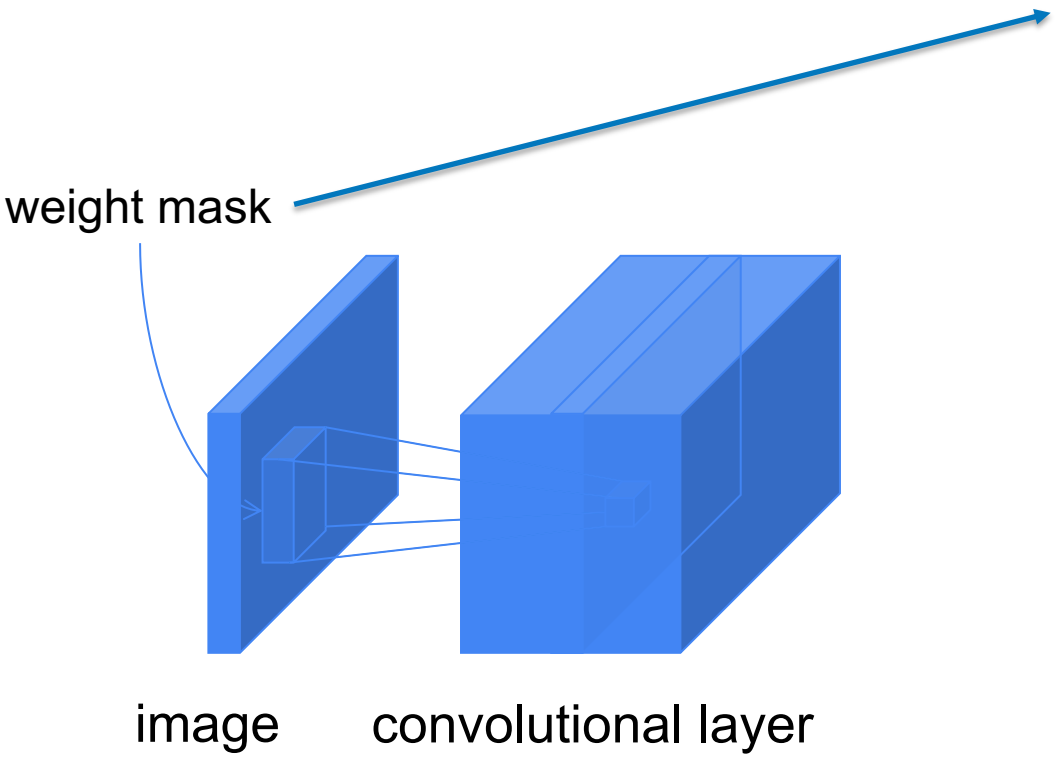

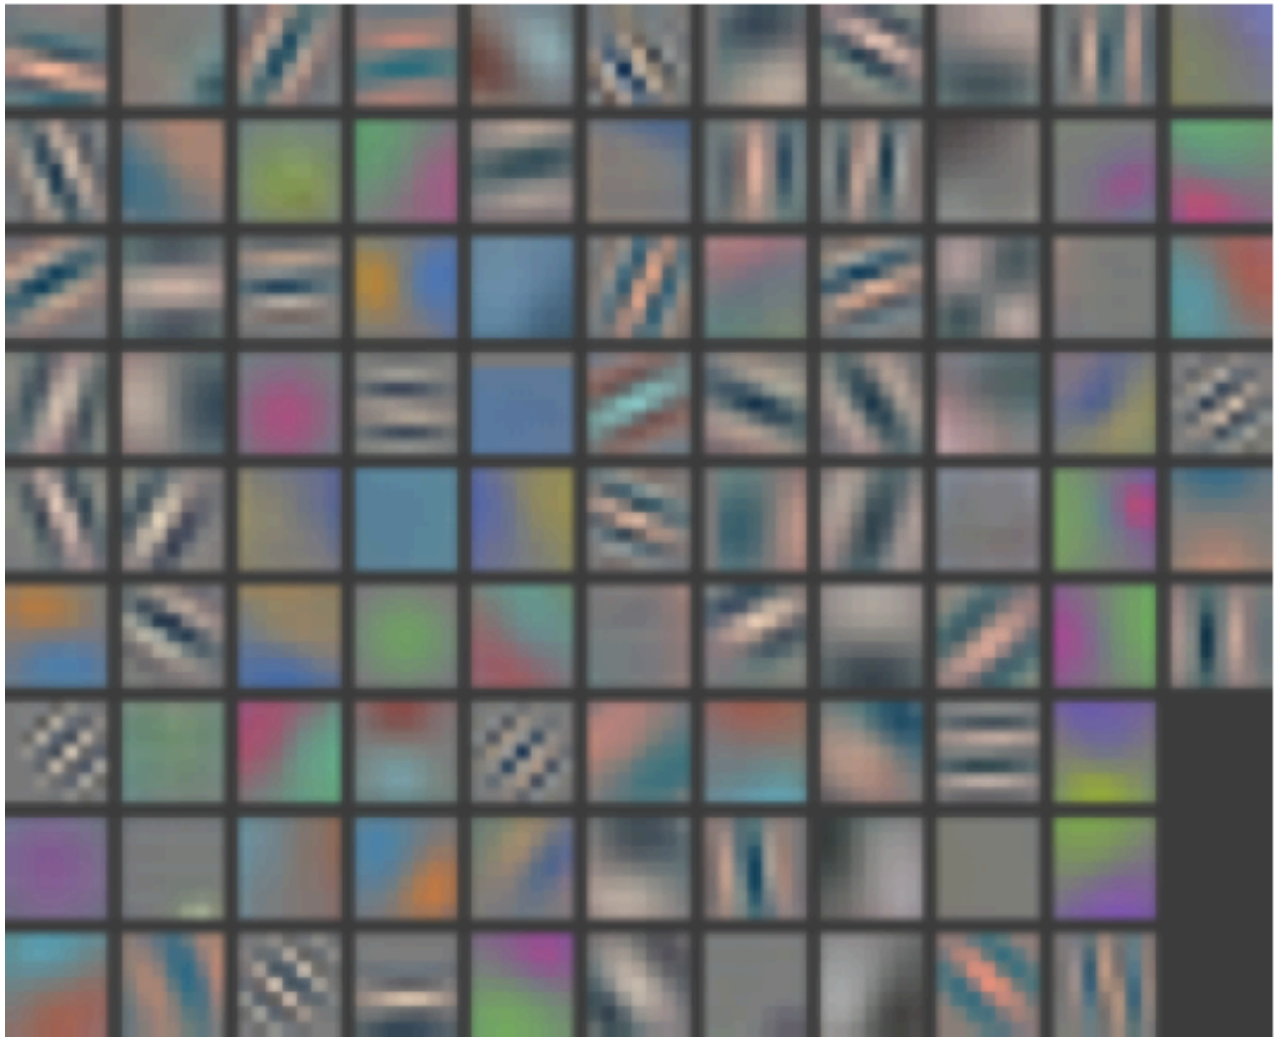

M. Zeiler and R. Fergus, [Visualizing and Understanding Convolutional Networks](), ECCV 2014 (Best Paper Award winner)

Layer 1: Top-9 Patches (patches that give maximal activation)

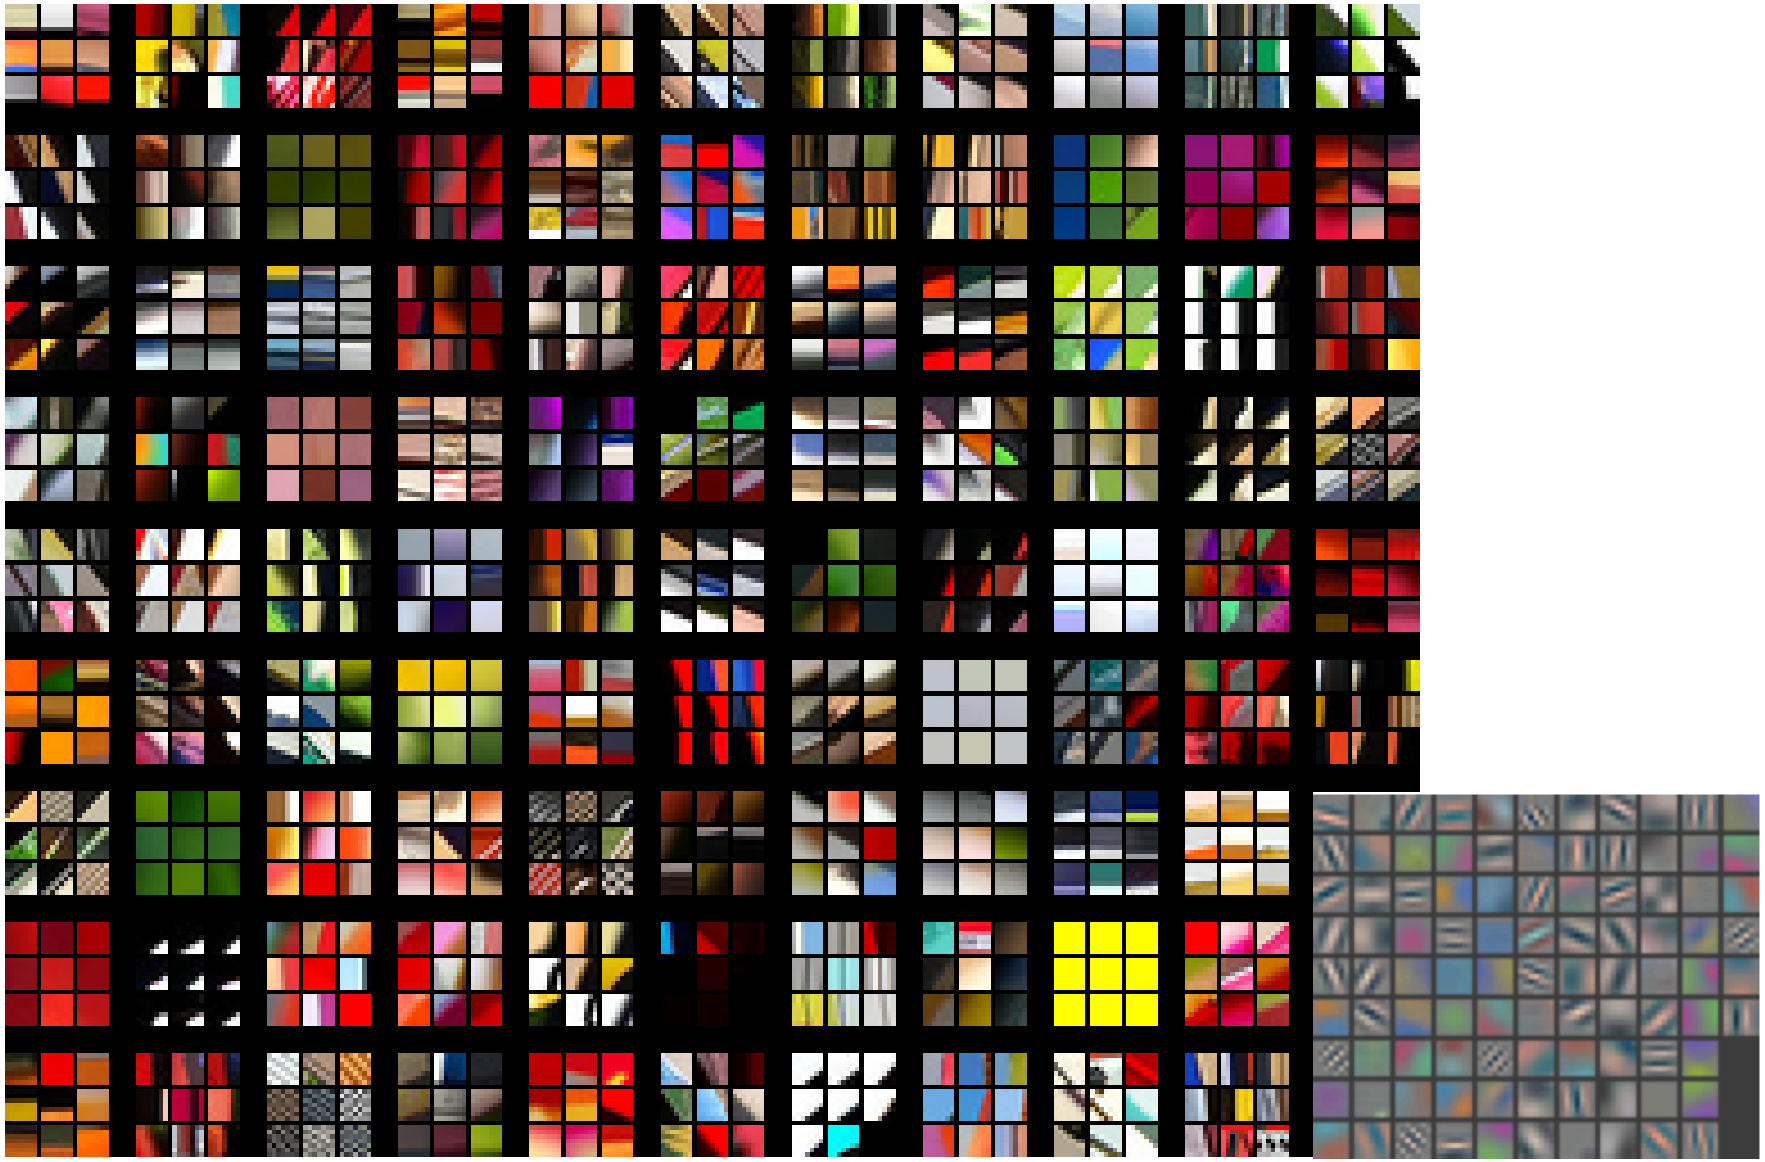

Layer 2: Filters

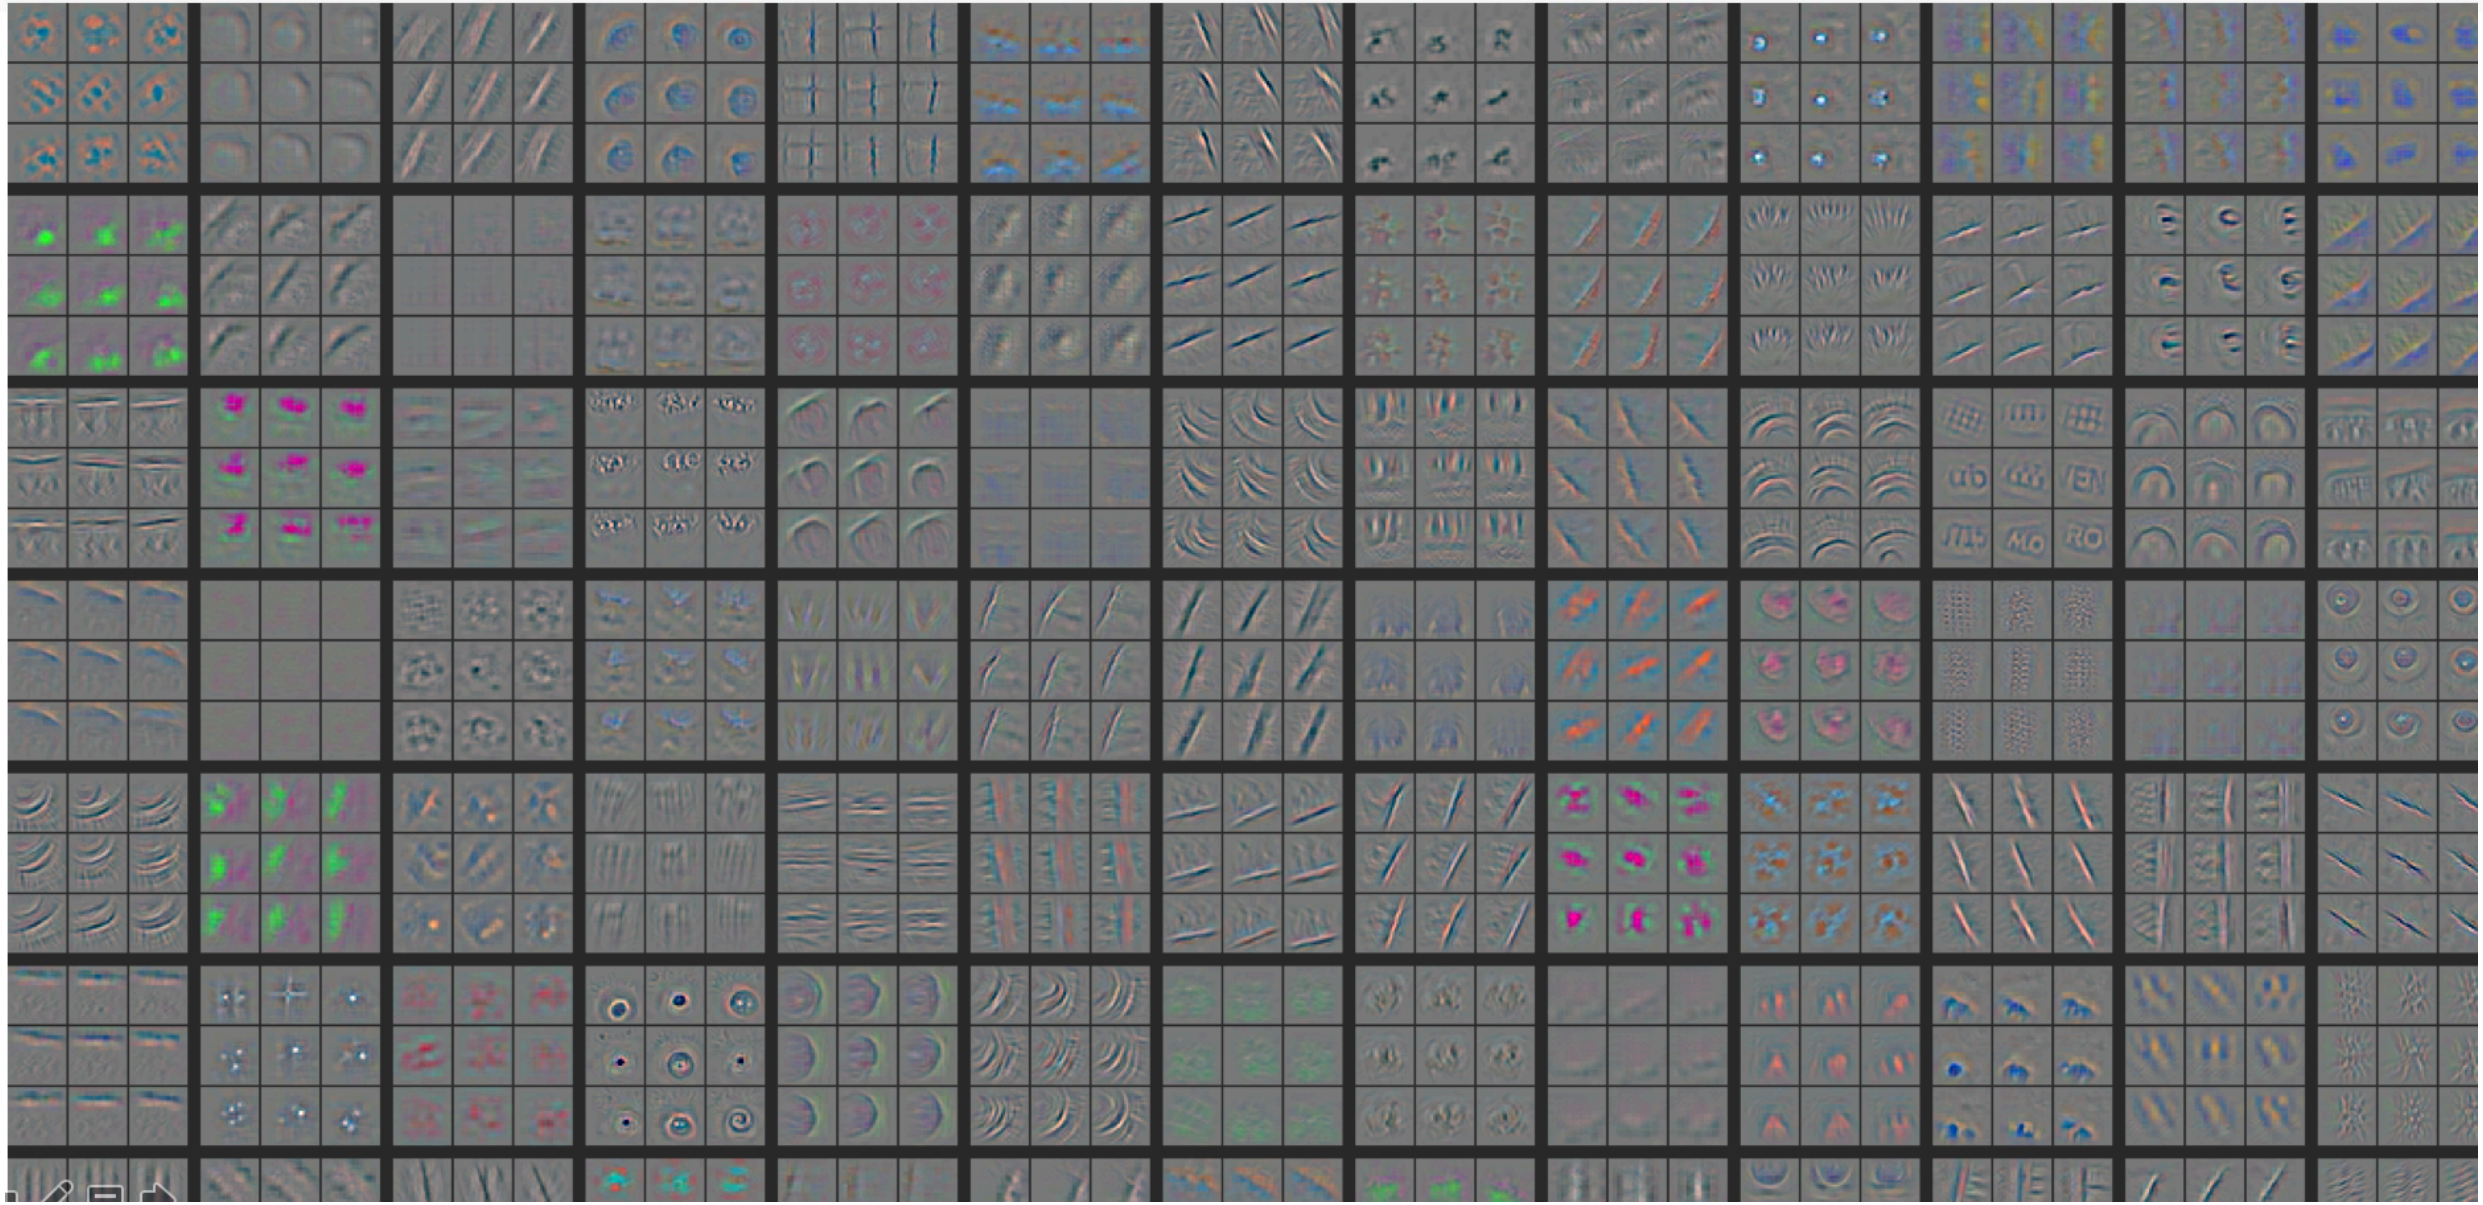

Layer 2: Top-9 Patches

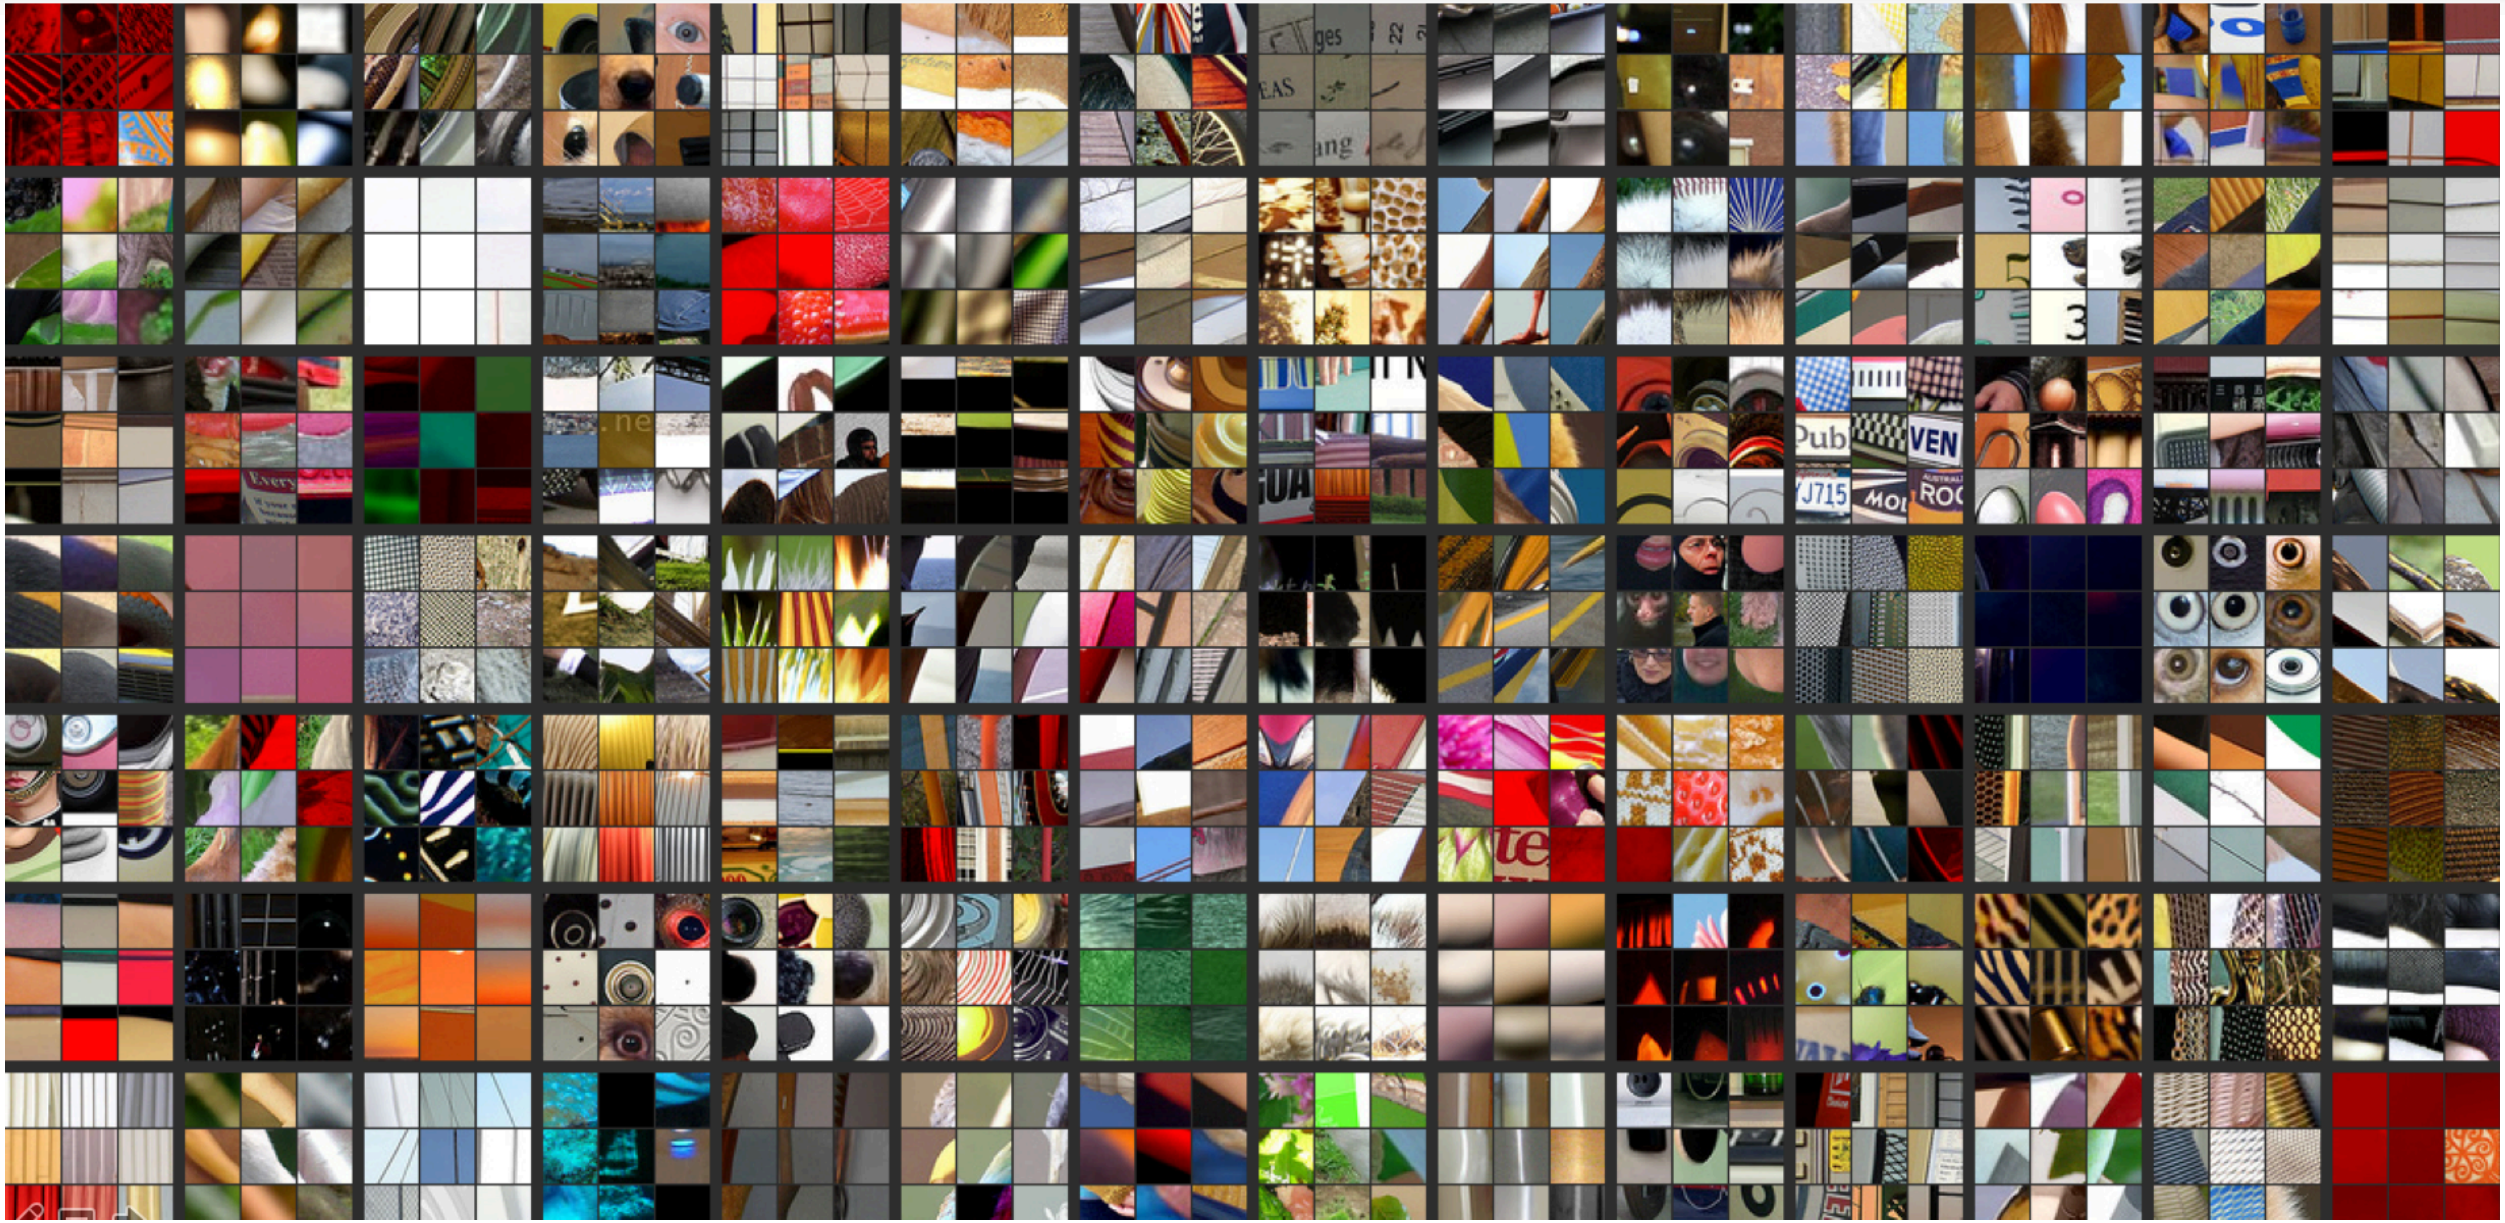

Layer 5: Filters

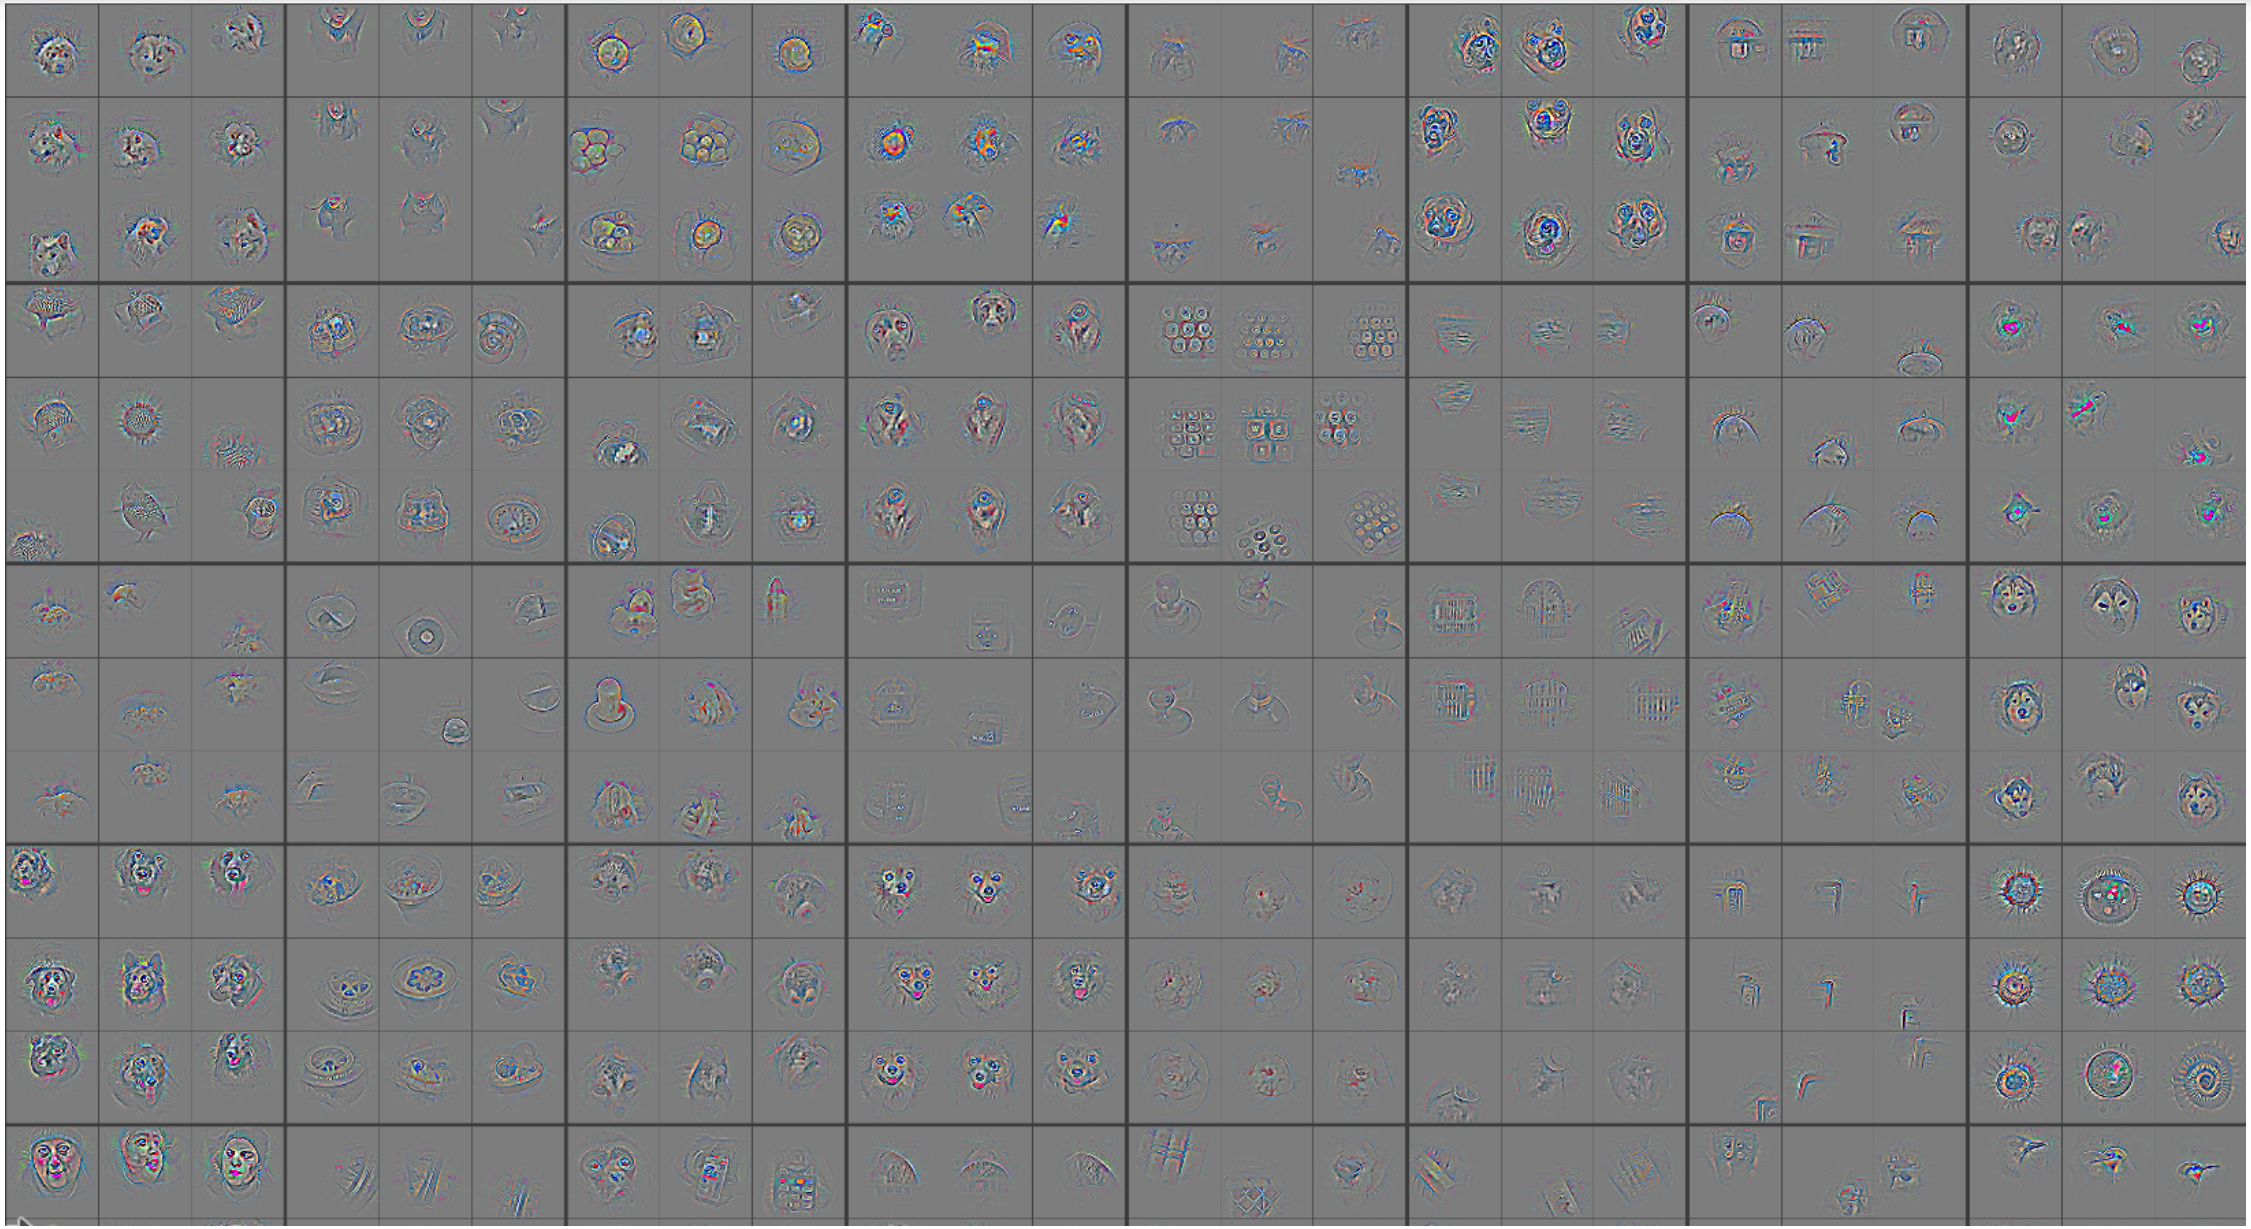

Layer 5: Top-9 Patches

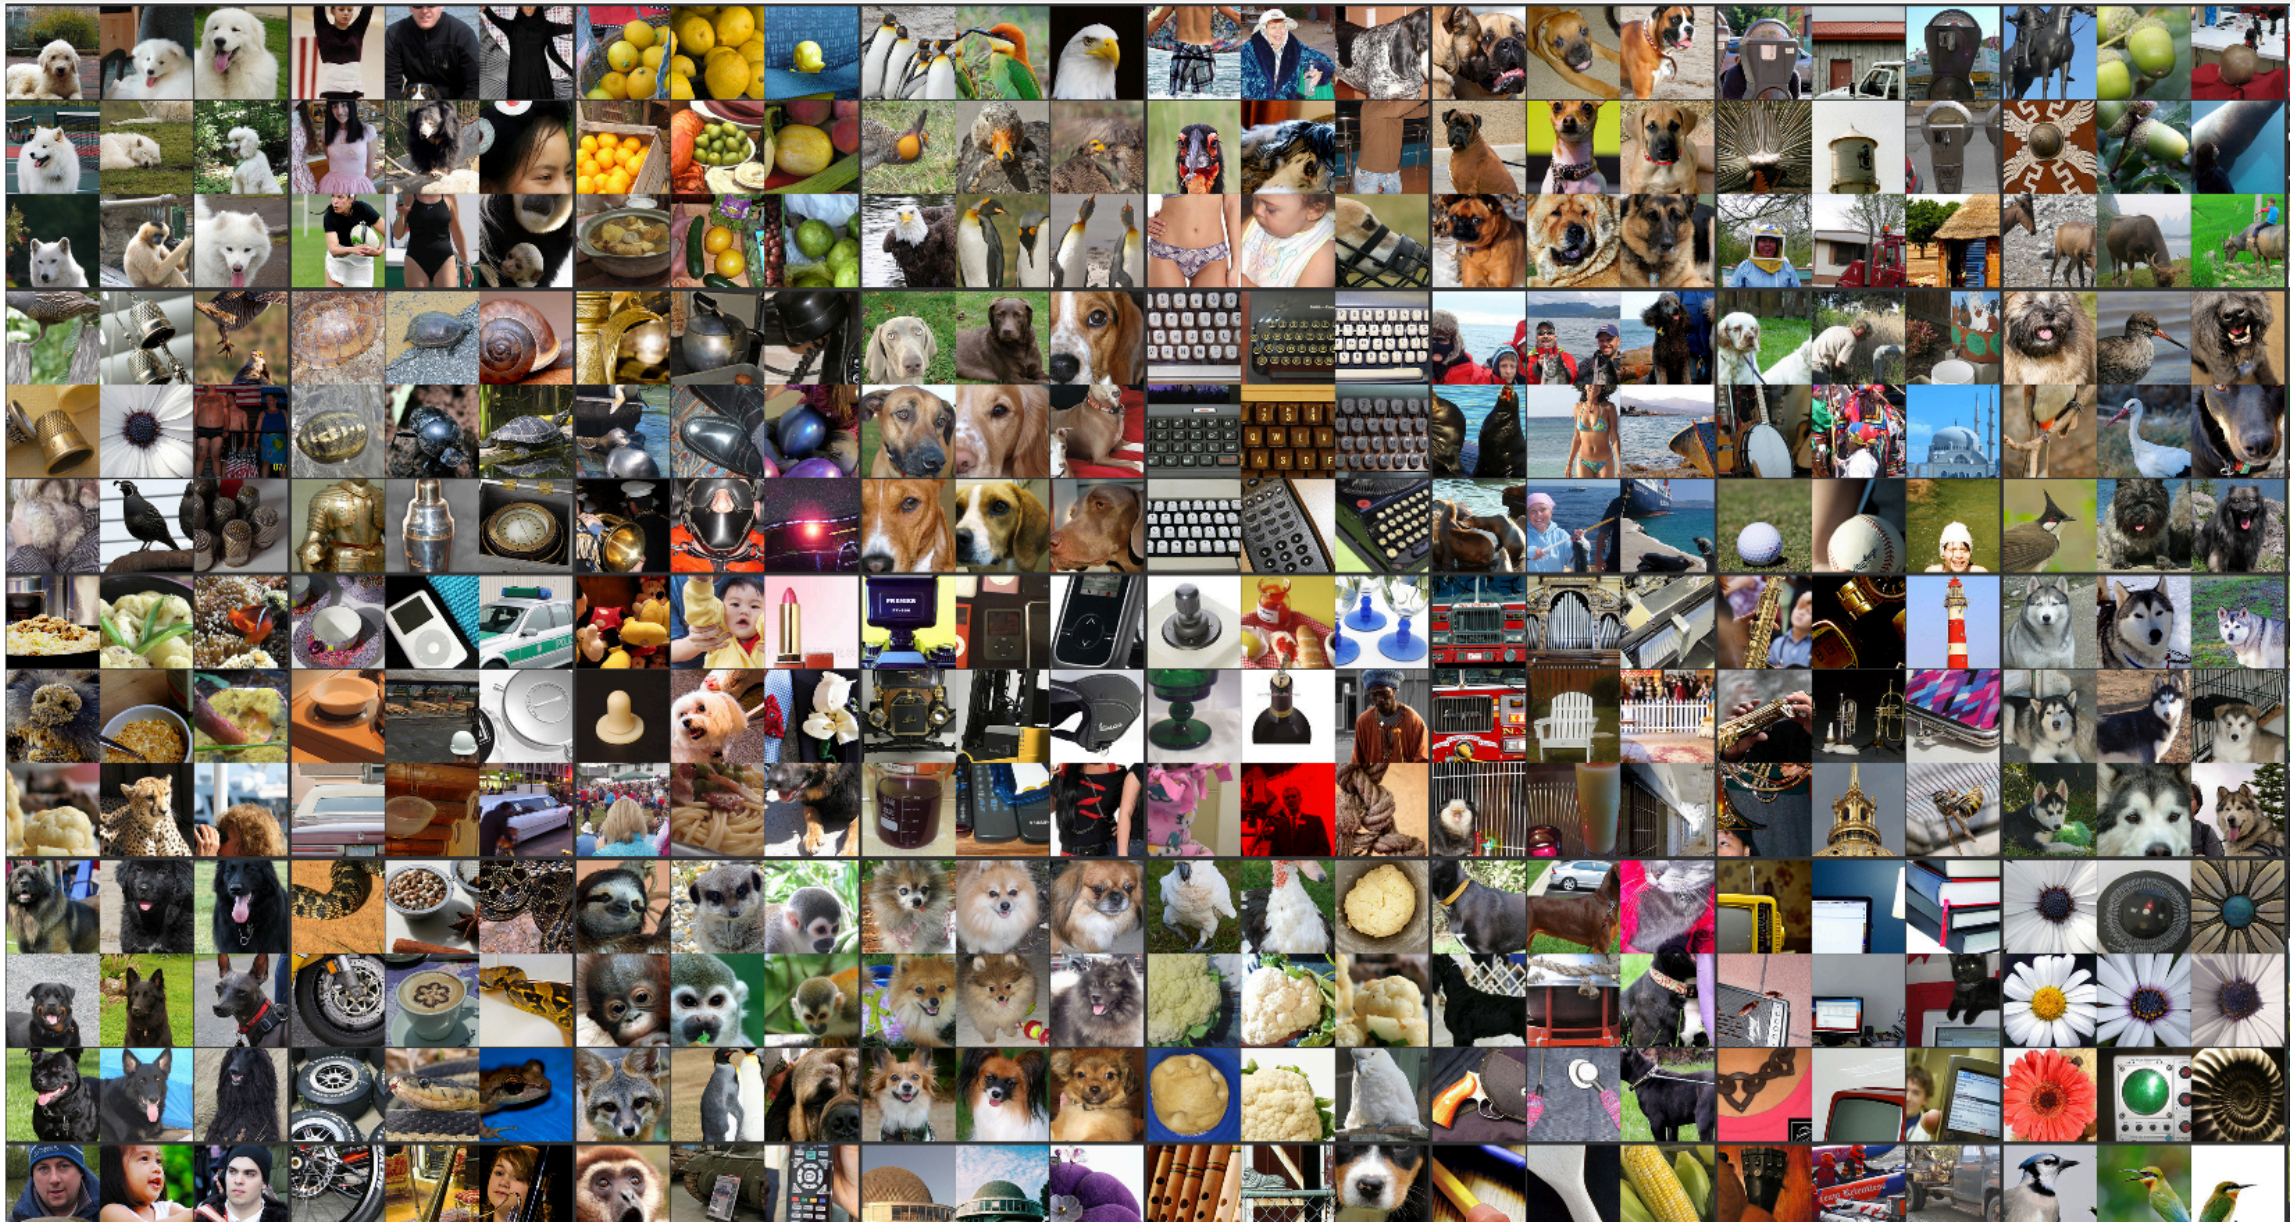

ImageNet Challenge

- Approx. 14 million labelled images, 20k classes
- Image gathered from Internet
- Human labels via Amazon Mturk
- Challenge:
  1.2 million training images and 1000 classes
- Metrics: top-5 and top-1 errors

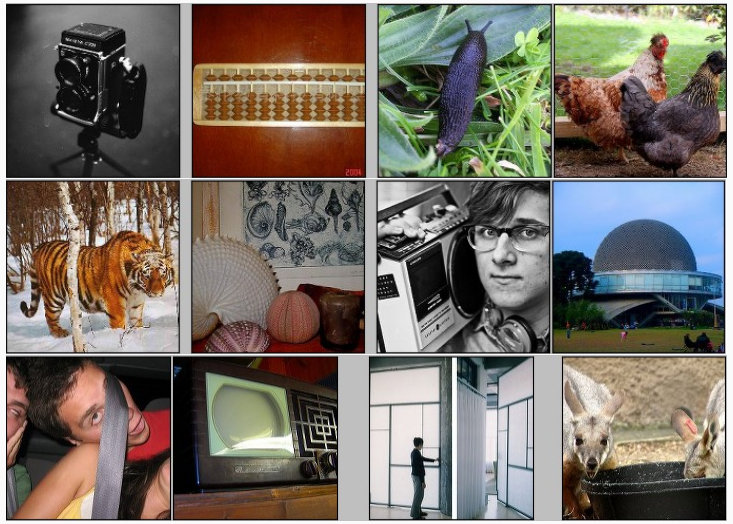

www.image-net.org/challenges/LSVRC/

Typical results

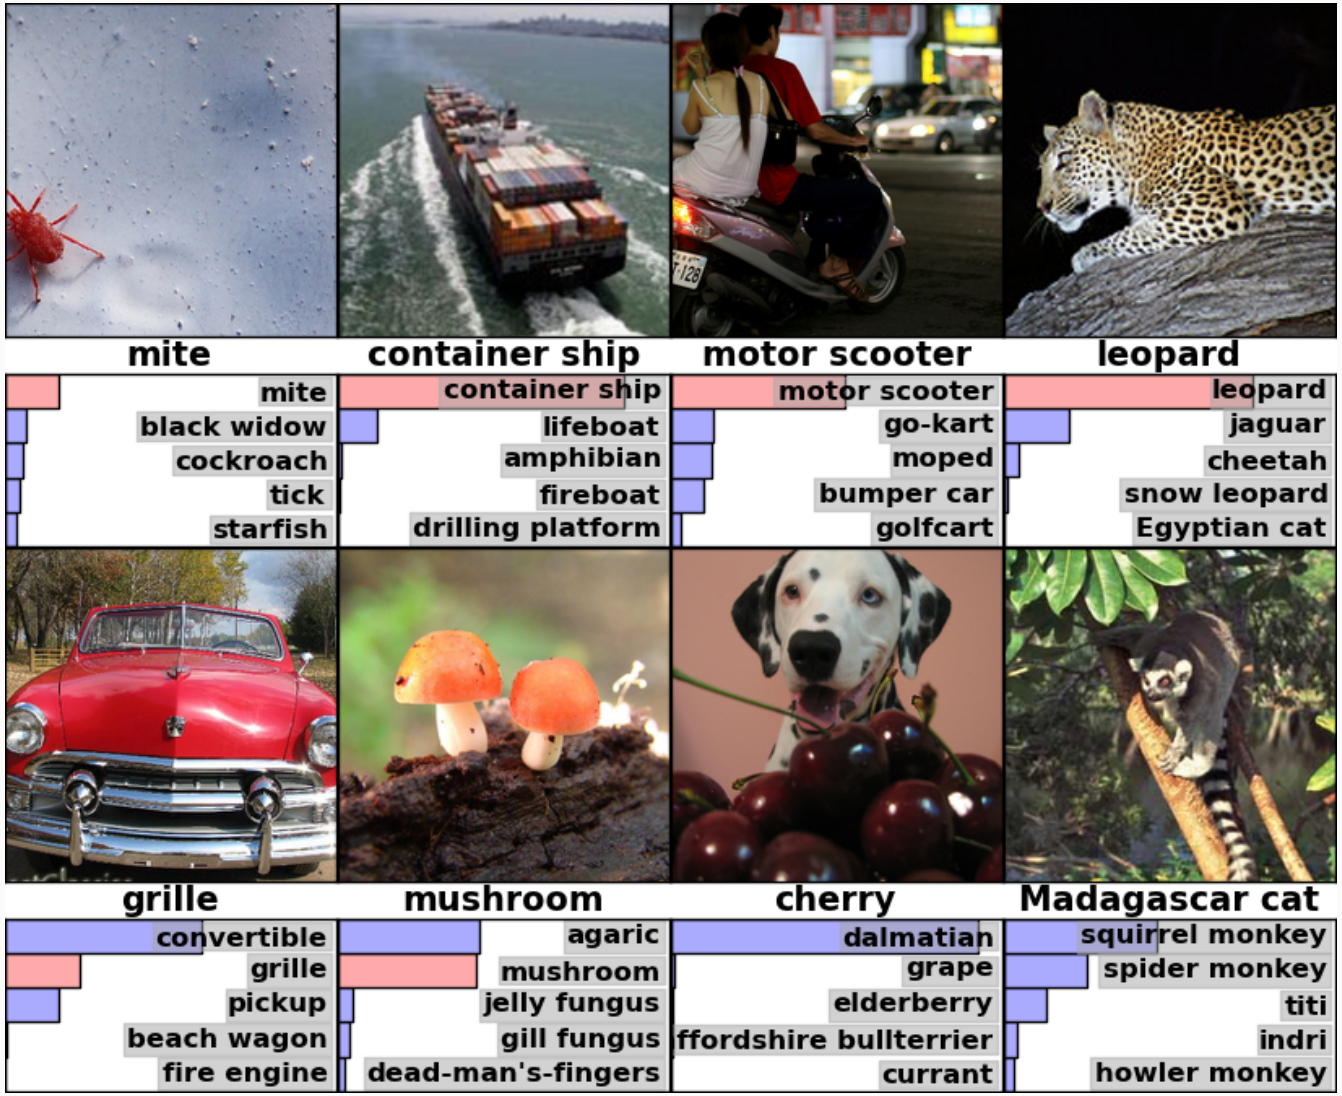

ImageNet classification with deep convolutional neural networks, Krizhevsky et al. NIPS 2012

ImageNet Challenge 2012-2013

Team	Year	Place	Error (top-5)	External data
SuperVision – Toronto (7 layers)	2012	-	16.4%	no
SuperVision	2012	1st	15.3%	ImageNet 22k
Clarifai – NYU (7 layers)	2013	-	11.7%	no
Clarifai	2013	1st	11.2%	ImageNet 22k

http://karpathy.github.io/2014/09/02/what-i-learned-from-competing-against-a-convnet-on-imagenet/

ImageNet Challenge 2014

Team	Year	Place	Error (top-5)	External data
SuperVision – Toronto (7 layers)	2012	-	16.4%	no
SuperVision	2012	1st	15.3%	ImageNet 22k
Clarifai – NYU (7 layers)	2013	-	11.7%	no
Clarifai	2013	1st	11.2%	ImageNet 22k
VGG – Oxford (16 layers)	2014	2nd	7.32%	no
GoogLeNet (19 layers)	2014	1st	6.67%	no
Human expert*			5.1%	

http://karpathy.github.io/2014/09/02/what-i-learned-from-competing-against-a-convnet-on-imagenet/

VGG-16 (Simonyan & Zisserman 2014)

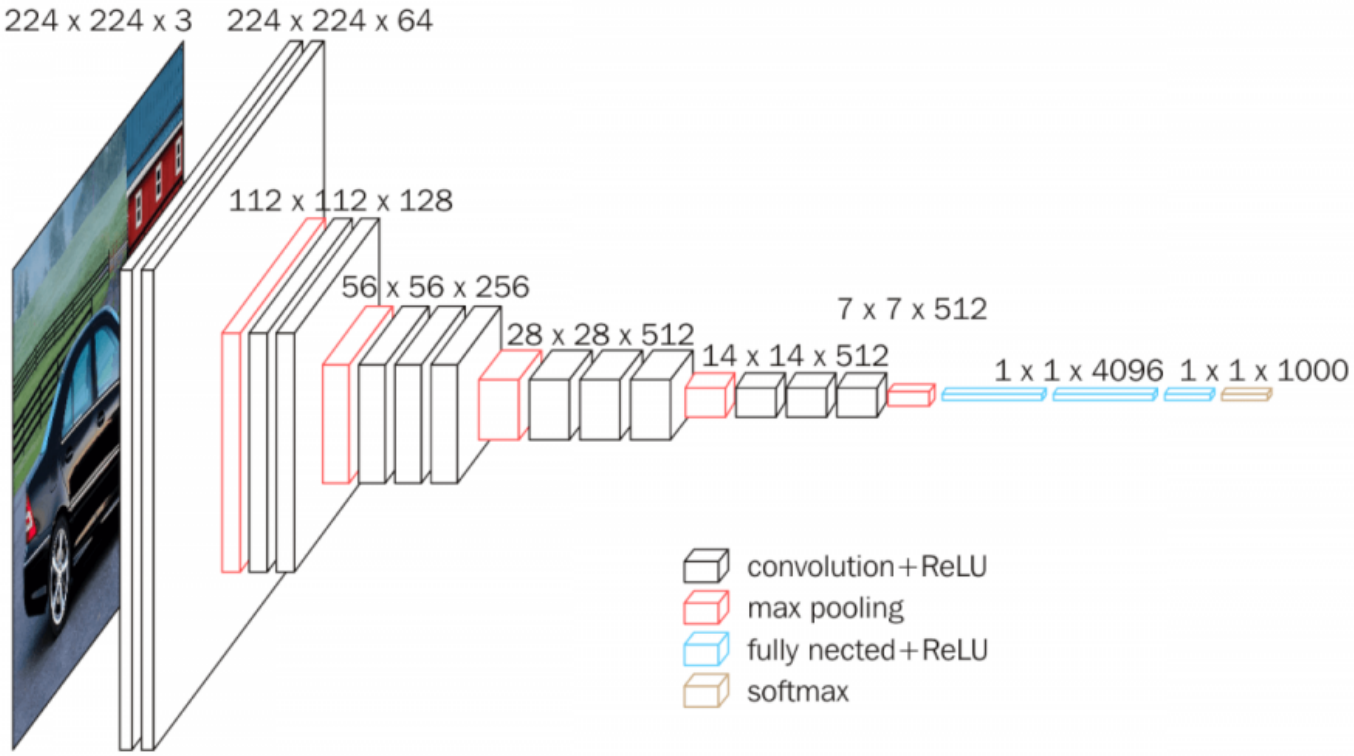

138 Million parameters
Top-5 error in ImageNet 7%

Very deep convolutional networks for large-scale image recognition, Simonyan et al. arXiv 2014

What about after 2014?

AlexNet, 8 layers
(ILSVRC 2012)

VGG, 19 layers
(ILSVRC 2014)

Deep Residual Nets

Revolution of Depth

AlexNet, 8 layers
(ILSVRC 2012)

VGG, 19 layers
(ILSVRC 2014)

ResNet, 152 layers
(ILSVRC 2015)

Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun, Deep Residual Learning for Image Recognition, arXiv 2015

Deep Residual Nets

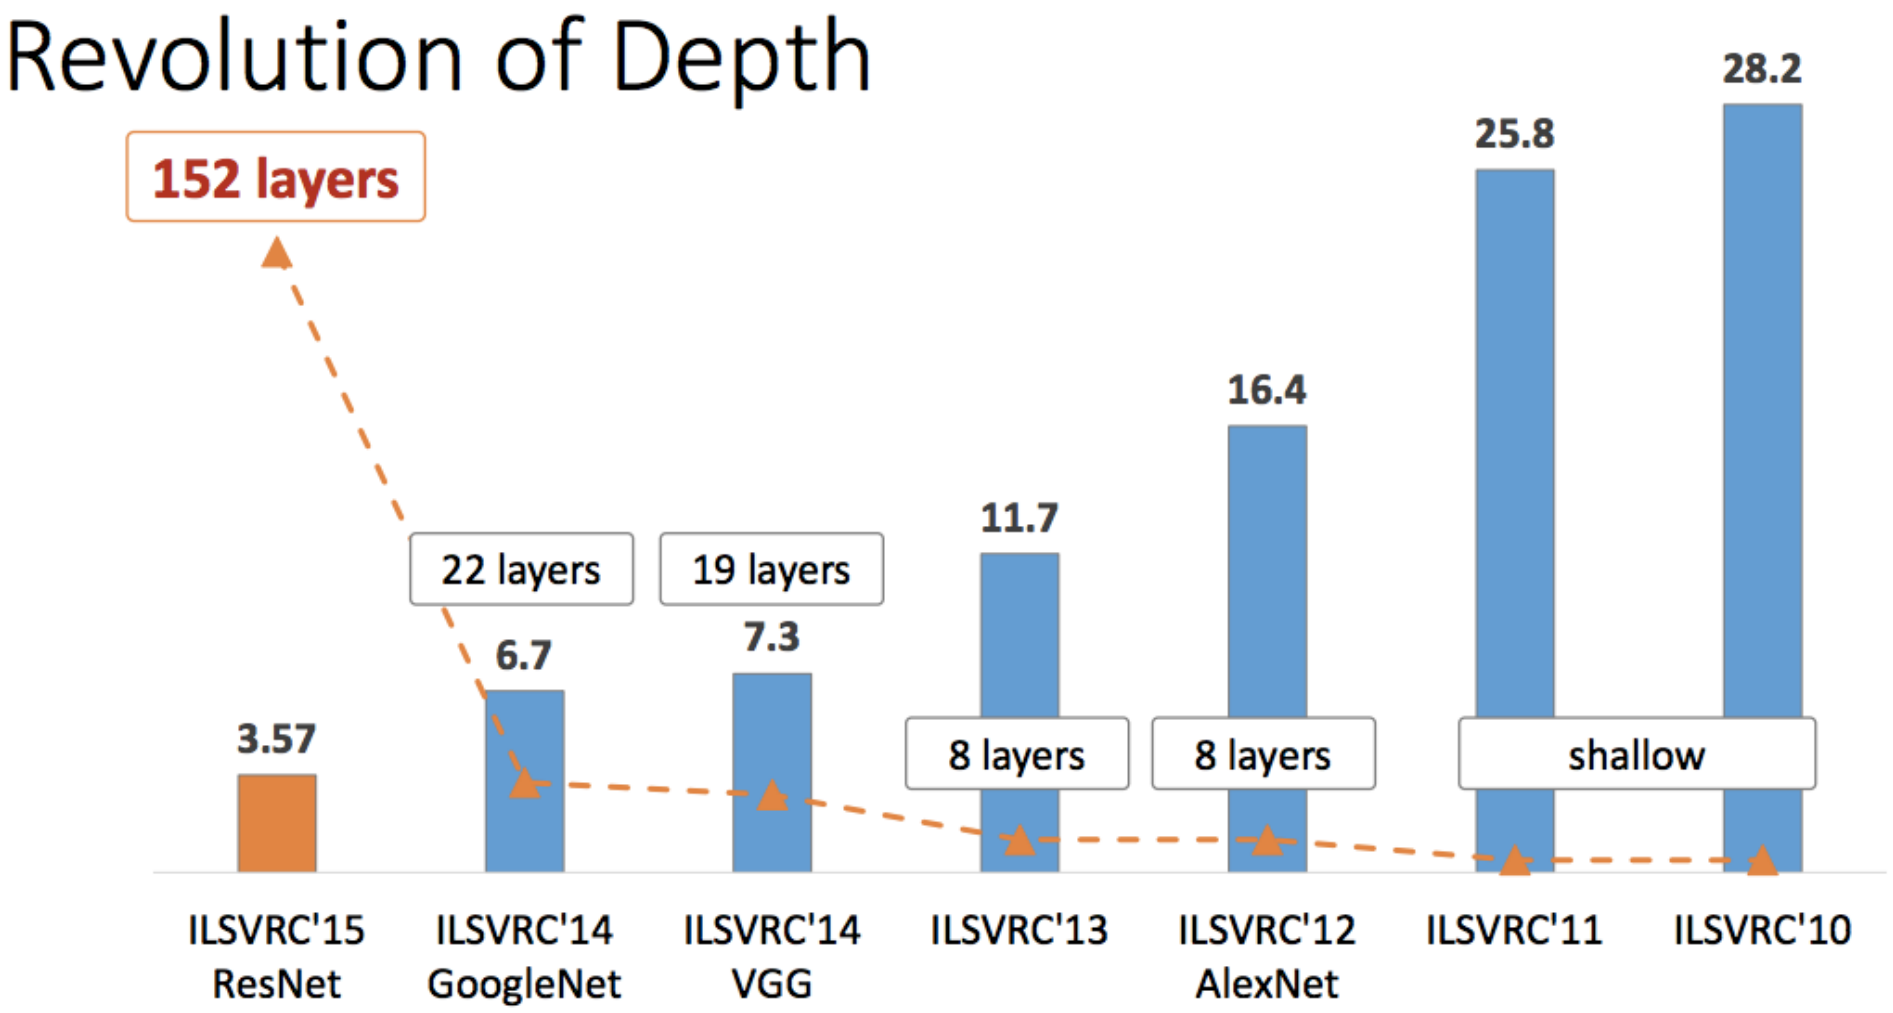

Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun, [Deep Residual Learning for Image Recognition](), arXiv 2015

Squeeze & Excitation

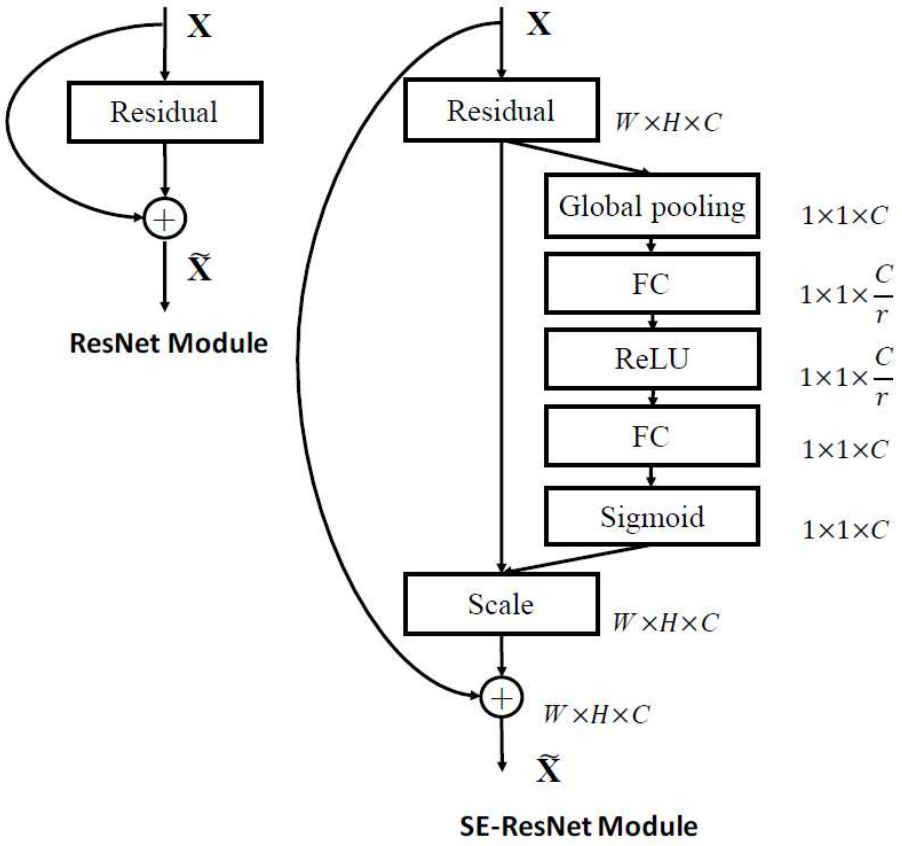

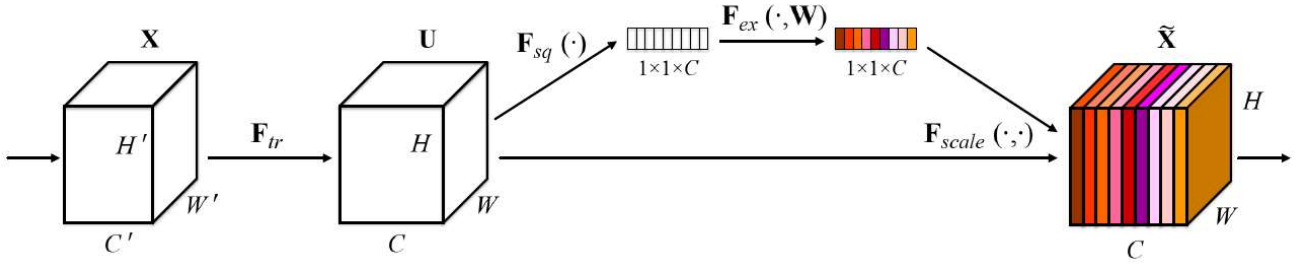

Hu et al. Squeeze-and-excitation networks CVPR 2018

ImageNet classification results (CLS)

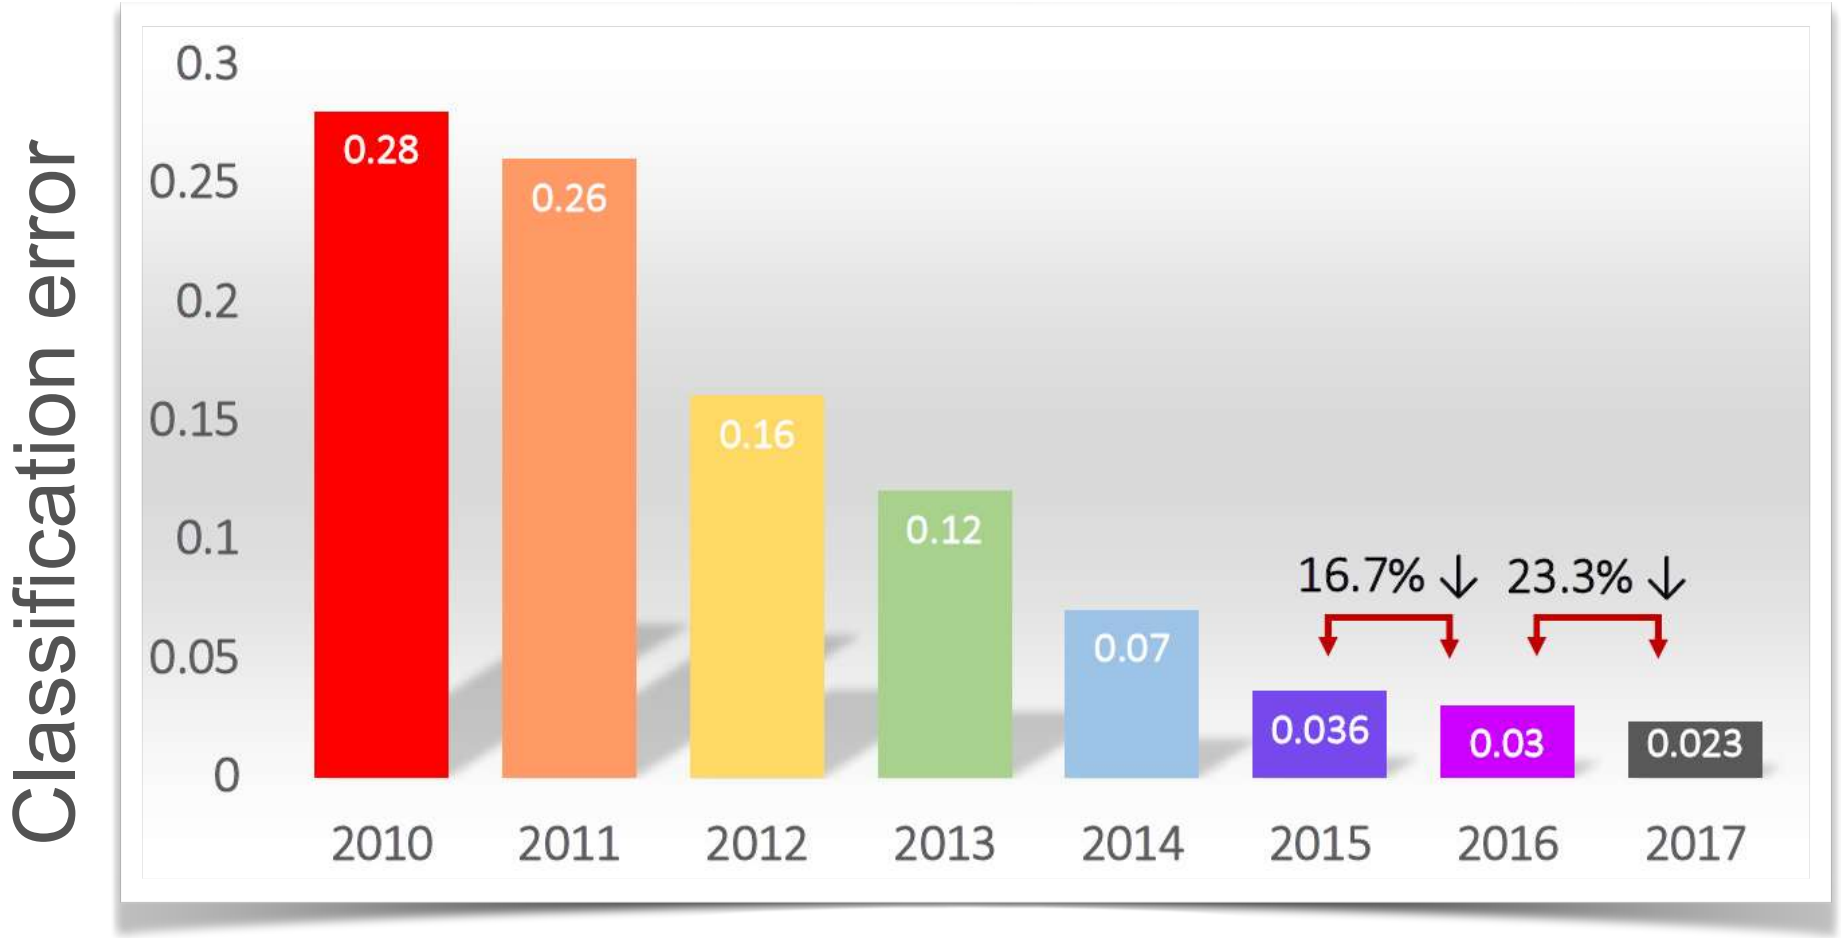

EfficientNet

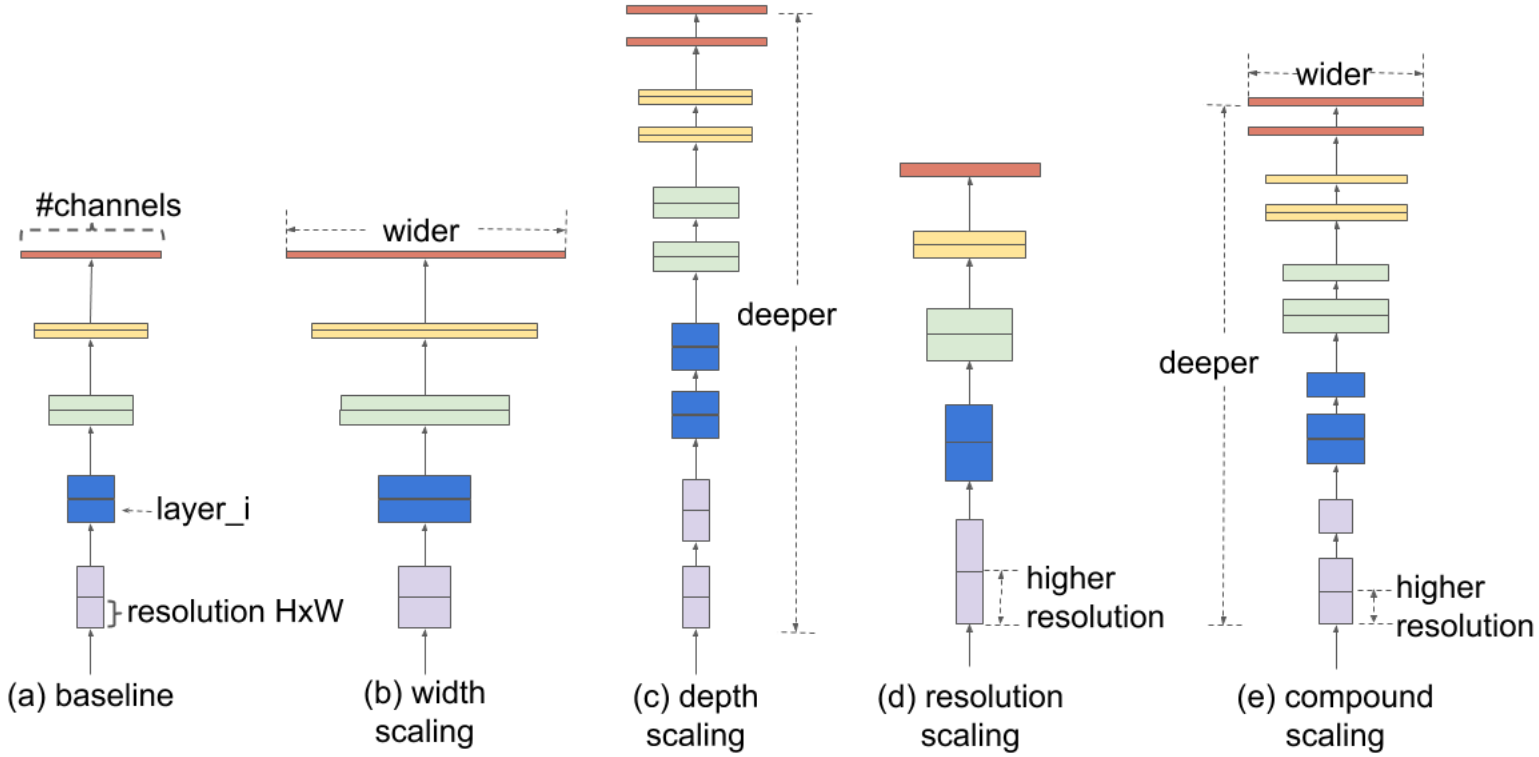

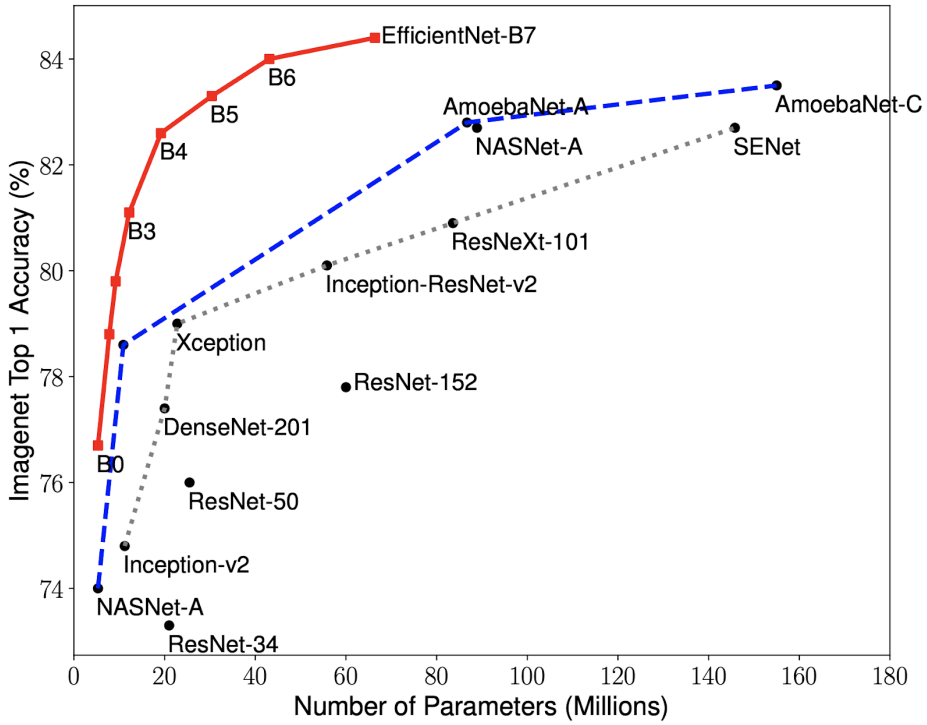

Tan et al. EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks, ICML 2019

Deep learning packages

- PyTorch
- TensorFlow
- Keras
- Caffe
- Theano
- Matconvnet
- …

https://en.wikipedia.org/wiki/Comparison_of_deep-learning_software

Breaking CNNs

Breaking CNNs

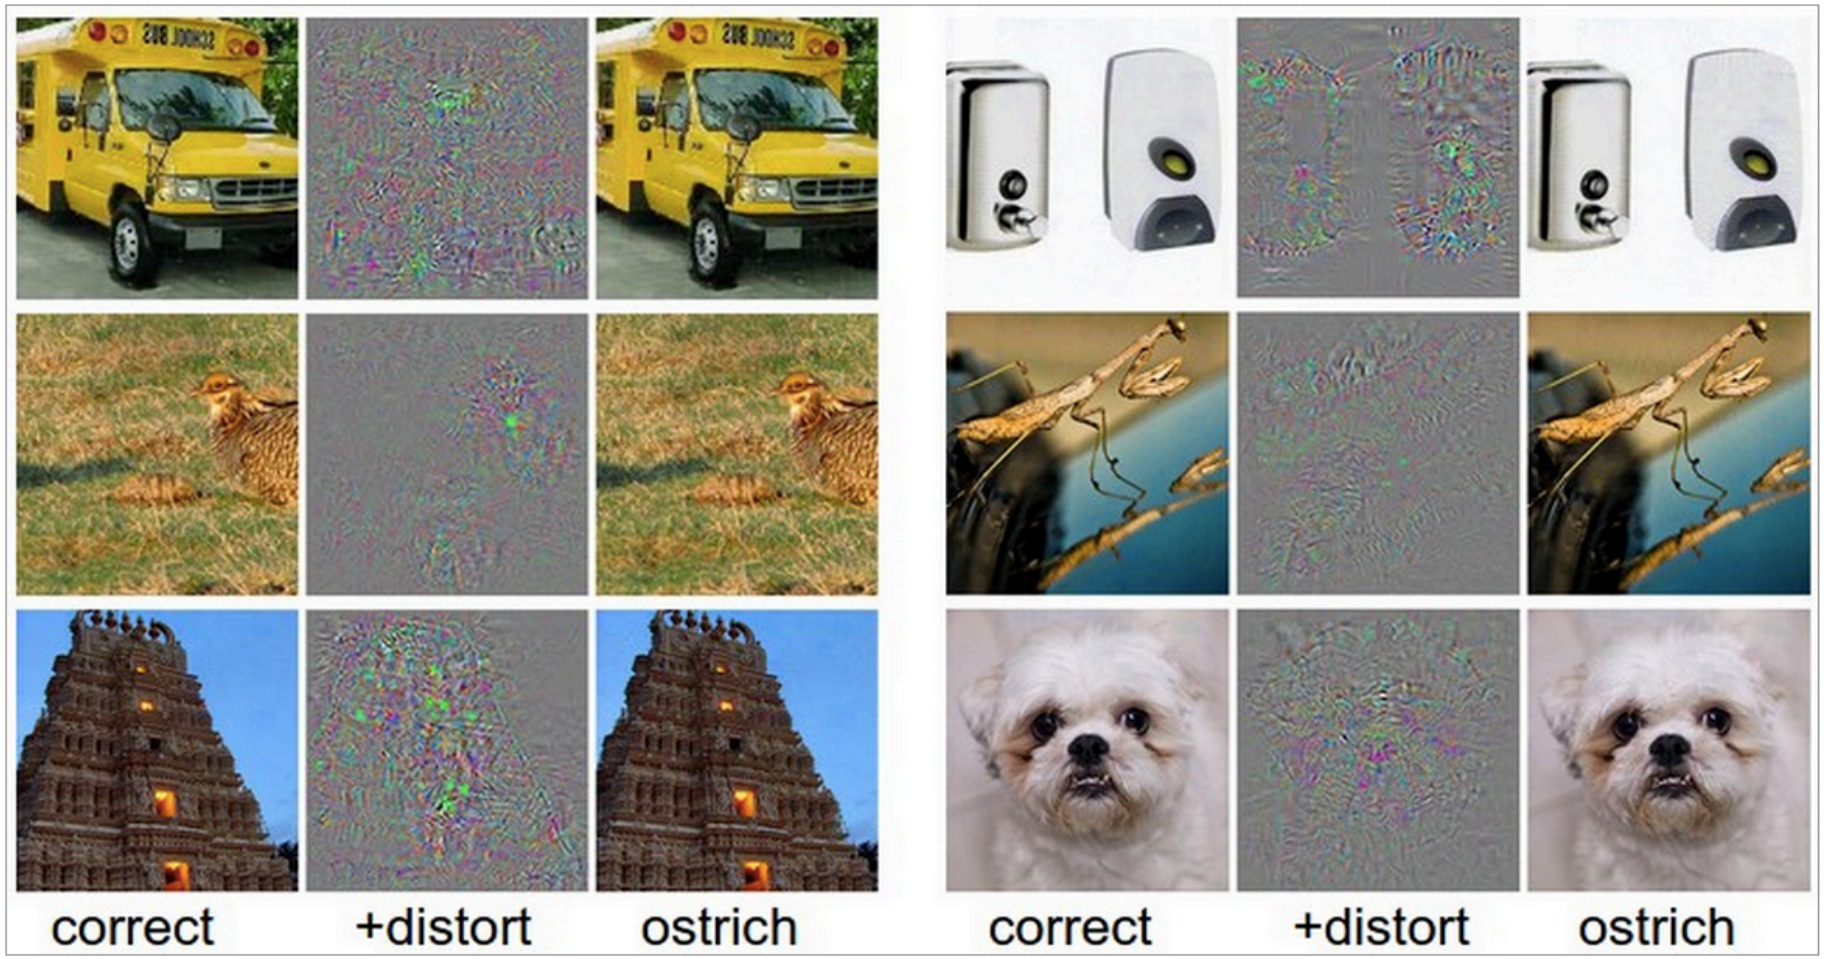

Take a correctly classified image (left image in both columns), and add a tiny distortion (middle) to fool the ConvNet with the resulting image (right).

http://arxiv.org/abs/1312.6199

http://karpathy.github.io/2015/03/30/breaking-convnets/

Breaking CNNs

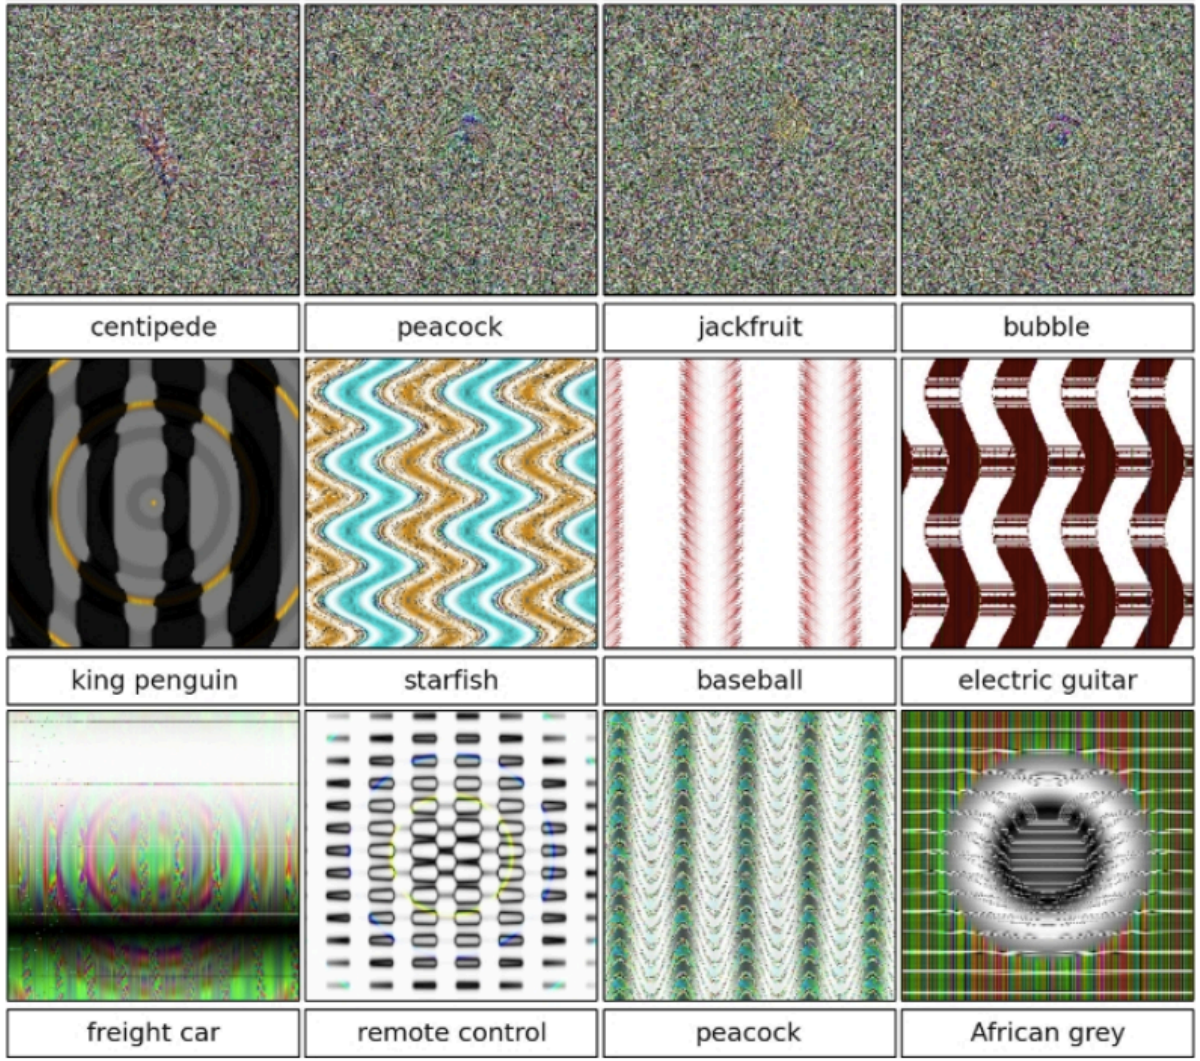

http://arxiv.org/abs/1412.1897

http://karpathy.github.io/2015/03/30/breaking-convnets/

What is going on?

- Recall gradient descent training: modify the weights to reduce classifier error

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial E}{\partial \mathbf{w}}$$

- Adversarial examples: modify the image to increase classifier error

$$\mathbf{x} \leftarrow \mathbf{x} + \alpha \frac{\partial E}{\partial \mathbf{x}}$$

http://arxiv.org/abs/1412.6572

http://karpathy.github.io/2015/03/30/breaking-convnets/

What is going on?

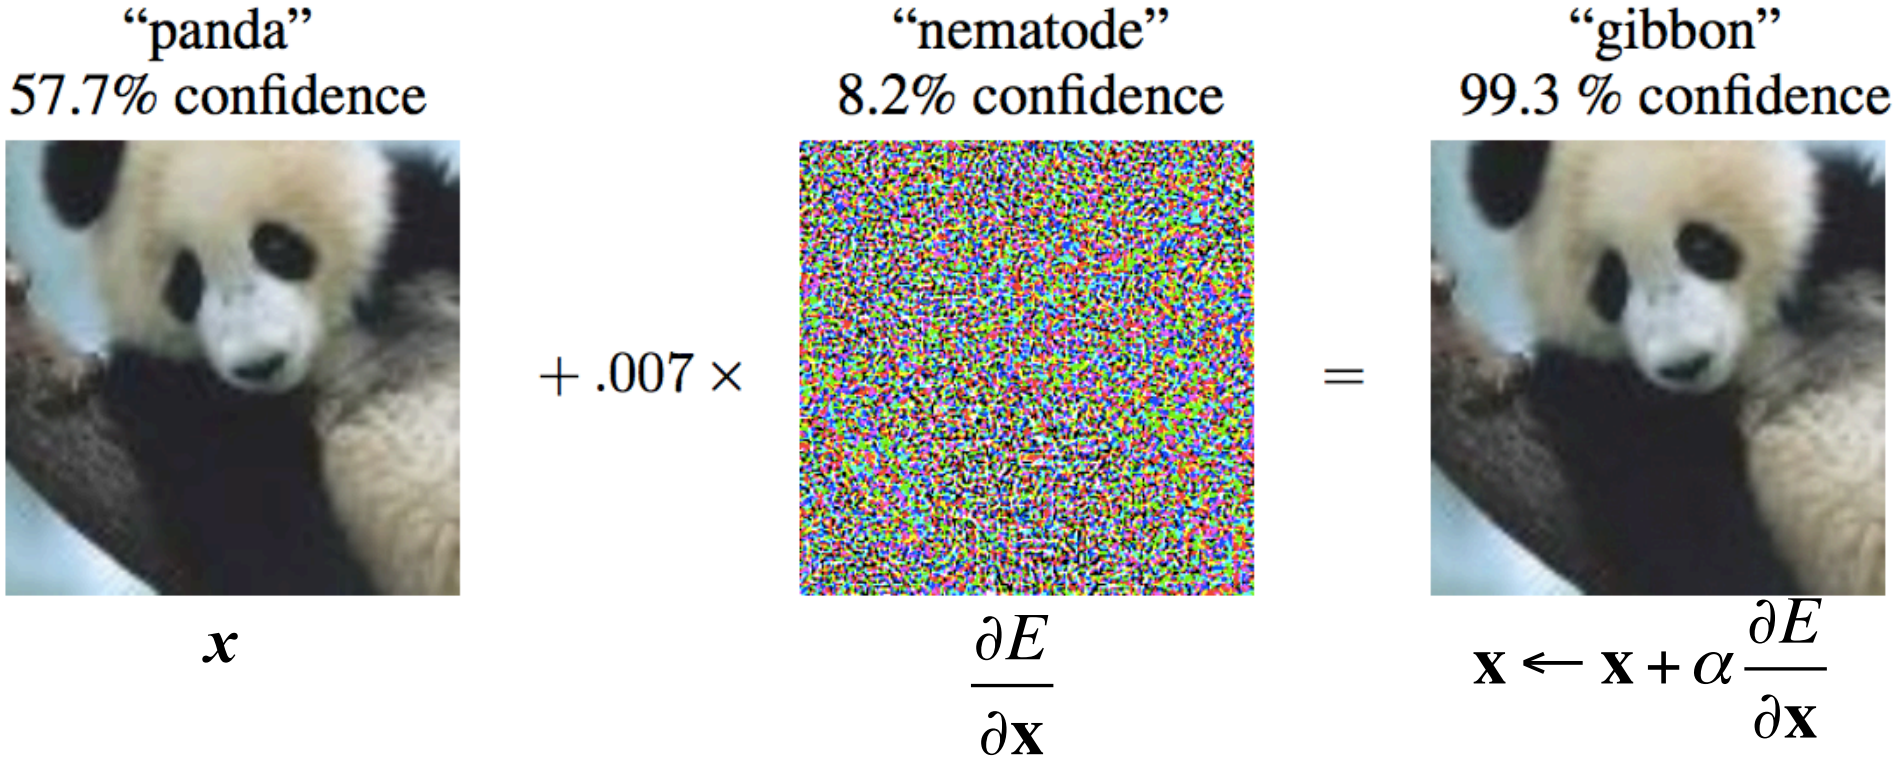

http://arxiv.org/abs/1412.6572

http://karpathy.github.io/2015/03/30/breaking-convnets/

Fooling a linear classifier

- Perceptron weight update: add a small multiple of the example to the weight vector:

$$w \leftarrow w + \alpha x$$

- To fool a linear classifier, add a small multiple of the weight vector to the training example:

$$x \leftarrow x + \alpha w$$

Fooling a linear classifier

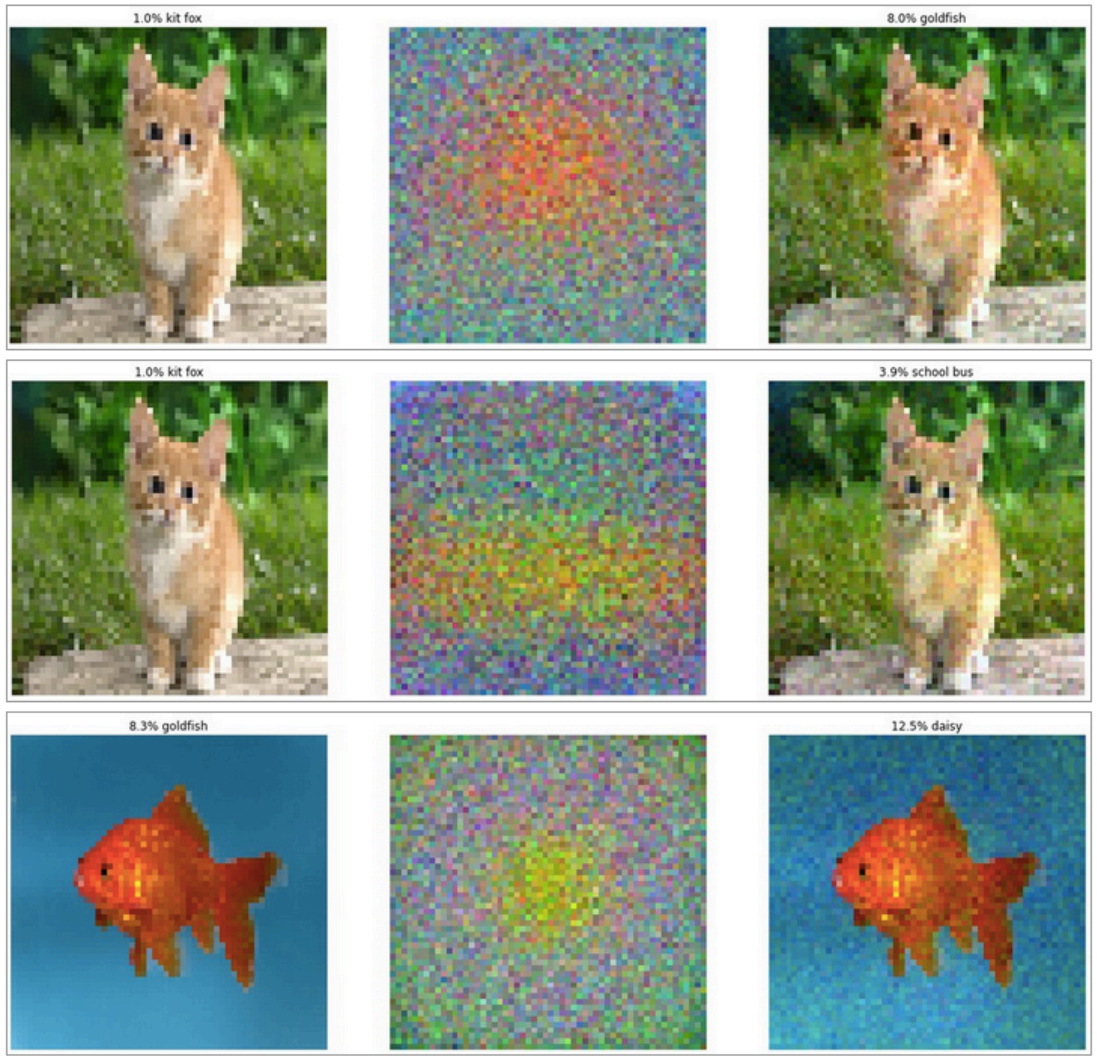

Fooled linear classifier: The starting image (left) is classified as a kit fox. That's incorrect, but then what can you expect from a linear classifier? However, if we add a small amount "goldfish" weights to the image (top row, middle), suddenly the classifier is convinced that it's looking at one with high confidence. We can distort it with the school bus template instead if we wanted to.

http://karpathy.github.io/2015/03/30/breaking-convnets/

Google DeepDream

- Modify the image to maximize activations of units in a given layer

https://github.com/google/deepdream/blob/master/dream.ipynb

Google DeepDream

https://github.com/google/deepdream/blob/master/dream.ipynb

Summary

- Currently most applications of deep neural networks in computer vision are based on supervised learning
- Need lots of annotated training data (and GPUs)
- Training via back-propagation takes time and is prone to local minima (many tips and tricks)
- Manual design of feature detectors and descriptors is replaced by manual design of network architectures (-> better automation needed)
- Lately there has been progress in weakly supervised and unsupervised learning (e.g. generative adversarial networks)

Further reading

- LeCun, Bengio, Hinton: Deep learning. Nature 2015.

http://pages.cs.wisc.edu/~dyer/cs540/handouts/deep-learning-nature2015.pdf

- Goodfellow, Bengio, Courville:

http://www.deeplearningbook.org/